How Shared Outputs Shape Source-Specific Evidence in In-Context Learning

Anonymous ACL submission

Abstract

An in-context learner can receive demonstrations from sources that use the same labels for different rules. How does it form the evidence needed to answer for one source? We find that the output vocabulary shapes learned relations during demonstration processing. Earlier answers can transmit one source's input–label associations into another source's evidence, giving that evidence a history that affects prediction. Causal interventions predict output-naming effects at the tested evidence interfaces and establish the historical pathway's contribution to ordinary judgments. This account also predicts different effects of source-separation policies. In shared-output tasks, selectively limiting answer-level sharing improves mean rule-use accuracy, while broader isolation can impair joint use. Public source–output information also affects separately formed rules before query reading, and equivalent formation backgrounds produce different joint usability. Controlled and real-text experiments across three model families, including independent 32B confirmation, connect these formation processes to source-conditioned behavior. The results explain how learned relations acquire a source provenance and how their formation constrains subsequent evidence selection.

1 Introduction

A review can praise a restaurant's food while criticizing its service. Two annotators may use the same binary labels but apply different judgment rules to that review. An in-context learner receives both annotators' examples and is asked to predict for one of them. It must infer the relevant input–output relation and use it on a new comment. Similar questions arise whenever demonstrations encode several tasks, conventions, or perspectives in one context. Understanding which evidence governs the answer requires understanding how those demonstrated relations become available to the model.

Cho et al. (2025) traced how input information enters label states and is retrieved by a query. Bakalova et al. (2025) showed that cross-demonstration contextualization shapes task information before aggregation. In a multi-source context, these processes create a distinction between a demonstration's positional owner and the history of relations represented there. Source A's tokens can contain information influenced by B's examples. The fact that a query reads A's positions leaves open how B's supervision shaped what it reads.

This distinction motivates our question: **how do shared output identities affect the formation and use of source-specific evidence?** Output words participate in the representations formed during demonstration processing. When different sources share them, changing one source's rule may change the evidence available at another source's positions. At the same time, surrounding context can supply information that helps several rules work together. An account of source-conditioned ICL should explain both the relations entering that evidence and the effects of sharing or separating them.

We change one source's demonstrated relation while holding the other source and the query fixed. Whole-query access restrictions and message interchanges distinguish positional access from causal dependence on a rule. A historical answer-to-label pathway predicts output-namespace counterfactuals at these interfaces in new contexts with a fixed coefficient; removing the operation tests its contribution to native predictions.

We then fix output words, positions, and input multisets while changing input–label associations. These counterfactuals identify relational content carried through the pathway under full-source reading and on real sentiment text. Formation yields different functional predictions for selective answer restriction, complete private-history isolation, and query-only restriction, which we test for both requested sources.

Finally, public source–output relations can help independently formed rules through formation, and information-equivalent public backgrounds yield different joint usability. Independent Qwen3-32B confirmation connects formation, native operation, and real-text function. Together, the experiments establish three findings: output identities affect historical formation, messages carry learned associations, and formation and reading impose different functional constraints. Source-conditioned computation depends on the provenance and use of learned relations, including when the model answers correctly.

2 Measuring which source's relations influence a prediction

2.1 Sources, rules, and materials

The current sentiment and dual-aspect experiments contain eight demonstrations from each source and eight novel queries. In shared-word conditions both sources use `no/yes`, with private label polarity balanced across contexts; separate-word conditions keep A's words and give B `pass/flag`; earlier Measuring Hate Speech experiments use `safe/toxic`. A's raw examples and labels are held fixed; paired counterfactuals change B. A single-aspect sentiment task separates association from output naming. A dual-aspect task independently varies food and service sentiment, B's inferred criterion, and its private output convention. Original SST5 reviews supply a natural-text test with independently sampled demonstrations and official held-out queries.

All models are frozen: Qwen3-8B, Qwen3-14B, and Qwen3-32B post-trained hybrid checkpoints, Mistral-7B-v0.3 base, and Llama-3.1-8B base. We score the serialized task as a continuation using FP32 eager attention. Confirmation uses new materials after discovery, without selecting models, contexts, or seeds by correctness. Model revisions, actual token streams, and complete split rules are recorded in the appendix.

2.2 Positional access and functional provenance

For a fixed common prompt background, let D_A and D_B denote the demonstrations owned by sources A and B. H_A and H_B denote their contextualized key/value states, grouped by positional ownership. We call these available demonstration states evidence. Formation is their computation during prefix processing, and reading is their subsequent use by query tokens. Let q denote the requested source and new input, and z the output-label logit difference. We write

$$H_A = F_A(D_A, D_B), \quad H_B = F_B(D_A, D_B),$$
$$z = R(q, H_A, H_B). \tag{136}$$

F_A and F_B denote prefix computations at the two source groups, and R denotes query computation under the specified access graph. For A's no/yes outputs, z is yes minus no. Relational provenance is measured as causal dependence on a source's demonstrated input–label association. We change B's pairing while fixing A and the query, then prevent every query position from accessing B positions. A remaining B-rule effect reaches the prediction through the available contextualized graph.

Under this restriction, changing B's rule changes 17.6% [12.1, 23.8] of A's answers in 32 independent Qwen3-8B Measuring Hate Speech contexts (E90). Separating B's output vocabulary reduces the inherited effect even though A's raw demonstrations and labels are unchanged. Source identity also has native causal effects in earlier interventions. Thus the result concerns the provenance of contextualized evidence, rather than a model completely ignoring source names.

2.3 Evaluation and reproducibility

Intervals use 10,000 context-bootstrap samples, with contexts as the sampling units. We retain every predefined condition and check consumed tokens and masks, cached/full-token scores, self interventions, and graph zeros. Independent analyses reconstruct raw quantities and gold/access layouts before reading summaries. Appendix A specifies serialization, analysis, and reproducibility details.

3 How output identities shape historical evidence

3.1 A frozen formation-path prediction

We compare shared-output and separate-output prefixes, keeping A's examples and output words identical. Counterfactual donor states replace historical B-answer-region K, V, or KV **only in messages consumed at A's label positions** during prefix processing. Subsequent states evolve normally. Interchanging the resulting A-label caches into one otherwise pristine recipient tests the specified formation interface.

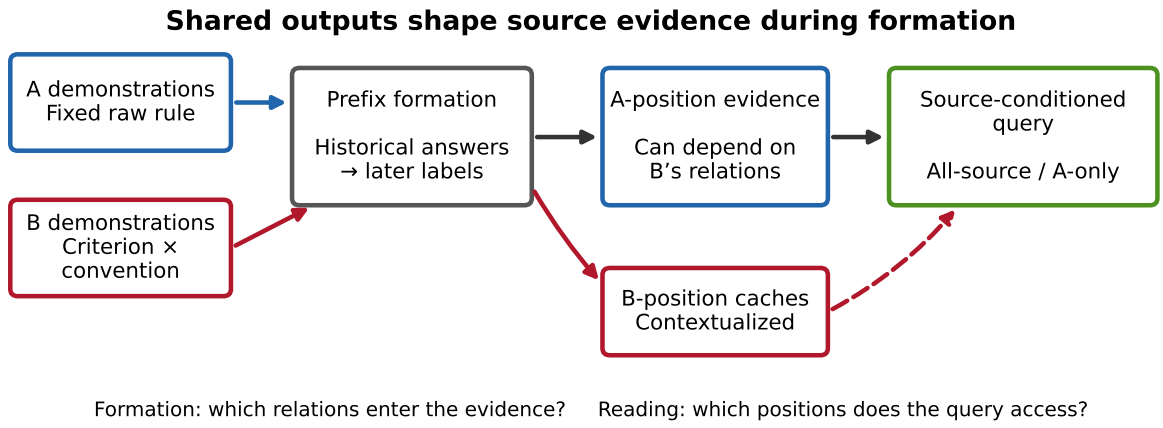


Figure 1: Historical formation creates contextualized evidence for both sources. Changing B can affect evidence at A positions, while full-source reading can also access B-position caches. Arrows denote tested causal interfaces and information access.

The prediction has coefficient one: the message-intervention effect predicts the actual namespace counterfactual at that interface. No head, layer, strength, or regression coefficient is selected. We report raw effects and rule double contrasts, which subtract the namespace effect between two B-association worlds. This separates relation-dependent naming effects from a common output shift.

In new contexts, K reconstruction gives rule-double-contrast correlations .969 and .966 for Qwen3-8B and Mistral on the controlled Measuring Hate Speech task; KV gives .985 and .977. The same frozen mechanism remains predictive on controlled sentiment and real SST5 text. These are actual-donor causal reconstructions: their target is a specified interface effect, and residuals are measured against the coefficient-one identity prediction.

3.2 Testing the native operation

Reconstruction alone does not establish native contribution. We therefore block historical B-answer access at A-label receivers during ordinary prefix processing, using the same fixed mask at every layer. There is no donor activation. A complementary test blocks all B-to-A prefix access, then reopens only the historical answer-to-label path. Within-A answer removal and single-source private functions control for disturbance to general historical learning.

For each context, the signed B-rule response averages A-correct-label-oriented query margins, D = mean[g_A(q)(z_cross(q) − z_origin(q))]. Removal reports D_native − D_blocked; restoration reports D_restore − D_all-B-off. Targeted removal reduces B-rule influence in every checkpoint. Restoration increases it under A-only reading, with negative, positive, or uncertain full-source effects. The operation therefore has a reading-dependent contribution. Masks renormalize attention and later states remain live. Own-history removal also impairs Mistral's private function; Qwen's cross-versus-within specificity interval includes zero (Appendix D).

4 Historical messages carry learned associations

4.1 Fixed-word relational counterfactuals

Namespace changes might alter output-related signals without changing learned relational content. We therefore hold B's label stream, source names, positions, and input multiset fixed, and reassign comments across opposite sentiment classes. A matched within-class reassignment moves the same number of inputs without reversing the association. A's raw evidence and novel queries stay fixed; paired token geometry is matched before model scoring.

For each world, we measure the correct-label-oriented query effect, then compare its magnitude between cross-class and within-class reassignment. Historical answer KV has an excess effect of .618 [.446, .813] nats in Qwen and .404 [.313, .499] in Mistral under A-only reading. V carries substantial content, while the namespace reconstruction is more accurately reproduced through K/KV. This

Model, n=32	Full-source operation removal	A-only restricted restoration	Full-source restricted restoration
Qwen3-8B	1.649 [1.108, 2.250]	1.455 [0.999, 1.945]	-0.806 [-1.365, -0.266]
Mistral-7B-v0.3	1.054 [0.886, 1.227]	0.647 [0.463, 0.835]	0.140 [0.053, 0.231]
Llama-3.1-8B	0.808 [0.517, 1.109]	0.958 [0.635, 1.276]	-0.233 [-0.336, -0.122]
Qwen3-14B	1.329 [0.967, 1.717]	1.120 [0.673, 1.646]	0.039 [-0.217, 0.287]

Table 1: Native operation removal and restricted restoration, n32 confirmation. Signed B-rule contrasts in nats with context-bootstrap 95% CIs.

dissociates the intervention's relation content from its output-identity effect without assuming independent internal components.

The relation effect also holds when queries can read both sources. Moreover, crossing A-position caches formed under shared/separate namespaces with B-position caches from the same backgrounds changes native B-rule dependence. The signed interaction is positive in the three smaller checkpoints and negative in Qwen3-14B ($-.686$ [-1.133, $-.237$] nats). The model therefore consumes jointly formed evidence; its source caches are not independent stores whose effects can be assumed additive.

The answer-region pathway reproduces part of the relation effect. Full label interchanges have additional effects, and the signed relation contrasts vary by model. These distinctions motivate measuring complete functional behavior in the next experiment, rather than equating a message effect with a fully transferred rule.

4.2 Reliable private conventions on natural text

We repeat the frozen 43-condition engine on original extreme-polarity SST5 reviews. Demonstration pools are disjoint text-hash partitions of training; discovery and confirmation queries use official validation and test splits. Gold sentiment and token geometry determine construction, without screening model outputs. Positive and inverted private conventions have single-source accuracies of 88.3–99.2% across the four checkpoints (Appendix D).

The same coefficient-one KV formation reconstruction gives rule correlations .931/.948/.930/.966 in Qwen3-8B, Mistral, Llama, and Qwen3-14B, respectively. Fixed-output association excess magnitudes are positive in every family. Llama's signed relation contrast is negative ($-.109$ [$-.194$, $-.024$] nats): transmitted content is not uniformly a correction of the requested rule. SST5 provides a functionally usable natural-text replication; these constructed conventions do not recover individual human preferences.

5 How formation constrains source-conditioned rule use

5.1 Competing formation and reading policies

The formation account predicts different consequences for three operations: selectively limiting other sources' answer messages, removing all cross-source formation, and restricting only query access. The first tests whether limiting answer sharing protects private mappings while retaining other contextual paths; the second can also remove useful shared information; the third acts on evidence already formed.

We fix three prefix policies—native, symmetric other-source answer removal at label receivers, and all cross-source removal—and two whole-query scopes. Both source requests use the same prefix. On dual-aspect comments, B's food/service criterion and same/opposite private convention vary independently. The primary averages two-candidate accuracy over both sources, worlds, and queries within each context. Single-source functions and a one-sentence instruction are evaluated separately.

5.2 Independent functional confirmation

Selective formation improves mean accuracy by 2.93 [1.42, 4.35], 3.42 [1.17, 5.76], 7.37 [5.37, 9.43], and 2.78 [0.49, 4.79] points in Qwen3-8B, Mistral, Llama, and Qwen3-14B. It outperforms complete cross-source formation removal in every checkpoint. Gains concentrate on conflicting correct labels; agreeing-label cases lose 2.34 [0.78, 4.00] points in Qwen3-8B and 3.61 [1.56, 5.86] in Qwen3-14B. Individual-source gains are not uniformly confirmed. Query-only isolation gives larger gains in Llama and Qwen3-14B, so selective formation is not a universally optimal policy.

Complete formation isolation can increase convention influence under full-source reading. Combining it with source-only reading gives the graph-

required zero B-world effect on A. These policies use known source ownership and remove unequal edge sets; they test functional consequences. Appendix C retains all policies, source-specific scores, agreement tradeoffs, and private controls.

5.3 A prospective test on real-text source functions

We next test the functional policy prediction on new SST5 contexts. A and B infer the same sentiment criterion but can adopt opposite private conventions. Origin and cross-convention worlds receive equal primary weight; a within-class movement remains a separate control. The same three formation policies and two query scopes are evaluated for both sources from each common prefix. A keeps `no/yes`; B uses either those shared words or `pass/flag`. Single functions and one-sentence source instructions are retained in all 85 conditions.

Selective formation improves the shared-word primary in all four checkpoints, with independently positive A and B gains. Candidate and full-vocabulary primary scores coincide here. Benefits concentrate on opposite conventions; same-convention effects are small and uncertain. The predefined B-convention response also decreases under selective formation while A's output words remain fixed (Appendix F), connecting the accuracy change to relational influence. Query-only isolation remains more effective in Llama and less effective in the two Qwen checkpoints; Mistral's point ordering is uncertain.

Separate words expose a further distinction. Complete cross-source formation isolation retains about 94–97% own-word candidate accuracy, yet full-vocabulary accuracy falls to about 51–54%: the model frequently outputs the other source's word. Thus learning within a label space and using that convention jointly require separate measurements. Native separate-word candidate accuracy is already high, so its smaller selective gains also reflect headroom and word priors. Mistral loses 1.27 [0.20, 2.44] full-output points under selective formation despite unchanged candidate accuracy. Appendix F preserves every namespace, private direction, source, actual output, and instruction control.

5.4 Association influence and criterion-related function

Source-conditioned use also depends on what a learned association does to a new input. To test this, we independently reverse B's private convention under food and service criteria, holding A fixed. Responses are projected onto food, service, bias, and their interaction. The resulting two-by-two criterion–feature matrix distinguishes changing the amount of influence from changing its functional direction.

In Qwen3-14B, native food and service selectivity are .282 [.199, .371] and .404 [.294, .528] nats. A prospective within-context normalized determinant is .410 [.330, .493], rejecting a restricted account in which both criteria only rescale one common functional direction. With the query prohibited from reading B, it is still .190 [.073, .310]. Local label and answer-message interfaces yield different responses. Historical answer KV only partially reconstructs the **actual label-interface** criterion counterfactual (r=.668, RMSE=.646, zero RMSE=.746).

Thus association transfer, full-history criterion dependence, and correct source-conditioned use are distinct measurable properties. The complete response matrices and negative effects are retained in the appendix; a positive aspect response alone is not sufficient evidence for a complete abstract criterion program.

6 Coordinating independently learned source rules

6.1 Sharing identities while keeping supervision private

What can be shared without exchanging private supervision? With cross-source demonstration formation blocked, we cross access to an unlabeled source-name list during formation and whole-query reading, using unrelated names as control. Qwen3-8B has a formation-specific full-output gain of 4.59 [2.73, 6.64] points, with uncertain candidate gain; Mistral's effect is negative and Llama's uncertain (Appendix G). Names alone therefore provide a partial, model-dependent account of the missing shared information.

6.2 Sharing an output schema before private rule learning

We next share a source-to-word-set table before independent formation, for example Alex: no, yes; Sam: pass, flag. It specifies output spaces, without review–label pairs or private polarity. An unrelated-name table has the same words and structure; a swapped table reverses the real sources' word as-

Model, confirmation n=32	Shared native accuracy, %	Selective formation, %	Selective-minus-native, pp [95% CI]	Query-only-minus-native, pp [95% CI]
Qwen3-8B	84.86	92.19	7.32 [5.37, 9.28]	-0.39 [-3.12, 2.44]
Mistral-7B-v0.3	88.77	94.14	5.37 [2.93, 7.71]	5.96 [3.22, 8.59]
Llama-3.1-8B	81.25	90.72	9.47 [7.52, 11.52]	16.70 [14.36, 19.04]
Qwen3-14B	89.94	97.36	7.42 [5.57, 9.38]	3.81 [1.27, 6.25]

Table 2: Real-text private rules with shared words: policy consequences, confirmation n32. Percent accuracy and paired changes with context-bootstrap 95% CIs.

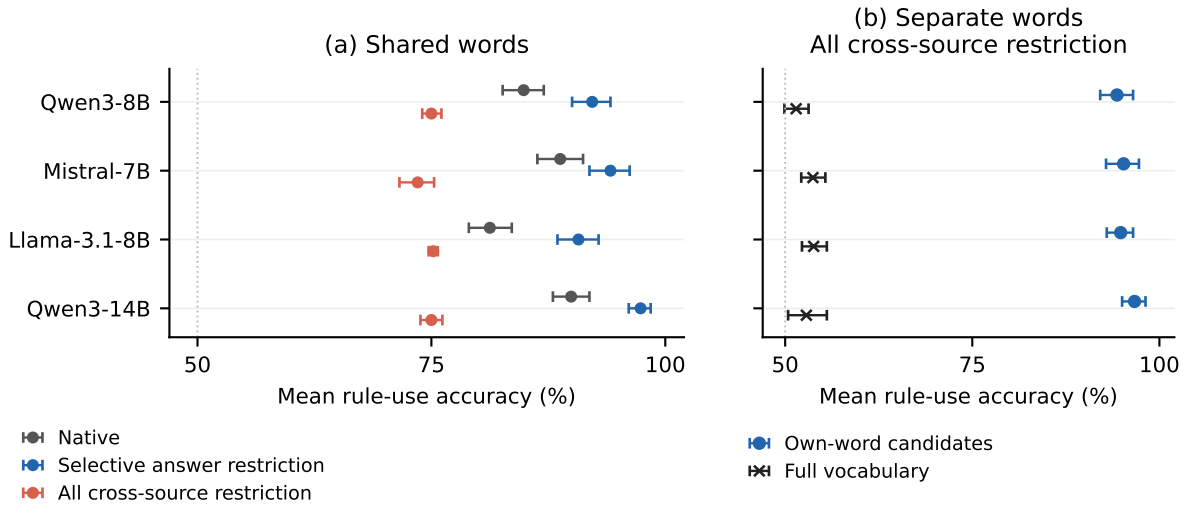


Figure 2: Independently confirmed real-text private functions, n=32 with 95% context-bootstrap intervals. (a) Selective answer restriction improves shared-word accuracy; complete cross-source restriction removes other useful formation paths. (b) With separate words, complete formation restriction preserves own-word candidate accuracy but impairs full-vocabulary output. Both panels allow full-source query access; Appendix F reports all six policies.

signments. Formation and query access are crossed while all private cross-source formation remains blocked.

True-minus-unrelated formation-only gains in full-vocabulary accuracy are 26.66 [23.05, 30.47], 7.71 [5.08, 10.45], and 8.30 [6.05, 10.84] points in Qwen3-8B, Mistral, and Llama. Direct reading-only contrasts include zero in all three. Public relational information therefore affects jointly usable private evidence through formation, even when no query position can read it.

Literal table routing is insufficient: correct tables outperform swapped tables, but swapped tables themselves improve full accuracy and reduce other-source outputs in every checkpoint. Qwen's candidate accuracy also falls, and some source-specific gains remain uncertain (Appendix H). The result identifies useful shared formation structure, with costs; the next test asks whether information-equivalent structures form equally compatible evidence.

6.3 Common formation backgrounds constrain joint rule use

We keep every public relation and private demonstration fixed, reversing only the two rows of the correct schema. Private cross-source formation remains blocked, and no query position can read the table. Each ordering supplies the same facts but forms its own pair of source caches. We combine A and B caches from one ordering (00 or 11), or from different orderings (01 or 10), interchanging all private K/V states at unchanged global positions.

With separate word sets, coherent composition improves full-output use much more than within-source candidate judgments (E126; Appendix I). A post-hoc diagnosis finds that most wrong-source words also match that other source's private convention. This motivates a prospective test with identical output words: a model must select the appropriate learned mapping even though the words themselves cannot identify a source.

E127 fixes four checkpoints and both shared/separate namespaces before scoring.

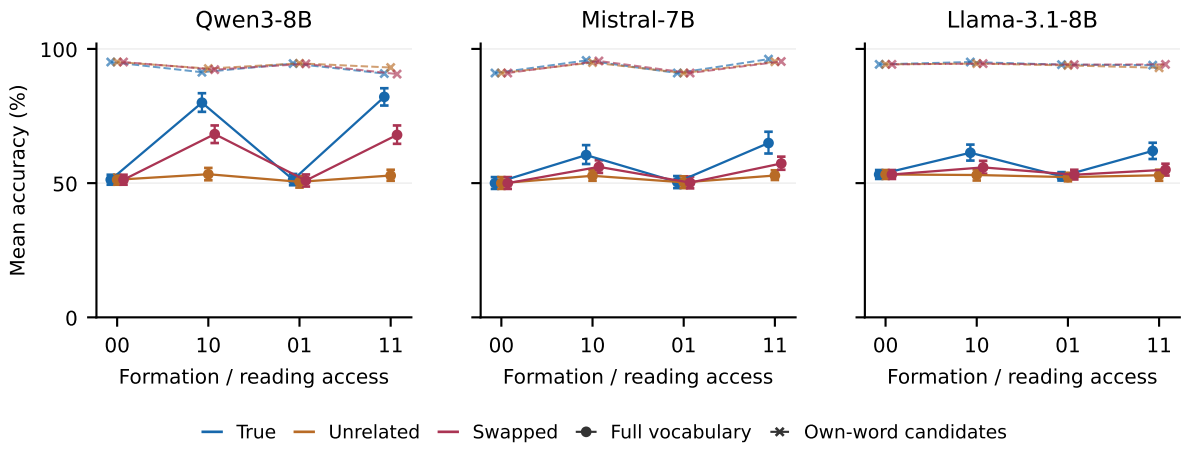


Figure 3: Public source–word-set access with private historical supervision isolated, n=32. Codes give formation/query access to the schema (1=open, 0=blocked). Solid curves show full-vocabulary accuracy with 95% context-bootstrap intervals; dashed curves show own-word candidate accuracy. Correct and swapped relations both help through formation, with different magnitudes.

Its primary uses shared words and opposite private conventions, where the two sources require different answers. The coherent-minus-mixed contrast averages 00 and 11 against 01 and 10. It cancels additive differences in individual-cache quality, while retaining the joint reader's nonlinear behavior.

Coherent composition improves shared-word, opposite-convention accuracy by 8.89 [6.64, 11.13], 5.86 [4.10, 7.71], 0.78 [0.10, 1.56], and 9.38 [7.32, 11.52] points in Qwen3-8B, Mistral, Llama, and Qwen3-14B. The source-conditioned margin also increases in each checkpoint (Appendix J).

The same dependence therefore extends beyond selecting a distinguishable word set. The paired source-conditioned margin is one half of the gold-A-oriented difference between A and B requests. A single request-independent output shift cancels in this difference and cannot explain its positive change. This is an equivalent projection of the two-source margins, not an additional independent mechanism measure.

The caches retain strong requested-source-only functions, yet coherent joint-reading mean accuracy under conflicting conventions is only 53–65%. The small Llama effect and nonuniform private-direction gains remain reported. With food/service criteria instead of opposing codes, coherent formation yields a smaller gain: Qwen3-14B improves by 2.93 [0.59, 5.47] points, while the other accuracy intervals touch or cross zero; joint-reading mean accuracy is about 50–53% (Appendix K).

Local private-function accuracy therefore does not determine joint usability. Shared-word effects exceed a word-set-only explanation; request margins exclude a common output shift. Whole-cache compatibility remains consistent with contextual binding and competitive retrieval (Appendices I–K).

7 The same formation and operation chain in a larger model

We repeat the formation, native-operation, and real-text function tests in Qwen3-32B, with a fresh Qwen3-8B anchor and byte-identical contexts within each stage. All conditions, channels, reading backgrounds, private directions, and instructions are retained. Discovery n16 precedes confirmation n32. Models remain frozen with FP32 eager attention and fixed interventions.

The 32B formation batch has 100% single-source accuracy under both private directions, while the historical relation-formation effect remains. Coefficient-one label-interface KV reconstruction has raw r=.883 and rule-double-contrast r=.922, with raw RMSE .864 versus zero-prediction RMSE 1.857. Native restricted restoration is positive under A-only reading (+.775 [.321, 1.284] nats) and negative under full-source reading (−.544 [−.833, −.276]), as in the fresh 8B anchor. The same formation operation therefore contributes in the larger model, with its function depending on the subsequent reading environment.

On real-text shared-word functions, 32B native full accuracy is 87.99%, rising to 95.70% with

Qwen3, matched confirmation n32	A-formation signed, nats	Native operation removal, nats	Selective-minus-native full, pp	Query-only-minus-native full, pp
8B	2.351 [1.857, 2.837]	1.762 [1.264, 2.294]	7.62 [5.18, 9.96]	−2.54 [−6.54, 1.46]
32B	2.676 [2.266, 3.096]	1.972 [1.467, 2.500]	7.71 [5.76, 9.57]	5.86 [4.20, 7.52]

Table 3: The same formation, native-operation, and real-text function chain across model sizes. Within-family matched confirmation contexts, n32; full-output accuracy and signed native effects.

selective formation; its advantage over complete cross-source formation removal is 20.41 [18.16, 22.36] points. The paired selective-gain difference is .10 [−3.22, 3.32] points, while the query-only gain differs by 8.40 [3.81, 13.09]. Formation and reading contributions thus vary separately across sizes. Training recipes also differ, so these are comparisons of checkpoint computation. Appendix M retains all sources, tradeoffs, outputs, and reconstructions.

8 Related work and positioning

Cho's circuit (Cho et al., 2025), information removal (Cho et al., 2026), Contextualize-then-Aggregate (Bakalova et al., 2025), and How Few-Shot Examples Add Up (Wang et al., 2026) establish input merging, information selection, contextualization, and QK/V contributions. Sun et al. (2025) trace input–label mappings, including 70B/72B checks. These explain important components of our setting. We add a measurement of whose learned relations enter evidence, an output-dependent formation prediction, and its native and functional consequences.

Local task vectors (Zheng et al., 2026) and just-in-time representations (Li et al., 2025) establish heterogeneous task states. Feng and Steinhardt (2024), Dai et al. (2024), and Mixing Mechanisms (Gur-Arieh et al., 2026) explain binding, ordering, and competing retrieval mechanisms; Test, then Route (Thavarasa and Sukumar, 2026) separates predicate testing from token-dependent routing. We examine how these frameworks apply to relations learned from multiple sources sharing output words. Parallel Context Windows (Ratner et al., 2023) and CacheBlend (Yao et al., 2025) already address independently computed prefixes and KV context dependence. Our counterfactual fixes public facts and private positions, blocks private supervision exchange, and tests joint function under equivalent formation backgrounds.

Task recognition/learning (Yang et al., 2026) separates label-space alignment from within-space mappings; Li et al. (2024) test task composition and label remapping. Demonstration Conflict (Jiao et al., 2026) studies corrupted demonstrations, and Jurisdiction (Li et al., 2026b) studies contextual authority over prior knowledge. Here both source rules are legitimate and must be inferred from demonstrations. Yu and Ananiadou (2024), Crosbie and Shutova (2025), and Racing Thoughts (Lepori et al., 2025) motivate operation-specific causal tests. Lepori et al. (2026) and CoSToM (Li et al., 2026a) separate internal information from deployment; our frozen-model tests connect historical relation formation to source-conditioned function.

9 Discussion

Source-conditioned ICL is a problem of evidence provenance as well as evidence access. We identify how shared outputs mediate the historical formation of learned associations, reconstruct naming counterfactuals, and test this operation in native judgments and larger checkpoints.

This account explains why separation policies have different consequences. Selective answer isolation can preserve private rules, while broader isolation can discard useful shared structure. Public source–output relations aid independently learned rules through formation, and equivalent formation backgrounds yield different joint usability. Which source a query reads and how its evidence was formed jointly shape source-conditioned behavior, including when its answer is correct.

Limitations

The checkpoint comparisons couple model capacity with training recipe. Their results describe the computation of the tested models. Dual-aspect comments always place food before service, so the criterion-dependent responses are tied to that serialization. Cache interchanges and renormalized masks test specified interfaces while allowing other paths and subsequent state changes. Policy masks remove unequal edge sets, and reconstructions use actual donor activations.

The source-specific conventions are constructed. They let us balance legitimate conflicting rules and measure their provenance, but do not estimate how frequently such conflicts occur in deployment or their practical error cost. SST5 text may have appeared in pretraining; novelty of the private convention does not imply novelty of each review. Intervals use context resampling and are not adjusted for every comparison. Complete conditions and methodological details accompany the reported effects in the appendix.

References

Aleksandra Bakalova, Yana Veitsman, Xinting Huang, and Michael Hahn. 2025. Contextualize-then-aggregate: Circuits for in-context learning in gemma-2 2b. *Preprint*, arXiv:2504.00132.

Hakaze Cho, Mariko Kato, Yoshihiro Sakai, and Naoya Inoue. 2025. Revisiting in-context learning inference circuit in large language models. *Preprint*, arXiv:2410.04468.

Hakaze Cho, Haolin Yang, Gouki Minegishi, and Naoya Inoue. 2026. Mechanism of task-oriented information removal in in-context learning. *Preprint*, arXiv:2509.21012.

Joy Crosbie and Ekaterina Shutova. 2025. Induction heads as an essential mechanism for pattern matching in in-context learning. In *Findings of the Association for Computational Linguistics: NAACL 2025*, pages 5049–5111, Albuquerque, New Mexico. Association for Computational Linguistics.

Qin Dai, Benjamin Heinzerling, and Kentaro Inui. 2024. Representational analysis of binding in language models. In *Proceedings of the 2024 Conference on Empirical Methods in Natural Language Processing*, pages 17468–17493, Miami, Florida, USA. Association for Computational Linguistics.

Jiahai Feng and Jacob Steinhardt. 2024. How do language models bind entities in context? *Preprint*, arXiv:2310.17191.

Yoav Gur-Arieh, Mor Geva, and Atticus Geiger. 2026. Mixing mechanisms: How language models retrieve bound entities in-context. *Preprint*, arXiv:2510.06182.

Difan Jiao, Di Wang, and Lijie Hu. 2026. Understanding the dynamics of demonstration conflict in in-context learning. *Preprint*, arXiv:2603.04464.

Michael A. Lepori, Tal Linzen, Ann Yuan, and Katja Filippova. 2026. Language models struggle to use representations learned in-context. In *Proceedings of the 64th Annual Meeting of the Association for Computational Linguistics (Volume 1: Long Papers)*, pages 14841–14857, San Diego, California, United States. Association for Computational Linguistics.

Michael A. Lepori, Michael Curtis Mozer, and Asma Ghandeharioun. 2025. Racing thoughts: Explaining contextualization errors in large language models. In *Proceedings of the 2025 Conference of the Nations of the Americas Chapter of the Association for Computational Linguistics: Human Language Technologies (Volume 1: Long Papers)*, pages 3020–3036, Albuquerque, New Mexico. Association for Computational Linguistics.

Jiaoda Li, Yifan Hou, Mrinmaya Sachan, and Ryan Cotterell. 2024. What do language models learn in context? the structured task hypothesis. In *Proceedings of the 62nd Annual Meeting of the Association for Computational Linguistics (Volume 1: Long Papers)*, pages 12365–12379, Bangkok, Thailand. Association for Computational Linguistics.

Mengfan Li, Xuanhua Shi, and Yang Deng. 2026a. CoS-ToM: Causal-oriented steering for intrinsic theory-of-mind alignment in large language models. In *Proceedings of the 64th Annual Meeting of the Association for Computational Linguistics (Volume 1: Long Papers)*, pages 9302–9317, San Diego, California, United States. Association for Computational Linguistics.

Pei-lin Li, Qingle Liu, Junyang Feng, Siyu Li, Sunqi Fan, Xin-Sheng Chen, and Shuojin Yang. 2026b. When context misleads: In-context learning with jurisdiction in large language models. *Preprint*, arXiv:2609.27603.

Yuxuan Li, Declan Campbell, Stephanie C. Y. Chan, and Andrew Kyle Lampinen. 2025. Just-in-time and distributed task representations in language models. *Preprint*, arXiv:2509.04466.

Nir Ratner, Yoav Levine, Yonatan Belinkov, Ori Ram, Inbal Magar, Omri Abend, Ehud Karpas, Amnon Shashua, Kevin Leyton-Brown, and Yoav Shoham. 2023. Parallel context windows for large language models. In *Proceedings of the 61st Annual Meeting of the Association for Computational Linguistics (Volume 1: Long Papers)*, pages 6383–6402, Toronto, Canada. Association for Computational Linguistics.

Pratik S. Sachdeva, Renata Barreto, Geoff Bacon, Alexander Sahn, Claudia von Vacano, and Chris

Kennedy. 2022. The measuring hate speech corpus: Leveraging rasch measurement theory for data perspectivism. In *Proceedings of the 1st Workshop on Perspectivist Approaches to NLP @LREC2022*, pages 83–94, Marseille, France. European Language Resources Association.

Richard Socher, Alex Perelygin, Jean Wu, Jason Chuang, Christopher D. Manning, Andrew Ng, and Christopher Potts. 2013. Recursive deep models for semantic compositionality over a sentiment treebank. In *Proceedings of the 2013 Conference on Empirical Methods in Natural Language Processing*, pages 1631–1642, Seattle, Washington, USA. Association for Computational Linguistics.

Chenghao Sun, Zhen Huang, Yonggang Zhang, Le Lu, Houqiang Li, Xinmei Tian, Xu Shen, and Jieping Ye. 2025. Interpret and improve in-context learning via the lens of input-label mappings. In *Proceedings of the 63rd Annual Meeting of the Association for Computational Linguistics (Volume 1: Long Papers)*, pages 3873–3895, Vienna, Austria. Association for Computational Linguistics.

Luxshan Thavarasa and Sivasuthan Sukumar. 2026. Test, then route: How language models execute in-context conditional rules across models and languages. *Preprint*, arXiv:2608.04183.

Entang Wang, Yiwei Wang, Aleksandra Bakalova, and Michael Hahn. 2026. How few-shot examples add up: A causal decomposition of function vectors in in-context learning. *Preprint*, arXiv:2605.16591.

Haolin Yang, Hakaze Cho, and Naoya Inoue. 2026. Localizing task recognition and task learning in in-context learning via attention head analysis. *Preprint*, arXiv:2509.24164.

Jiayi Yao, Hanchen Li, Yuhan Liu, Siddhant Ray, Yihua Cheng, Qizheng Zhang, Kuntai Du, Shan Lu, and Junchen Jiang. 2025. Cacheblend: Fast large language model serving for rag with cached knowledge fusion. *Preprint*, arXiv:2405.16444.

Zeping Yu and Sophia Ananiadou. 2024. How do large language models learn in-context? query and key matrices of in-context heads are two towers for metric learning. In *Proceedings of the 2024 Conference on Empirical Methods in Natural Language Processing*, pages 3281–3292, Miami, Florida, USA. Association for Computational Linguistics.

Bowen Zheng, Ming Ma, Zhongqiao Lin, and Tianming Yang. 2026. Label words as local task vectors in in-context learning. In *Findings of the Association for Computational Linguistics: ACL 2026*, pages 6319–6332, San Diego, California, United States. Association for Computational Linguistics.

A Experimental details and reproducibility

A.1 Models and numerical execution

Experiments use frozen Qwen3-8B and Qwen3-14B (both post-trained hybrid models), Mistral-7B-v0.3 base, and Llama-3.1-8B base, FP32 eager attention, PyTorch 2.8.0 and Transformers 4.57.6. The controlled sentiment task contains eight demonstrations from each source and eight new queries per context. Source names, private output direction, and demonstration order vary; the two private directions are balanced in confirmation. Shared outputs are `no/yes`; separate B outputs are `pass/flag`. A's outputs remain unchanged. Models process the serialized task as a continuation, without adding a chat template or generating reasoning. Qwen3-14B is a same-family size comparison; training data and recipes are not controlled between sizes. The two-aspect templates always mention food before service, so this experiment measures functional aspect distinctions under fixed order and does not separate semantic identity from position selection.

The checkpoint identifiers and revisions are fixed before scoring. The 70B rows in Appendix M are discovery results until the prescribed confirmation completes; they are kept separate from the confirmed main-text chain.

All reported interventions use frozen weights, evaluation mode, and inference without gradient updates. The cross-size runs use FP32 weights, attention states and matrix products with TF32 disabled, eager attention, and no quantization or CPU offload. The 32B/70B models use contiguous layer placement over two/four GPUs on one node. Every layer receives the same prescribed intervention; no head or layer is selected. The official 70B checkpoint revision `349b2ddb53ce8f2849a6c168a81980ab25258dac` and the public Nous mirror revision `beb678ba5bd7eb1aafeffa01e2a2b3b5e93d1dd3` have identical hashes for all 30 weight shards.

A.2 Materials and serialization

Controlled experiments include single-aspect sentiment templates and dual-aspect restaurant comments. The dual-aspect task crosses food and service polarity and the source-specific criterion. Each context contains eight demonstrations per source and eight novel comments. Source names and demonstration order vary; private output directions are balanced. The model infers the source's rule from examples. Shared outputs are `no/yes`; separate outputs retain A's words and replace B's with `pass/flag`. Label words including their leading space are checked to be single tokens in the tokenizer used for that run.

Natural sentiment text comes from the Stanford Sentiment Treebank (Socher et al., 2013), distributed as `SetFit/sst5`, Arrow build `e51bdcd8cd3a30da231967c1a249ba59361279a3`. We keep only original classes 0 and 4, mapped to negative and positive, and deduplicate identical text within each split. Discovery and confirmation demonstrations are drawn from opposite SHA256 text-hash parities of training. Discovery queries use official validation; confirmation queries use official test, excluding text in training and, for confirmation, validation. The command-line mode named `validation` denotes confirmation and selects the SST5 test split. Comments contain 5–64 tokens under the tested tokenizer. Selection uses gold sentiment and token geometry, with no model-output screening.

The earlier Measuring Hate Speech (Sachdeva et al., 2022) task uses the public `ucberkeley-dlab/measuring-hate-speech` corpus. We group by comment ID, average hate-speech and target-category annotations, keep comments of at most 50 words with mean hate-speech score at least 1, and separate race-targeted texts (race score at least .75, gender at most .1) from gender-targeted texts (gender at least .75, race at most .1). A fixed shuffle seed 5600 gives a 70/30 demonstration/query partition. Sources adopt opposite category-to-label conventions, with both private directions evaluated. These constructed rules test source-specific supervision rather than reconstructing individual annotator preferences.

For B's eight demonstrations we sample two length/offset-matched quartets, each containing two positive and two negative comments. Cross-class reassignment exchanges the positive and negative comments while retaining the label stream; within-class reassignment swaps within each polarity and moves the same number of comments. This preserves the input multiset, source names, label words, anchor positions, and all tokens outside comment content. A's eight comments and all eight queries stay fixed. Four query comments have each sentiment; the private reversal alternates across contexts. Paired-model materials are byte-identical within family and stage, while stage generators and

Checkpoint	Training stage	Hugging Face revision
Qwen/Qwen3-8B	Post-trained hybrid	`b968826d9c46dd6066d109ea bc6255188de91218`
Qwen/Qwen3-14B	Post-trained hybrid	`40c069824f4251a91eefaf28 1ebe4c544efd3e18`
Qwen/Qwen3-32B	Post-trained hybrid	`9216db5781bf21249d130ec9 da846c4624c16137`
mistralai/Mistral-7B-v0.3	Base	`caa1feb0e54d415e2df31207 e5f4e273e33509b1`
NousResearch/Meta-Llama-3.1-8B	Base	`1f47e50cdbe801ad8a517415 6ec3a0655108fb9f`
NousResearch/Meta-Llama-3.1-70B	Base	`beb678ba5bd7eb1aafeffa01 e2a2b3b5e93d1dd3`

Table 4: Fixed model checkpoints, training stages, and exact acquisition revisions. Llama mirrors are verified against the official weight hashes.

query samples remain separate.

The input is a task continuation, with no added chat template or generated reasoning. The common header is `Below are comments and the labels that individual annotators gave them.` followed by two newlines. Each demonstration has the following literal field order, followed by two newlines:

```
Comment: <original comment>
Annotator: <source name>
Label: <one label word>
```

The query uses the same fields and ends immediately after `Label:`. The block ending at that colon, the leading-space label word, and the two-newline separator are encoded separately without special tokens. This segmentation is part of the protocol. The first answer-token position is the label anchor. The historical answer region includes that word and its separator tokens. Source ownership extends across every token of its demonstration; common header tokens have no source owner.

The actual serialized token stream is part of the method. Blocks, answer words, and separators are encoded separately. For Mistral, this introduces fixed extra spaces relative to literal whole-string encoding: the first native confirmation context has 351 rather than 335 tokens. Its B answer-message region includes the answer word plus three separator tokens per demonstration, totaling 32 sender positions; Qwen uses 16. Paired token positions and boundaries remain identical across the relevant manipulations. Native controls validate the actual stream, not a claimed equivalent whole-string encoding. Source code and independent tokenizer reconstruction preserve this disclosure.

A.3 Causal interfaces and access graphs

Formation and reading are manipulated separately. In a formation-message interchange, donor K, V, or KV at historical B-answer positions replace the corresponding inputs to attention only for A-label receivers at each layer. The implementation calls the model's native eager attention with the full original matrix geometry, retaining recipient queries and the causal mask; only selected receiver outputs are substituted. Other receiver outputs remain native, and subsequent states evolve normally. Donor keys are taken from the actual layer cache, with its positional transformation already applied. Cached-state interchanges instead replace selected A/B K/V positions in a deep copy of a completed prefix; they leave global positions and the other recipient cache positions unchanged.

Native-operation tests have no donor. They set the specified historical attention edges to negative infinity before softmax. Targeted removal blocks earlier B-answer regions at A-label receivers; within-source removal blocks A-answer regions there. Complete B-to-A removal excludes all earlier B-owned positions from every A-owned receiver, and restricted restoration reopens only earlier B-answer-to-A-label edges in that graph. Symmetric functional policies apply the source-swapped rule to both groups in one prefix. Masks renormalize attention, remove unequal edge sets, and permit downstream state changes.

Reading restrictions apply to every live query token, including source-name and comment tokens, across every layer. They exclude the unrequested source's complete prefix positions while retaining the common background, requested-source positions and causal query history. A-only reading therefore restricts direct access; it retains any ear-

lier B influence already formed in the allowed evidence. The eight queries are batched with left padding and explicit positions starting after the prefix. Each scoring call uses a fresh deep copy of the pristine cache; padded tokens are masked and cache sequence length and storage independence are checked.

Public-schema experiments block private cross-source formation, independently crossing access to the public source–word-set table during prefix formation and whole-query reading. The table has no comment–label pairs or private polarity. Equivalent-background experiments reverse its two rows, then cross A/B private caches from the two backgrounds at identical positions. Coherent compositions use 00/11 and mixed compositions use 01/10. The query cannot read the public table; these tests concern the joint usability of independently formed private evidence.

A.4 Scores, aggregation and uncertainty

Let z_{cq} be the second-minus-first candidate-label logit difference and $g_{cq} \in \{-1, +1\}$ the requested source's gold direction. For a B-rule intervention, the context response is $D_c = Q^{-1} \sum_q g_{cq}(z_{cq}^{cross} - z_{cq}^{origin})$. Signed effects compare these context responses; absolute effects first take $|D_c|$, then form the paired contrast. Likewise, the fixed-word relational excess compares the magnitudes of context-averaged cross-class and within-class effects, rather than averaging absolute per-query differences.

Candidate accuracy restricts the decision to the requested source's two words. Full-output accuracy requires the highest-logit token in the entire vocabulary to equal the correct word. Joint correctness requires both source requests to be correct for the same input/world. Mean accuracy over sources is a different measure. Real-text function-policy primaries (E123/E131) give equal weight to origin/cross worlds, both sources and all queries within each context; the within-class world remains a separate control. Controlled criterion–convention policy tests (E122) average all four B criterion/convention worlds. Single-source and instruction controls are evaluated separately. Gold-agreement strata are determined by the two requested-source answers, with no filtering by observed correctness.

Formation reconstruction compares the actual namespace-donor interface effect with the historical message-interchange effect using coefficient one. Raw predictions retain both rule worlds; rule-double contrasts subtract the origin effect from the cross effect. Pearson correlation and RMSE summarize those predefined targets, with the zero-prediction RMSE as reference. No slope, offset, channel weight or intervention strength is fitted.

We use 10,000 percentile bootstrap resamples of whole contexts. Each draw preserves all queries, worlds, source requests and paired model conditions belonging to that context. Context-level contrast seeds are 910; reconstruction uses 109910. Correlations are computed from flattened targets with whole-context resampling, not an average of context-wise correlations. Main comparisons use fresh confirmation n32 after discovery n16. The original native/formation/function discovery seeds are 120001/121001/123001 and confirmation seeds are 120101/121101/123101. E131 retains the discovery seeds and uses new confirmation seeds 131101/131102/131103. These are finite checkpoint comparisons, with training recipes allowed to differ across sizes.

A.5 Implementation checks and complete records

Effects use 10,000 context-bootstrap resamples, seed 910; reconstruction intervals follow the frozen E109 analysis. Queries within a context are not treated as independent observations. Raw logits, full-vocabulary argmax, binary candidate-label accuracy, all conditions, and failed controls are retained. Confirmation uses new word/template pools; discovery-to-confirmation text overlap is zero for E113 and the Qwen E114 batches, but vocabulary/templates recur across contexts within a pool.

Instruction controls differ between experiment engines. E122/E123 add one sentence to native formation. In E124–E127, the stored `instruction` condition adds that sentence while also blocking private cross-source formation; it is a composite control, and its difference from native cannot be attributed solely to prompting. E124/E125 additionally include an instruction-free independent-formation reference. The primary formation, schema, and cache-composition contrasts do not use this instruction condition.

Self interchanges produce zero error. Cached versus full-token native query scoring differs by at most 4.96e−5 in Qwen and 7.63e−6 in Mistral confirmation. Cache fingerprints, layer consumption, and matched token geometry are checked. An inde-

pendent session reconstructs the token layouts and recomputes 559 numerical checks from raw records before reading author summaries, with maximum discrepancy 4.77e−15. Separate bootstrap resampling confirms the signed interactions; it is not represented as an independent GPU rerun.

The earlier E113/E114 package comprises five runs and 36,736 scores. The completed E115–E119 expansion comprises 24 discovery/confirmation runs, 576 contexts, 181,632 scores, and 2.393 recorded GPU-hours. Independent raw reviews recompute 36,750 numerical quantities (maximum discrepancy 1.13e−14), reconstruct actual token layouts, and check recorded all-layer consumption. This is independent analysis, not an independent GPU rerun. The E120/E121 native-operation and real-text additions comprise 16 runs, 384 contexts, 164,352 scores, and 2.540 recorded GPU-hours. Their independent raw audits recompute 18,064 scalars (maximum discrepancy 5.00e−15), reconstruct 2,112 actual token streams, and verify the recorded masks and reachable edges. No weights are trained. All model/seed/channel results remain available. Cards specify revisions, script hashes, and rebuilding commands; small summaries, figures, code, and audits are versioned, while raw JSONL and models remain at the documented asset paths.

B Complete criterion–convention functional response

B.1 Separating criterion from private output convention

A learned association can change a classifier's confidence without specifying a complete criterion for a new input. We test this distinction with comments containing independently balanced food and service sentiment. Source A's demonstrations establish a food criterion. Source B can evaluate either food or service, under an output convention agreeing with or opposite to A's. The criteria are inferred from demonstrations rather than stated in the multi-source prompt (E117/E118).

The four B worlds are constructed by bijectively reassigning inputs to fixed demonstration positions. A's raw demonstrations, B's output stream, source names, positions, and input multiset remain unchanged. Eight novel queries cover all four sentiment combinations. Every interface is evaluated on the same queries: complete native histories, histories whose queries cannot read B positions, A-label cache interchange into one recipient, and historical B-answer messages consumed during A-label formation. Separate-output label donors test the effect of output coding during formation while restoring the shared-output recipient for querying.

Let $f_q, s_q \in \{-1, +1\}$ be food and service polarity, oriented to A's private convention. For criterion t, define $\delta_t(q) = \frac{1}{2}[z_{t,+}(q) - z_{t,-}(q)]$, the response to changing B's convention. The context-level matrix

$$M_c = \begin{pmatrix} \mathbb{E}_q[\delta_{food}(q) f_q] & \mathbb{E}_q[\delta_{food}(q) s_q] \\ \mathbb{E}_q[\delta_{service}(q) f_q] & \mathbb{E}_q[\delta_{service}(q) s_q] \end{pmatrix}$$

separates the two criteria's effects on the two aspects. Bias and food–service interaction are also retained. These are B-relation effects while A is requested, not B's own accuracy or a gain in correctly isolating A.

B.2 Association magnitude and criterion selectivity give different answers

Across three model families, E117 confirmation finds that changing B's pairing alters both food and service sensitivity through A-label caches. A positive service response alone therefore cannot establish transport of B's service criterion. E118 independently crosses criterion with B's convention and preserves the entire response matrix. Full native histories show more criterion selectivity than local label interchanges in several settings, with substantial family variation. The mean criterion–feature interaction is positive in all three families, but Qwen3-8B's interaction mainly reflects service selectivity; unequal gains along one common direction can also produce a positive mean interaction.

Separating B's output namespace during formation can lower both aspect responses without appreciably changing their mean selectivity contrast. In Llama, food-criterion responses decrease by .267/.243 nats on food/service, and service-criterion responses by .110/.079; all four paired intervals are negative. Namespace effects on evidence formation thus include a change in relational influence magnitude, not only which aspect dominates.

B.3 A prospective functional-direction test in Qwen3-14B

Before any E119 scores, we freeze an additional observer: the determinant of each context's response matrix and its normalized determinant,

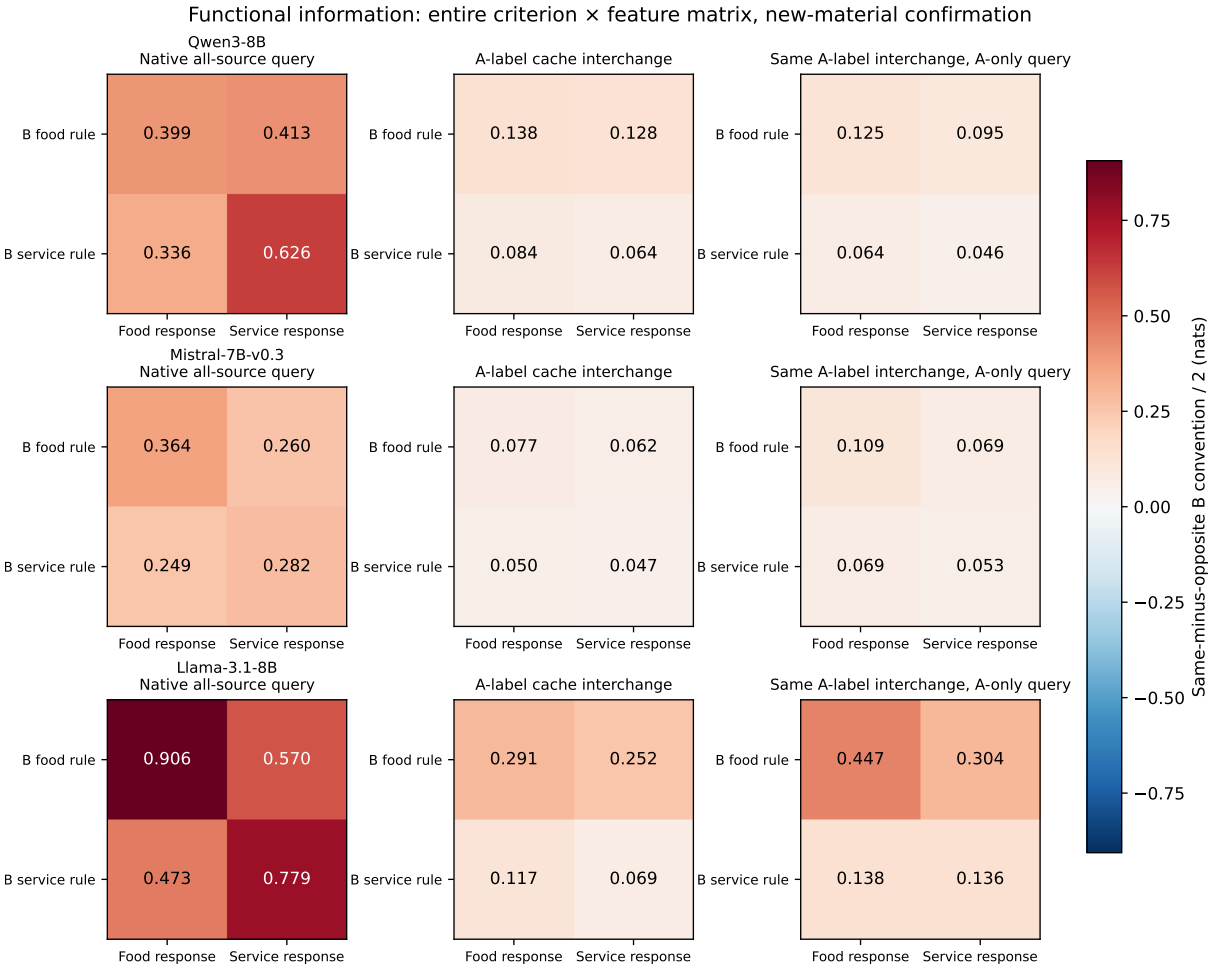


Figure 4: New-material confirmation, 32 contexts per family. Rows vary B's inferred criterion; columns measure the induced food/service response under independent B-convention reversal. The requested source remains A. Full matrices avoid equating one positive aspect response with complete criterion transfer.

$\det(M_c)/(\|M_{c,food}\|\|M_{c,service}\|)$. If the two criteria merely rescale a common functional direction, each determinant is zero. These quantities are computed within context before averaging. A row norm of at most .001 nats sets the normalized statistic to zero while retaining the context; no confirmation context meets this threshold.

Qwen3-14B confirms criterion-specific native responses: food selectivity is .282 [.199, .371] nats and service selectivity .404 [.294, .528]. Native signed determinant is .568 [.408, .751] nats² and normalized determinant .410 [.330, .493]. When every query position is prohibited from reading B, the A cache still formed from the full history yields .0289 [.0032, .0553] and .190 [.073, .310]. Thus criterion-related history can have distinct functional effects through the remaining computation.

The local label interface does not reproduce the full native responses. Its negative signed determinant also rules out describing every formation interface as a common-direction gain; the gated label interface has a small positive normalized response. The signed and normalized statistics weight contexts differently, so both are reported. Nonzero determinants reject the restricted functional common-direction account, not every nonlinear latent representation or an alternative retrieval policy.

Historical answer KV only partially predicts the actual criterion-label counterfactuals: native/gated raw correlations are .668/.665, RMSE .646/.878, compared with zero-prediction RMSE .746/1.018. This is a functional boundary of the formation reconstruction, not a reason to discard its successful namespace prediction. Local association transfer, whole-history criterion dependence, and source-conditioned use are different causal questions.

Qwen3-14B interface, n=32	Mean interaction J, nats [95% CI]	Signed determinant, nats² [95% CI]	Normalized determinant [95% CI]
Complete native histories	.343 [.265, .433]	.568 [.408, .751]	.410 [.330, .493]
Full history, A-only query	.058 [.029, .089]	.0289 [.0032, .0553]	.190 [.073, .310]
A-label interchange	−.0238 [−.0478, −.0025]	−.0201 [−.0374, −.0052]	.006 [−.121, .140]
A-label interchange, A-only query	.038 [.012, .068]	.0074 [−.0106, .0271]	.124 [.017, .242]
Historical answer KV at formation	.0060 [.0029, .0093]	.00068 [.00012, .00130]	.085 [−.042, .214]
Historical answer KV, A-only query	.0126 [.0076, .0188]	.00194 [.00031, .00474]	.072 [−.040, .190]
Separate-output label donor	−.0192 [−.0348, −.0060]	.00192 [−.00123, .00556]	.011 [−.145, .164]

Table 5: Criterion-dependent functional directions in Qwen3-14B confirmation, n32. Context-wise signed and normalized determinants distinguish direction changes from common rescaling; all prescribed interfaces are retained.

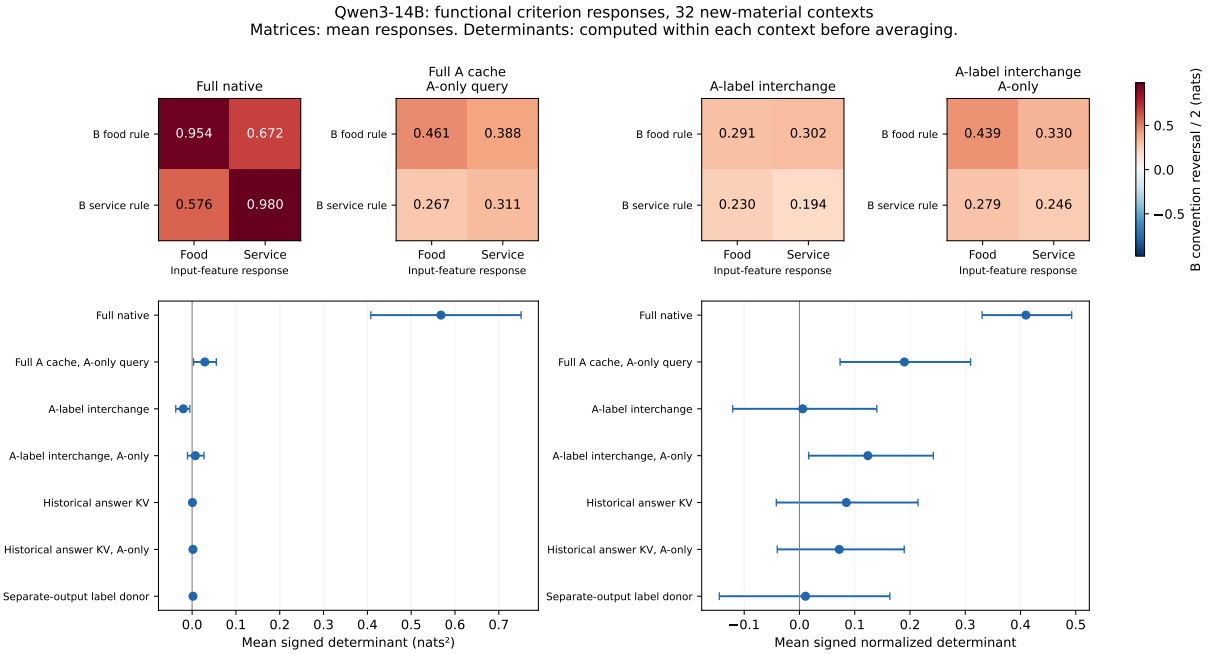


Figure 5: Qwen3-14B confirmation. Matrix means summarize functional responses; determinant intervals use determinants computed separately within each context. All seven predefined interfaces, including reverse and weak effects, are retained. Different cache interfaces do not establish different hidden algorithms.

C All formation/query policies and functional controls

Policy results average the two source requests, four B worlds, and eight queries within each context. Candidate accuracy and full-vocabulary argmax remain separate in raw records. `normal`, `answer_cross_off`, and `all_cross_off` use no trained parameters; the generated source ownership determines the mask. Both actors consume the same prefix cache.

All entries are accuracy percentages with context-bootstrap95%CI. Separate A/B, agreeing/disagreeing correct labels, all four criterion/convention worlds, single functions, one-sentence instructions, Food/Service/bias/FS responses, and all signed contrasts are in `results/e122/*_confirmation/analysis.json` and the independent raw audit. The single-aspect SST ability does not establish perfect dual-aspect primitive mastery.

Bars show the source-average primary; black circle/triangle show A/B means. These are not joint both-correct rates. Each paired contrast uses its context-level difference, not differences between separately constructed intervals.

Independent confirmation: each source's accuracy change, agreement/conflict of the correct labels, and alternative policies. All combined formation/query conditions and full-vocabulary scores are reported in the appendix.

Model, confirmation n=32	Native	Answer formation	All cross formation	Query only	Answer + query	All formation + query
Qwen3-8B	71.58 [69.43, 73.44]	74.51 [72.66, 76.17]	62.60 [61.62, 63.67]	70.46 [67.87, 73.05]	73.78 [71.58, 75.73]	73.29 [70.80, 75.68]
Mistral-7B-v0.3	69.14 [67.63, 70.65]	72.56 [70.31, 74.85]	62.74 [60.99, 64.55]	70.85 [69.14, 72.46]	71.97 [69.87, 74.12]	73.54 [71.04, 75.88]
Llama-3.1-8B	73.83 [72.22, 75.44]	81.20 [79.10, 83.25]	68.55 [66.75, 70.26]	85.89 [84.08, 87.74]	86.77 [85.11, 88.53]	84.18 [81.30, 86.87]
Qwen3-14B	78.12 [75.68, 80.76]	80.91 [78.86, 82.86]	67.58 [66.02, 69.19]	84.23 [81.64, 86.82]	83.30 [80.81, 85.69]	82.91 [80.22, 85.64]

Table 6: Controlled dual-aspect rule use under all formation and reading policies, n32 confirmation. Entries are mean candidate accuracy in percent over both sources and the predefined worlds and queries.

Model, n=32	Single A	Single B service/same	Single B service/opposite	Single B food/same	Single B food/opposite
Qwen3-8B	67.97 [64.06, 71.88]	83.59 [79.69, 87.50]	80.08 [76.17, 83.98]	70.70 [67.58, 73.44]	72.66 [69.14, 75.78]
Mistral-7B-v0.3	72.66 [69.53, 75.39]	72.27 [68.36, 75.78]	75.78 [72.27, 79.30]	76.95 [73.83, 80.08]	76.56 [73.44, 79.69]
Llama-3.1-8B	79.30 [75.00, 83.59]	87.11 [83.59, 90.62]	86.33 [82.03, 90.62]	82.42 [78.12, 86.72]	81.64 [76.95, 85.94]
Qwen3-14B	82.42 [78.91, 85.94]	85.55 [81.64, 89.45]	88.67 [84.77, 92.58]	83.59 [80.08, 87.11]	84.77 [80.86, 88.67]

Table 7: Single-source functional controls on dual-aspect comments, n32 confirmation. Percent candidate accuracy is reported separately for each B criterion and convention.

D Native operation and real-text completeness

Paired signed B-rule-response differences in nats with95%CI. The operation acts at every prefix layer, masks strict historical senders, preserves self attention, and renormalizes the remaining attention. The restricted restoration opens only the historical answer→label edges; it does not restore original numerical messages or all native paths. E120 single within-answer-off controls, far-namespace effects, residuals, and signed/magnitude contrasts are in the card/results.

Each SST private direction has n16contexts; natural-text pool/split/source provenance is independently audited. The word-plus-separator sender regions differ by tokenizer. Llama fixed-word association signed effects are negative; all K/KV/V raw/rule reconstructions and zero predictions remain in the summaries.

Paired signed B-rule effects after operation removal and restricted restoration. The contrast in restoration direction shows why historical sharing cannot be described as uniformly harmful.

E Registered evidence map

F Real-text function policies and complete output readouts

Primary accuracy averages origin/cross worlds, two requested sources, and eight queries within each context. Within-class movement is reported separately. The original primary is candidate accuracy; full-vocabulary output and word-space diagnostics are separate secondary readouts. Shared-word candidate/full-vocabulary aggregates coincide in all six confirmation policies. Private polarity is balanced across contexts; source-specific conventions are inferred from examples, not provided as explicit rules.

These are paired context-bootstrap95%CI. The namespace contrast is not headroom matched; word meanings/priors also differ. No layer, head, intervention strength or context is selected to improve these results.

This classification uses the actual argmax token and the other source's two word IDs. The average other-source-word rates are 44.6–46.9%, but their A/B composition differs. In the shared-word case, surface words cannot identify their source: recording zero identifiable other-space words is not evidence of provenance isolation.

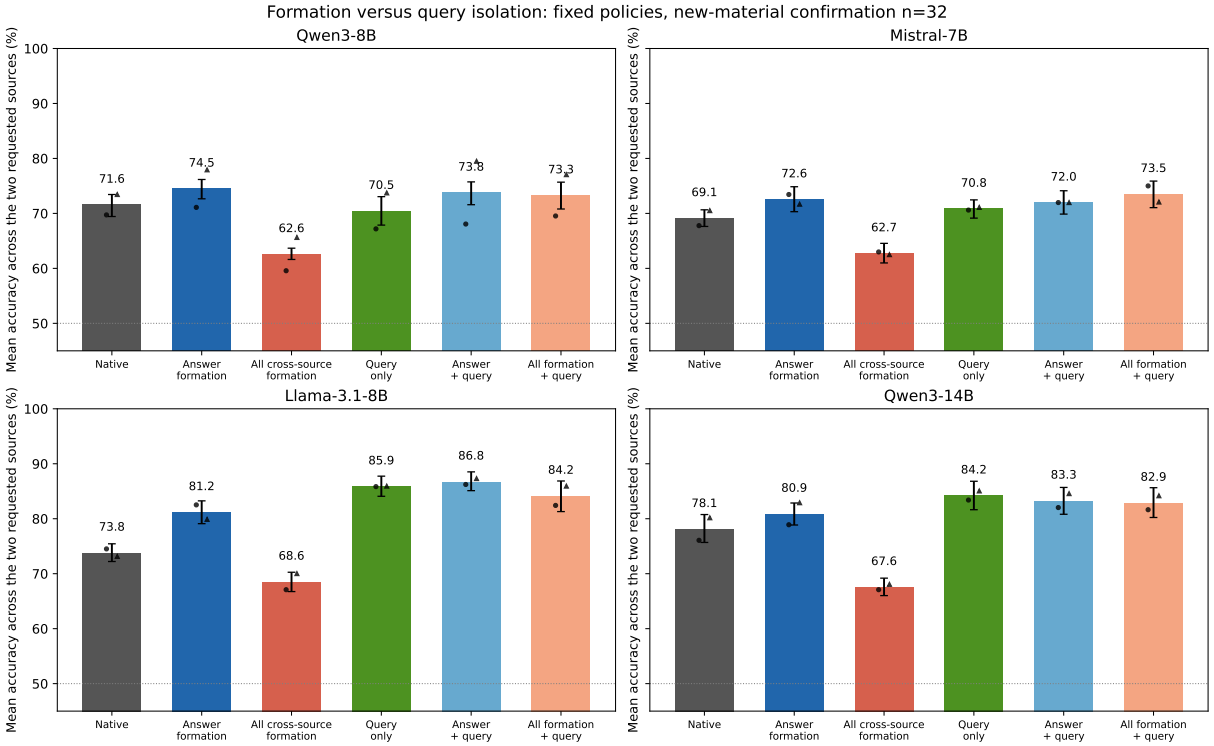


Figure 6: Complete policy accuracy

Model	Selective-minus-native A, pp	B, pp	Agreeing golds, pp	Disagreeing golds, pp	Selective-minus-all formation, pp
Qwen3-8B	1.37 [-0.68, 3.32]	4.49 [2.15, 6.84]	-2.34 [-4.00, -0.78]	8.20 [5.86, 10.45]	11.91 [10.25, 13.53]
Mistral-7B-v0.3	5.66 [2.54, 9.08]	1.17 [-1.27, 3.52]	-1.95 [-4.98, 0.88]	8.79 [6.15, 11.62]	9.81 [6.74, 13.04]
Llama-3.1-8B	8.01 [4.79, 11.33]	6.74 [4.49, 8.89]	-1.86 [-4.30, 0.39]	16.60 [13.87, 19.43]	12.65 [10.40, 15.04]
Qwen3-14B	2.83 [0.00, 5.47]	2.73 [0.10, 5.37]	-3.61 [-5.86, -1.56]	9.18 [6.15, 12.30]	13.33 [11.33, 15.38]

Table 8: Paired selective-formation effects by requested source and gold-label agreement, n32 confirmation. Changes are in percentage points with context-bootstrap 95% CIs.

The private positive/inverted halves, actual full-vocabulary controls, all one-sentence instructions, complete A convention-response/signed/within movement readouts, and all 85 condition statistics remain in the experiment summaries and independent audit. Aggregate private functions are usable but not uniformly perfect; the distinct words also have asymmetric direction priors. Formation masks remove unequal sets of edges, renormalize attention, and allow states to evolve. Consequently the result identifies functional consequences of fixed policies, not isolated representation components or a tuned performance method.

All six formation/query policies in both output namespaces, independently confirmed n32 for every checkpoint. Circles score within the requested source's two candidate words; black crosses score the full-vocabulary output. Error bars use the original context-bootstrap intervals. This complete view accompanies the focused Figure 2 in the main text.

The observer is the absolute context mean of the A-gold-oriented half contrast between B conventions, computed in A's fixed output space. It is not attention-based source ranking. Signed and matched within-class movement remain alongside it.

E123 freezes the scientific source at91374879 before the first score. Four checkpoints each execute discovery16 and confirmation32, totaling192contexts/130560scores. Scientific inputs, masks, full-output tokens, and small report data are independently checked; see E123 and its raw

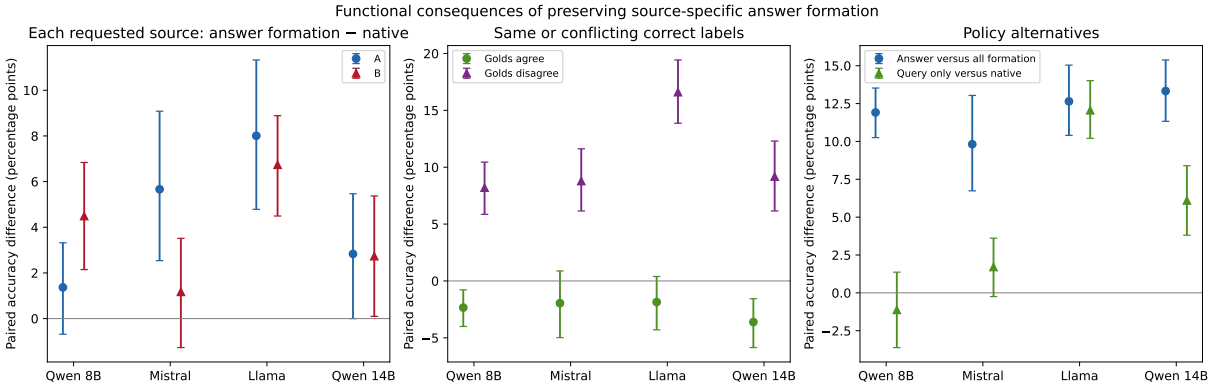


Figure 7: Functional source-policy effects and tradeoffs

Model, n=32	Full-source operation removal	A-only restricted restoration	Full-source restricted restoration
Qwen3-8B	1.649 [1.108, 2.250]	1.455 [0.999, 1.945]	-0.806 [-1.365, -0.266]
Mistral-7B-v0.3	1.054 [0.886, 1.227]	0.647 [0.463, 0.835]	0.140 [0.053, 0.231]
Llama-3.1-8B	0.808 [0.517, 1.109]	0.958 [0.635, 1.276]	-0.233 [-0.336, -0.122]
Qwen3-14B	1.329 [0.967, 1.717]	1.120 [0.673, 1.646]	0.039 [-0.217, 0.287]

Table 9: Native answer-path removal and restricted restoration, n32 confirmation. Signed B-rule contrasts are measured in nats, separately for full-source and A-only reading.

audit. The pinned SST downloader reproduces all three Arrow files and the Llama32 confirmation contexts byte-for-byte in the recorded environment, without modifying the frozen scientific modules.

G Shared identities and independent private supervision

E124 isolates both sources' private demonstration histories at every layer, retains original global positions, and places a public names-only list after the generic header but before all supervision. F controls demonstration access to the list; R controls every query position's direct access. All generic public tokens precede the list, preventing an indirect public relay when F is off. No output words, criterion or private convention are provided. True and unrelated lists have equal token length; numerical cache and full-output no-ops establish their complete invisibility under F0R0.

The primary is full-vocabulary accuracy, averaged uniformly over worlds, sources, and eight queries within each context. Identity-specific formation gain is [true F1R0−true F0R0]−[unrelated F1R0−unrelated F0R0]; reading gain substitutes F0R1. It is not a comparison to the shorter no-list prefix.

Qwen's formation identity gain is positive for A and B and both private directions. Mistral's true formation-only operation has opposite effects on A and B; Llama's positive candidate identity contrast does not transfer to the full-output primary. Mistral's reading identity contrast is positive but small. All 47 conditions, instructions, single-source functions, private directions, signed/absolute B-convention responses, and F×R contrasts are retained in the complete report and its summaries. There is no equivalence conclusion from a confidence interval crossing zero.

Three checkpoints each execute 16 discovery and 32 confirmation contexts, changing Alex/Sam to Robin/Jordan in confirmation. The 54,144 scores undergo independent numerical, token and graph review; 11,742 numerical comparisons match exactly. Whole-query mask consumption is checked in every layer. Independent review is not an independent GPU rerun. The names-only public context is not a full entity–attribute binding construction, and its partial/negative effects do not refute all binding mechanisms.

H Public output schemas and private rule composition

E125 keeps private cross-source demonstration access blocked at every layer. The public table precedes all supervision and specifies only source-to-word sets, not input–label pairs, criterion, or private polarity. F controls all demonstration access; R controls direct access by every query token.

Model, n=32	Single-source positive / inverted accuracy, %	Native A-formation magnitude, nats [95% CI]	KV rule r / RMSE / zero RMSE
Qwen3-8B	98.4 / 88.3	2.151 [1.487, 2.895]	0.931 / 1.363 / 3.658
Mistral-7B-v0.3	99.2 / 96.9	1.091 [0.925, 1.258]	0.948 / 0.336 / 1.048
Llama-3.1-8B	97.7 / 99.2	0.689 [0.488, 0.901]	0.930 / 0.309 / 0.840
Qwen3-14B	99.2 / 93.8	1.678 [1.340, 2.052]	0.966 / 0.519 / 1.949

Table 10: Natural SST5 private functions and historical formation, n32 confirmation. Single-source accuracy, formation magnitude, and fixed-coefficient KV rule reconstruction retain distinct units.

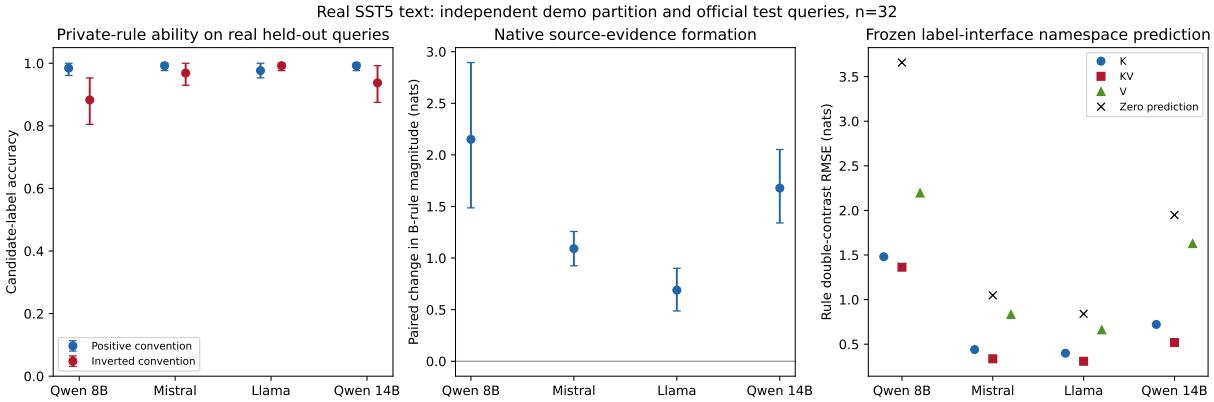


Figure 8: Real-text formation and private functions

Generic HEAD tokens precede the table, preventing a later public relay. True, unrelated-name, and swapped-relation tables have matched lengths and unchanged private token positions. Query reading otherwise sees both source caches. This is explicit relational supervision, not fully unsupervised coordination.

The primary difference-in-differences compares true and unrelated formation-only or reading-only gains. All three tables share identical relevant states and full-output argmax under F0R0. Below, confidence intervals are 10,000 context-bootstrap intervals; contexts, rather than individual queries, are the sampling units.

A correct schema helps more than its swapped counterpart, but the swapped schema also improves actual full-output accuracy and reduces other-source outputs in every model. Therefore a positive true−swapped contrast does not establish active wrong-space routing. Qwen's true formation operation costs 3.81 [1.86, 5.66] candidate-accuracy points; Mistral and Llama have uncertain A-specific identity gains. All private directions, A/B/world readouts, single functions, instructions, signed/absolute B-convention effects, and F×R interactions are retained in the complete report and per-run analysis files. A shared schema is not a lossless or universally sufficient composition method.

The frozen scientific source is `75f75665 82cbaf7cf80dd130b409b32436df281a262884d8 db381991437612b9` (preregistered commit a0a79918). Six runs include 144 contexts and 72,576 scores, recording 0.735239 GPU hours. The independent audit reconstructs 18,126 raw numerical entries with maximum error zero, actual tokenizer geometry, golds, 100,800 prefix and 302,400 query layer-mask calls, and 9,216 inert-table full-output comparisons. This is independent analysis rather than an independent GPU rerun. Fresh confirmation contexts and names do not constitute an entirely new benchmark.

I Equivalent-context cache composition

E126 changes only the order of two correct public source–word-set rows. Private examples, all public facts, source names, token lengths, and private global positions are unchanged. Private cross-source formation remains blocked at every layer, while every query position is denied direct schema access. Coherent pairs 00/11 combine whole source-owned K/V from one ordering; mixed pairs 01/10 combine one source from each ordering. Generic HEAD and unreadable schema cache states come from order0. Coherent11 is verified against the pristine order1 computation under the same query mask.

The primary is full-vocabulary accuracy, with

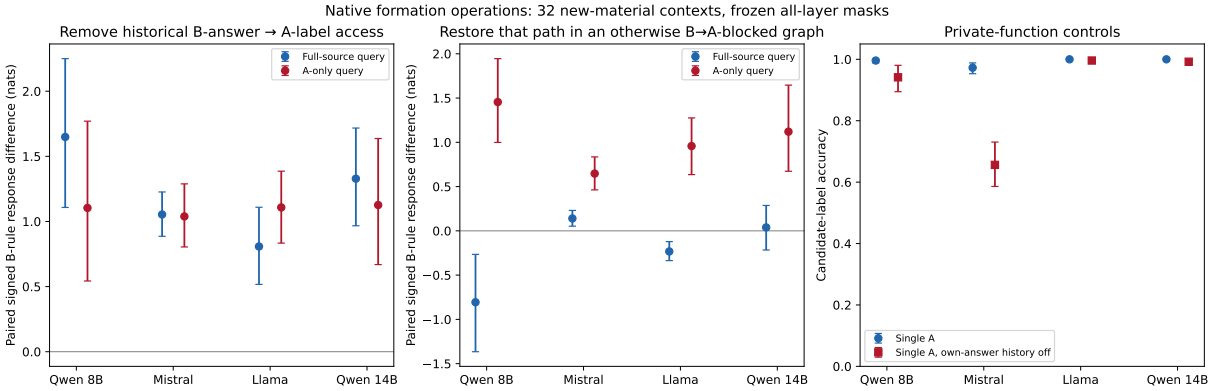


Figure 9: Native operation and restricted restoration

Manuscript claim	Records
Positional access does not isolate relation dependence	E90, E91
Output-dependent formation reconstruction, C21	E109, E111, E112, first independent audit
Fixed-label partial association content, C22	E111, E113, E114
Native operation contribution and restricted restoration	E120, independent operation audit
Reliable private functions and formation on natural text	E121, independent data and mechanism audit
Native conditional effects and interaction, C23	E114, new independent audit
Functional criterion response and its formation-interface boundary, C24	E117, E118, E119, E118 audit, E119 audit
Llama, larger checkpoint, and real-text coverage	E115, E116, E119; wave audit
Limits and failed functional interfaces	E50–E52, E99, E107, E110; full ledger in CLAIMS
Joint source-function formation and query policies	E122, independent raw audit

Table 11: Manuscript evidence map. The experiment records specify the full intervention conditions and raw-derived analyses supporting each claim.

equal weights over worlds, sources, and queries within a context. The paired coherent−mixed contrast is .5[(00+11)−(01+10)]. It cancels additive single-cache quality effects but can still reflect nonlinear output scoring, attention competition, or binding. It does not by itself identify nonlinear latent representations. Requested-source-only reading makes the balanced contrast zero by the combination design; individual function scores must be reported separately.

Both public orders have strong source-only function scores, but they are not functionally identical. Joint order11−00 is negative in Qwen and Mistral. Their primary correct-margin point estimates are also negative with intervals spanning zero; improved full-output use is not a uniform positive shift in every score. Mistral's A-specific joint full accuracy is only 42.58% in00 and 33.40% in11, despite strong requested-source-only scores. All actor, world, private direction, margin, convention-response, single-function, and one-instruction results are retained in the complete report.

Six runs comprise144contexts/49536scores and0.627887GPU hours. Source `ad2c6f7a f7ab81966ea1253a80e7006747f3485b182fddd1 453ffbb1ceb45975` was frozen at377de6c4 before scoring. The independent audit reconstructs12462 numerical entries with maximum error0,1584 actual token/mask layouts,49536 golds,52800 prefix and206400 query layer-mask calls,9216 mixed-own argmax controls, and4608 A own-world comparisons. No independent GPU rerun is claimed. The portable audit confirms Llama32 material bytes from the pinned public dataset.

I.1 Post-hoc selected-output compatibility

After the complete E126 confirmation, we classify actual argmax outputs in the opposite-convention world against both sources' private gold. The same-convention world cannot distinguish the two mappings and is excluded from this diagnostic. Every context and both requested sources remain included; the unit is still context. This is an additional observation, not a change to the frozen primary or a new model run.

These categories use the task's canonical code indices rather than the words' natural semantic equivalence. They show compatibility of an observed output with a demonstrated mapping, without claiming to observe an internal called rule. The

Model / output words, n=32	Native candidate / full, %	Selective candidate / full	All formation off candidate / full	Query-only candidate / full	Answer + query candidate / full	All formation + query candidate / full
Qwen3-8B / shared	84.86 / 84.86	92.19 / 92.19	75.00 / 75.00	84.47 / 84.47	93.85 / 93.85	94.34 / 94.34
Qwen3-8B / far	94.24 / 92.09	94.34 / 92.09	94.34 / 51.46	94.34 / 94.34	94.63 / 94.63	93.16 / 93.16
Mistral-7B-v0.3 / shared	88.77 / 88.77	94.14 / 94.14	73.54 / 73.54	94.73 / 94.73	96.39 / 96.39	93.16 / 93.16
Mistral-7B-v0.3 / far	98.05 / 96.00	98.05 / 94.73	95.21 / 53.71	97.27 / 97.27	97.46 / 97.46	95.61 / 95.61
Llama-3.1-8B / shared	81.25 / 81.25	90.72 / 90.72	75.20 / 75.20	97.95 / 97.95	99.12 / 99.12	99.12 / 99.12
Llama-3.1-8B / far	97.95 / 95.41	98.54 / 95.12	94.82 / 53.81	96.58 / 96.58	96.97 / 96.97	95.31 / 95.31
Qwen3-14B / shared	89.94 / 89.94	97.36 / 97.36	75.00 / 75.00	93.75 / 93.75	97.27 / 97.27	98.05 / 98.05
Qwen3-14B / far	95.90 / 95.80	97.27 / 97.07	96.68 / 52.83	92.97 / 92.97	94.63 / 94.63	96.68 / 96.68

Table 12: All real-text formation and reading policies, n32 confirmation. Candidate/full-output accuracy pairs separate within-word-space judgments from actual output choice; all namespaces are retained.

Model	Selective shared gain A, pp	B, pp	Opposite conventions, pp	Same conventions, pp	Shared-minus-separate gain, pp
Qwen3-8B	4.49 [1.56, 7.03]	10.16 [6.05, 14.26]	15.23 [11.91, 18.75]	-0.59 [-2.15, 0.98]	7.23 [5.08, 9.47]
Mistral-7B-v0.3	4.49 [0.98, 8.20]	6.25 [3.32, 9.38]	12.11 [8.20, 16.21]	-1.37 [-3.71, 0.20]	5.37 [2.73, 7.81]
Llama-3.1-8B	7.42 [4.69, 10.16]	11.52 [8.20, 14.65]	19.14 [15.04, 23.24]	-0.20 [-0.59, 0.00]	8.89 [6.84, 10.94]
Qwen3-14B	4.88 [2.15, 8.20]	9.96 [6.05, 14.06]	15.23 [11.52, 19.14]	-0.39 [-1.37, 0.59]	6.05 [4.30, 7.91]

Table 13: Real-text selective-formation gains by source, convention conflict, and output namespace, n32 confirmation. Percentage-point changes include context-bootstrap 95% CIs.

predominant change is in outputs compatible with the other source's convention. Own-space wrong-code outputs also increase in Qwen by 2.54 [0.68, 4.88] points and Llama by 0.68 [0.10, 1.27], so this is not a uniform reduction of every error type. These opposite-world categories are not a mediated fraction of the both-world primary. Full pair-specific rates and discovery results are in the post-hoc report and raw recomputation.

J Shared-output cache compatibility and source-conditioned mapping use

E127 holds the original comments, source identities, class-index conventions, and global private positions fixed across shared and separate namespaces. Two orders of the same correct public facts form whole source-owned K/V caches. Every private cross-source formation edge is blocked; every query position is denied schema access. The 00/11 pair uses one common ordering and 01/10 mixes orderings. Each cache is scored under joint and requested-source-only reading. Both private directions, sources, and all predefined worlds remain included.

The preregistered primary is **shared outputs, opposite private conventions, joint coherent-minus-mixed full-vocabulary accuracy**. Other worlds and separate word sets are secondary. Shared words cannot identify which source an answer came from: the stored `other_source_word=False` is unidentifiability, not zero leakage. Full argmax, candidate choice, and raw label logits are retained. Neither the source-average primary nor the local functions measure both answers being simultaneously correct.

J.1 Functional request differences and common shifts

For shared outputs in the opposite-convention world, define $S = \frac{1}{2}\mathbb{E}_q[g_A(q)(z_A(q) - z_B(q))]$ and $C = \frac{1}{2}\mathbb{E}_q[g_A(q)(z_A(q) + z_B(q))]$. Absolute common magnitude averages the absolute value per

Model	Far/all-formation-off full accuracy A, %	B, %	Other-source word output A, %	B, %
Qwen3-8B	48.24 [40.43, 56.05]	54.69 [45.70, 63.48]	51.17 [43.36, 58.98]	42.58 [33.40, 51.76]
Mistral-7B-v0.3	30.86 [21.68, 40.62]	76.56 [68.75, 83.98]	66.99 [56.84, 76.76]	22.27 [14.84, 30.08]
Llama-3.1-8B	64.65 [58.01, 71.29]	42.97 [34.96, 50.98]	34.96 [28.32, 41.60]	55.27 [47.46, 63.28]
Qwen3-14B	46.09 [39.65, 52.34]	59.57 [53.52, 66.02]	53.71 [47.46, 60.16]	38.48 [32.03, 44.73]

Table 14: Separate-word complete formation isolation, n32 confirmation. Full-output accuracy and other-source word rates are reported for both source requests.

Model	Single A, %	B opposite shared	B same shared	B opposite far	B same far
Qwen3-8B	95.31 [92.58, 97.66]	94.14 [90.62, 97.27]	97.27 [94.92, 99.22]	92.97 [88.28, 96.88]	94.14 [90.23, 97.66]
Mistral-7B-v0.3	98.83 [97.27, 100.00]	98.05 [94.92, 100.00]	98.44 [96.88, 99.61]	98.83 [97.66, 100.00]	98.05 [96.48, 99.61]
Llama-3.1-8B	98.44 [96.88, 99.61]	99.22 [98.05, 100.00]	99.61 [98.83, 100.00]	94.14 [89.84, 97.66]	89.06 [83.58, 94.14]
Qwen3-14B	98.05 [96.48, 99.61]	96.88 [93.75, 99.22]	98.44 [96.48, 100.00]	95.70 [93.75, 97.66]	96.48 [92.97, 99.22]

Table 15: Real-text single-source functional controls, n32 confirmation. Percent accuracy is shown for both private conventions and word sets.

query before averaging. S is also the mean correct margin over both requested sources, because their gold labels are opposite. Its change excludes a single request-independent logit offset as the complete explanation; it does not identify a hidden binding representation or a unique retrieval algorithm. Separate outputs instead use the logsumexp difference between word sets, so their functional units must not be treated as the same shared-word measure.

J.2 Local functions and incomplete joint use

Strong private functions coexist with incomplete joint selection. Llama's primary is small, its B-specific interval spans zero, and the positive effect is not uniform across private directions. Mistral's two coherent orders differ. The balanced own-only linear coherence contrast is zero by construction, whereas absolute-common observers may have nonzero contrasts because of their nonlinear definition. None of these identities demonstrates a separate latent module. All sources, worlds, private directions, single functions, native and instruction-plus-isolation controls, candidate/full scores, B-code response, and raw common shifts appear in the complete report.

Eight fixed runs contain 192 contexts and 132,096 scientific scores. The completed-run timers total 1.950864 GPU hours; they exclude preserved partial-run computation, the zero-score load failure, startup reservations, and model transfer. Scientific source `cef7952d4ef556ab ad45c9c2e6d6f351e3b0a1b3fffa9232b9eb7288 06c943de` was frozen before scoring at commit4b8ee609. The independent CPU record covers all 192 contexts and actual tokenizer/configuration layouts; the complete raw audit reports numerical recomputation, golds, masks, and actual per-layer consumption. These are independent analyses, not independent GPU reruns.

Qwen3-14B's first confirmation load failed before any scientific score after an external job occupied its GPU. The pinned checkpoint was then verified on another empty node. The initial discovery run's ten complete contexts and raw records remain preserved separately. Identical contexts were rerun in the full fixed discovery batch, without seed or outcome selection: all 6,880 matching argmax decisions and candidate signs agree across hardware; maximum absolute logit-difference and individual label-logit deviations are .001028 and .002842. These hardware comparisons are separate from the same-run FP32 controls and are not additional samples. Checkpoint identity and the operational records preserve the migration. Portable reconstruction reproduces the 32 Llama confirmation contexts byte-for-byte from the pinned public SST5 files.

Coherent-minus-mixed cache effects on held-out contexts. Left: identical words and conflicting conventions, the frozen primary. Right: separate words and both convention worlds, a secondary reference. The panels use different world scopes.

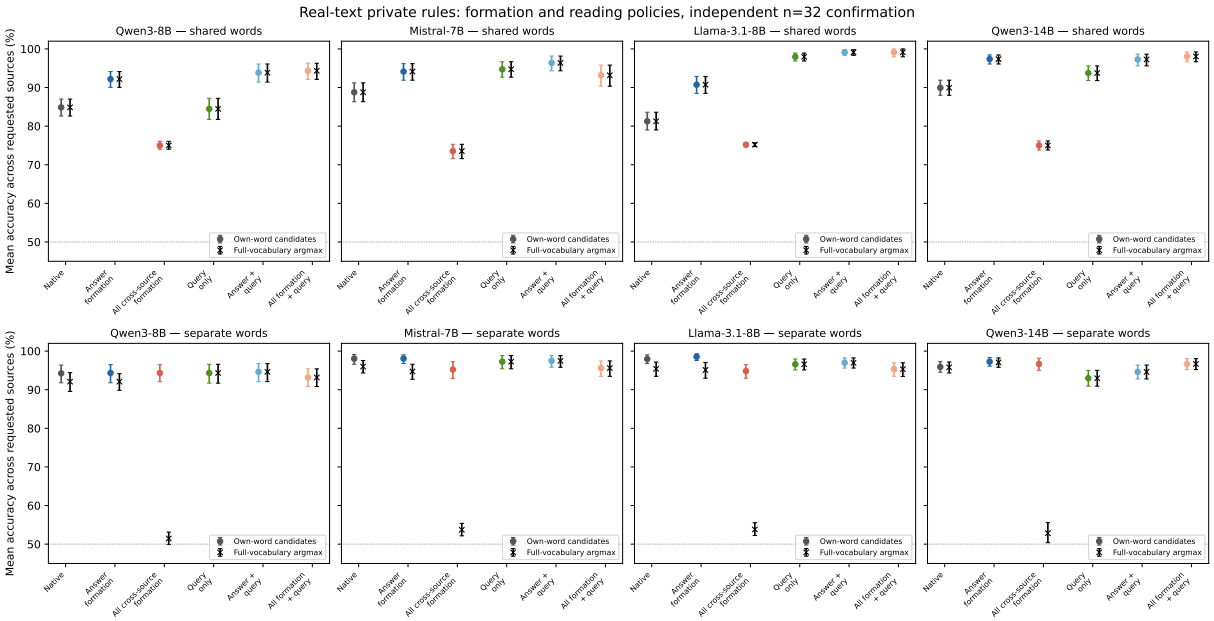


Figure 10: All real-text formation and reading policies

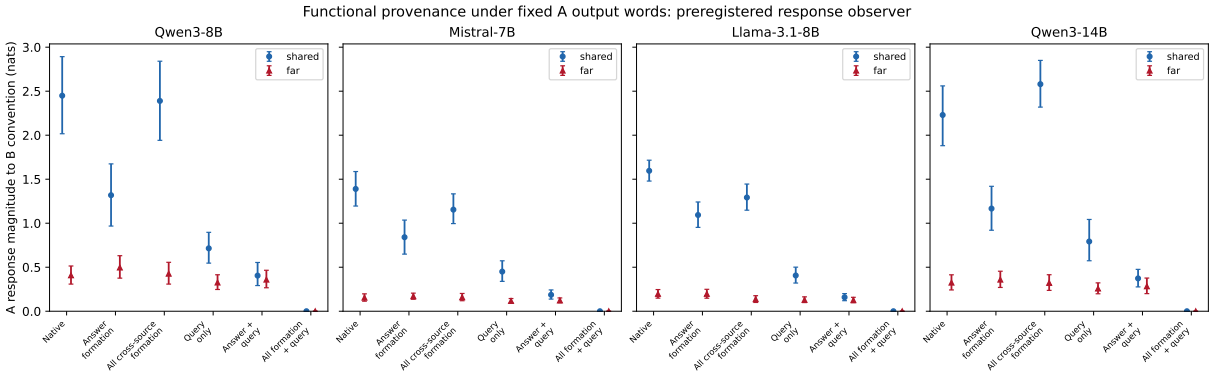


Figure 11: B convention response on fixed A words

K Compatibility when input criteria differ

E128 uses controlled restaurant comments with two independently varied aspects and both private code directions. A follows Food; B has four criterion×code worlds. Both sources share no/yes, public facts and private global positions match across orders, private cross-source formation is blocked, and the whole query cannot read the public table. The frozen primary is the origin world (A Food/B Service, same code), the four aspect-conflicting queries, averaged over both requested sources. This differs from E127's convention primary and does not inherit its stronger private-function scores.

For origin, let f and s denote the Food and Service signs oriented by A's private code. We report $T=\frac{1}{2}\,E[(zA-zB)(f-s)]$ over all eight balanced queries. This equals the average correct margin of the two sources on the four conflicting queries: it is an equivalent output projection, not an independent probe or evidence of input filtering. The complete Food/Service/bias/product basis and common shifts remain available.

K.1 Actual function scores on the primary domain

The single-function conflicting-query values are an explicitly **post-hoc descriptive scope audit** of already predefined controls, recorded by the independent reviewer. They neither change the primary nor add model runs. All-query controls in the full report have a different domain and cannot establish capability on the primary conflict cases. Qwen3-8B's Food is weak; Mistral's Food score is modest and its Service interval includes chance; Qwen3-

Model, n32	Formation identity gain, pp	Reading identity gain, pp	true F0R0 full/candidate, %	true F1R0 full/candidate, %
Qwen3-8B	4.59 [2.73, 6.64]	0.29 [-0.78, 1.46]	52.73/93.16	58.98/93.65
Mistral-7B-v0.3	-1.07 [-2.05, -0.20]	0.49 [0.10, 0.98]	53.32/91.41	52.83/92.97
Llama-3.1-8B	-0.68 [-1.95, 0.49]	-0.20 [-0.88, 0.39]	53.22/93.95	54.59/93.95

Table 16: Unlabeled source identities during independent private learning, n32 confirmation. Paired gains in percentage points and candidate/full-output controls distinguish formation from direct reading.

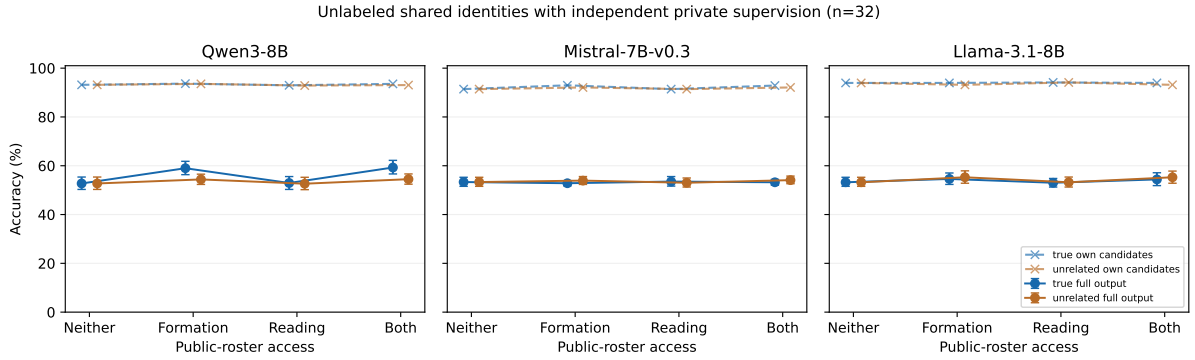


Figure 12: Shared identities at formation and reading

14B's matched Food/Service functions are above chance. Its small positive compatibility effect therefore coexists with a substantial private-to-joint gap. This prevents interpreting the weaker effects as absence of criterion information, or the positive margin as successful criterion composition.

All 85 interfaces, four worlds, private directions, source-specific intervals, both coherent orders, native and **pure one-sentence instruction** conditions, full/candidate outputs, and fixed-word response bases are preserved in the complete report and independent disclosure. Aspect order remains fixed, and whole-cache interchanges remain compatible with contextual binding, calibration, and criterion-conditioned retrieval.

Eight runs comprise192contexts/130560scores and1.237902 recorded GPU hours. Scientific source `f0a0146708e193e017eb33a1caf955d8e4d4bf22d91d73abb403b500be92265a` was frozen at5e9549ed before scoring. The independent raw audit reconstructs38760 numerical entries with maximum error0 and checks actual masks, golds, full argmax, cache origin, and the margin identity. The author report check adds38764 numerical entries and352 table rows; these comparison counts overlap the same observations and do not add experimental sample size. No independent GPU rerun is claimed.

All four cache pairs on the frozen criterion-conflict domain, confirmation n32 per checkpoint. Scores average separate source requests; they do not measure simultaneous two-source success.

L Does a pre-input source block deploy a source-specific criterion?

E129 retains the controlled Food/Service material and crosses B's criterion with its private code while fixing names, B label streams, input multisets, and global positions. Private demonstration histories form independently, without a public schema. Before any test comment, the complete `Annotator: name / Comment:` block reads both private histories; the cue stage does not use an oracle source mask. We exchange only this entire multilayer cue K/V block, preserving every recipient demonstration cache. The primary body graph then reads recipient B, HEAD, and the cue; joint-body reading is secondary. Two same-code recipients, four donor worlds, and CueA/CueB controls remain fixed.

Let $\delta t{=}\frac{1}{2}(\mathrm{zcue}(t,+){-}\mathrm{zcue}(t,-))$. Its projections on private-code-oriented Food and Service signs form a two-row functional matrix M. The two co-primary quantities are det(MCueB) and det(MCueB−MCueA), calculated separately for every recipient/context before averaging the two recipients and bootstrapping contexts. A positive difference determinant alone can result from subtracting two rank-one matrices, so the first co-primary and correct row selectivities cannot be omitted. These are output functions, not measured latent

Model / phase	Formation true−unrelated, pp	Reading true−unrelated, pp	Formation true−swapped, pp	Swapped formation gain, pp
Qwen3-8B / discovery	20.31 [16.99, 23.83]	0.39 [-0.39, 1.17]	10.16 [7.03, 12.89]	9.96 [6.25, 13.87]
Mistral-7B-v0.3 / discovery	10.94 [6.25, 16.60]	0.20 [-0.98, 1.37]	6.05 [2.34, 10.16]	6.64 [3.71, 10.35]
Llama-3.1-8B / discovery	8.98 [5.08, 12.70]	0.00 [-0.59, 0.59]	12.50 [8.79, 16.80]	0.20 [-2.54, 2.54]
Qwen3-8B / confirmation	26.66 [23.05, 30.47]	0.78 [-0.29, 1.95]	11.72 [7.91, 15.72]	16.99 [13.77, 20.31]
Mistral-7B-v0.3 / confirmation	7.71 [5.08, 10.45]	0.10 [-0.29, 0.49]	4.39 [2.15, 6.74]	6.05 [4.10, 8.11]
Llama-3.1-8B / confirmation	8.30 [6.05, 10.84]	0.29 [-0.20, 0.78]	5.47 [2.44, 8.59]	2.64 [0.10, 5.08]

Table 17: Public output-schema relations during independent private learning. Formation and reading contrasts compare correct, unrelated, and swapped relations; discovery and confirmation remain separate.

Model, confirmation n32	True F0R0 full / candidates, %	True F1R0 full / candidates, %	Formation identity gain A / B, pp	Swapped other-word change, pp
Qwen3-8B	51.27/95.12	79.98/91.31	29.88 [23.05, 36.72] / 23.44 [17.97, 29.69]	-18.75 [-21.78, -15.92]
Mistral-7B-v0.3	50.00/91.02	60.45/95.70	4.49 [-0.20, 9.38] / 10.94 [7.42, 14.45]	-4.79 [-6.74, -2.93]
Llama-3.1-8B	53.22/94.24	61.33/95.12	0.20 [-5.27, 5.66] / 16.41 [10.35, 23.05]	-3.22 [-5.47, -0.98]

Table 18: Public-schema confirmation, n32. Full-output/candidate accuracies, source-specific gains, and other-source outputs describe functional benefits and costs.

dimensionalities.

Both confirmation co-primary intervals span zero. Llama's small positive Source-specific determinant in discovery does not confirm. Qwen3-14B retains small criterion-related responses, including a mean row-selectivity interaction of0.00160 [0.00024,0.00300]nats. This does not establish correct deployment of both criteria; it also prevents claiming that the cue contains no task information.

The matched single functions are usable, while transplantation does not establish a complete criterion×code function. Every patch reports accuracy against both recipient and donor rules, so following the donor is not mistakenly judged only by recipient accuracy. The full report preserves all64 conditions, two body graphs, source cues, recipients, worlds, private directions, four response bases, normalized rotations, actual output tokens, and function controls.

Native early/late uses the same formation and reading graphs, changing only field order. Independent early prepares a cue with access to both histories and then uses an own-only body; independent late uses own-only reading throughout, so their difference also changes the graph. The pure one-sentence instruction retains native formation. In Qwen3-14B origin B it gives84.38% candidate accuracy but44.14% full accuracy, with55.08% of outputs in the label set; the actual-output cost is retained rather than described as recovery.

This is a test of cues prepared from **independently formed private histories**, not a negative result about every native task state. It remains compatible with criterion-conditioned retrieval, task information in other positions, and information becoming deployable after query input. CueA/B differs in lexical identity; donor-world subtraction cancels only additive, world-invariant name effects. General task-state transfer is already established by Hendel et al., just-in-time task representations (Li et al., 2025), and representative task-vector analysis. This interface test neither claims a first task vector nor retracts the demonstrated historical formation mechanism.

Four runs comprise96contexts/49152scores and0.413184 recorded GPU hours. Source `a177434cf425889df1fcc8b3607e6d5057a26080ad467e7999d002e6504b00bd` was frozen at57ca241a before scoring. The independent raw audit checks20580 numerical quantities, both gold definitions, actual masks and58752 prefix/27648 cue/221184 body layer-mask calls. Across all four

Model / phase	Joint coherence, pp	Candidate coherence, pp	Joint order11−00, pp	Coherence A / B, pp
Qwen3-8B / discovery	19.43 [15.82, 23.24]	1.07 [0.00, 2.15]	0.20 [-5.27, 5.47]	20.12 [13.28, 27.73] / 18.75 [14.26, 23.83]
Mistral-7B-v0.3 / discovery	15.33 [11.33, 19.53]	0.29 [0.00, 0.59]	4.88 [-2.73, 12.89]	25.39 [18.55, 33.20] / 5.27 [2.54, 8.20]
Llama-3.1-8B / discovery	7.91 [5.08, 10.84]	0.10 [-0.49, 0.68]	-0.98 [-4.88, 2.34]	9.18 [4.49, 14.65] / 6.64 [3.32, 10.16]
Qwen3-8B / confirmation	23.49 [20.80, 26.07]	1.12 [0.24, 2.05]	-4.10 [-7.62, -0.59]	21.19 [17.58, 24.90] / 25.78 [21.09, 30.47]
Mistral-7B-v0.3 / confirmation	10.40 [7.96, 13.04]	0.63 [0.29, 0.98]	-5.27 [-10.25, -0.49]	17.48 [13.18, 22.07] / 3.32 [1.46, 5.37]
Llama-3.1-8B / confirmation	5.42 [4.00, 6.79]	0.20 [-0.24, 0.63]	1.46 [-0.98, 3.91]	6.05 [4.20, 8.11] / 4.79 [2.73, 6.74]

Table 19: Equivalent public backgrounds with separate output words. Coherent-minus-mixed cache composition, order effects, and source-specific contrasts are reported in percentage points.

Model, confirmation n32	Own-only00 / 11, %	Own-only01 / 10, %	Joint other-word coherence, pp
Qwen3-8B	92.09 / 93.36	91.99 / 93.46	-25.88 [-28.66, -22.90]
Mistral-7B-v0.3	96.68 / 96.78	96.39 / 97.07	-10.64 [-13.33, -8.11]
Llama-3.1-8B	93.16 / 92.58	92.77 / 92.97	-5.76 [-7.18, -4.30]

Table 20: Local private functions under equivalent backgrounds, n32 confirmation. Own-source-only accuracies and the joint other-source output contrast separate local use from joint compatibility.

runs, the maximum full-token/cached difference is5.5313e−5nats; graph-zero and self-interchange checks are0. These are independent analyses, not independent GPU reruns.

Complete confirmation functional matrices, own-body scope. F/S denote Food/Service; arrows label the donor criterion row and query feature projection. Intervals use contexts rather than individual queries.

L.1 Source-cue transfer under different formation policies

E130 follows the E129 cue/body/readout design and holds the recipient history fixed within each newly sampled context, while preparing donors under three historical policies: native contextualization, selective removal of foreign-answer→own-label edges, and complete private-history isolation. All donors are transplanted into the same independently formed recipient caches. The two fixed co-primary determinants use native donors; every policy contrast and actual donor/recipient function is reported.

All four intervals span zero. Qwen3-14B nevertheless retains a native CueB Food row selectivity of .02144 [.00972,.03348] nats; its Service selectivity is −.01465 [−.02661,−.00353] nats. The response therefore cannot be summarized as no criterion information, or as correct transfer of both criteria. Changing formation policy did not establish the complete pre-input function at this interface. The full report retains all 136 interfaces, policy contrasts, private directions and recipient/donor full-output functions. This follow-up tests the source-cue interface; the historical formation, relation-content and native-operation results remain separate evidence.

The four fixed runs contain 96 contexts and 104448 scores, with .805577 recorded single-GPU run hours. Source 7ef59b3f… was frozen at fa41da09 before scoring. The independent audit verifies 81188 raw-derived numerical quantities and all actual prefix/cue/body mask calls; full-token/cached error is at most 2.48e−5 nats. The final report check verifies 28728 displayed numbers and retains the disclosed Llama discovery stdout reading-order exception. No independent GPU rerun is claimed.

M Matched cross-size formation, operation, and source-function tests

E131 preserves the original E121 formation and reconstruction engine, E120 native-operation engine, and E123 real-text function-policy engine. The preregistered matrix contains Qwen3-32B and Llama-3.1-70B discovery n16→confirmation n32,

Confirmation model	Coherent-minus-mixed other-source rule-compatible outputs, pp [95% CI]	Other-source outputs matching requested canonical code, pp [95% CI]
Qwen3-8B	−23.73 [−27.44, −20.02]	−1.37 [−2.25, −0.59]
Mistral-7B-v0.3	−10.06 [−12.79, −7.42]	−0.39 [−1.07, 0.20]
Llama-3.1-8B	−4.88 [−6.64, −3.22]	−0.59 [−1.17, −0.10]

Table 21: Post-hoc selected-output compatibility, confirmation contexts. Percentage-point contrasts classify other-source outputs by the private convention they match; these are diagnostic rather than prospective primary tests.

Model / phase	Primary, pp [95% CI]	Shared/all worlds, pp	Separate/all worlds, pp
Qwen3-8B / discovery	6.84 [4.10, 9.77]	3.61 [2.25, 5.08]	17.09 [14.26, 19.63]
Mistral-7B-v0.3 / discovery	5.47 [3.71, 7.81]	2.64 [1.76, 3.81]	12.60 [8.50, 16.99]
Llama-3.1-8B / discovery	0.78 [0.00, 1.76]	0.39 [0.00, 0.88]	5.47 [2.83, 8.01]
Qwen3-14B / discovery	2.54 [0.98, 4.10]	1.46 [0.59, 2.25]	20.51 [16.89, 24.71]
Qwen3-8B / confirmation	8.89 [6.64, 11.13]	4.64 [3.52, 5.71]	23.49 [21.29, 25.73]
Mistral-7B-v0.3 / confirmation	5.86 [4.10, 7.71]	2.98 [2.10, 3.91]	9.13 [6.84, 11.43]
Llama-3.1-8B / confirmation	0.78 [0.10, 1.56]	0.39 [0.05, 0.78]	6.49 [4.93, 8.06]
Qwen3-14B / confirmation	9.38 [7.32, 11.52]	4.64 [3.56, 5.76]	21.44 [18.75, 24.02]

Table 22: Prospective cache-compatibility tests with shared and separate words. The primary uses shared words and opposite private conventions; all-world contrasts retain both namespaces.

and fresh same-family 8B n32 anchors for all three stages. At this snapshot, all six 8B stages, all six 32B stages, and all three 70B discovery stages are complete and independently checked. The three remaining 70B confirmation stages are still running or pending. Confirmation statements in the main text concern the completed 32B chain.

Contexts are matched byte-for-byte **within each family and stage**. Formation and function use separately sampled SST5 contexts; native operation uses controlled restaurant contexts. Confirmation seeds are 131102/131101/131103 for these stages. Parent analyses and gold weighting remain unchanged, and comparisons do not pool different stages or treat a recipe change as an isolated capacity intervention. Each large model uses all its layers with frozen FP32 weights, eager attention, no CPU offload or quantization; Qwen3-32B uses two GPUs. Model pins, dataset revisions, actual consumed masks, and engine/parent AST checks are recorded in the E131 manifests.

These are signed B-rule contrasts, not absolute magnitudes. The restoration direction agrees across these checkpoints while the A-only restoration magnitude differs (large−small −1.368 [−2.299, −.474] nats). The 32B formation-stage signed interaction interval contains zero (−.037 [−.525, .462]), while the absolute interaction is positive (.953 [.162, 1.749]); cancellation of signed context effects does not remove the separately measured formation contribution. All twelve channel/target/background reconstructions remain reported, including partial residuals.

The 32B real-text function primary averages both requested sources, origin/cross convention worlds, and queries at the context level. Full-vocabulary native accuracy is 87.99 [86.23, 89.75]%, selective formation 95.70 [93.55, 97.36]%, and complete cross-source formation removal 75.29 [74.12, 76.27]%. The selective-minus-native gain is 7.71 [5.76, 9.57] points; A gains 7.42 and B 8.01 points. Opposite-convention gains are 15.82 points, whereas the same-convention contrast is −.39 with an interval containing zero. Selective mean accuracy is not joint correctness: the latter is 92.19%. Query-only isolation improves the mean by 5.86 [4.20, 7.52] points.

Separate-word complete isolation retains 95.31% candidate accuracy but has 52.64% full-output accuracy and 9.96% joint correctness. Thus larger size does not make within-word-space accuracy a sufficient measure of using both source functions. All directions, private controls, instructions, other-source outputs, discovery results, and 36 fresh paired contrasts are linked in the complete E131 report, raw-first audit, and each stage's analysis/full-output summary. No coefficients or conditions are chosen from the cross-size results.

The 70B discovery formation batch (n16) retains 100% A single-source full-output accuracy in both private directions while exhibiting a signed formation effect of .993 [.612, 1.373] nats. Its

Confirmation model	Shared S coherence, nats	Signed common coherence	Absolute common coherence
Qwen3-8B	0.63 [0.53, 0.74]	0.00 [-0.06, 0.06]	-0.05 [-0.10, -0.01]
Mistral-7B-v0.3	0.13 [0.11, 0.16]	0.04 [0.01, 0.07]	0.00 [-0.02, 0.02]
Llama-3.1-8B	0.03 [0.02, 0.03]	0.00 [-0.00, 0.01]	0.00 [-0.00, 0.01]
Qwen3-14B	0.50 [0.44, 0.56]	-0.00 [-0.07, 0.07]	-0.02 [-0.07, 0.03]

Table 23: Source-request differences and common shifts, n32 confirmation. Signed and absolute coherence contrasts are in nats; the request-dependent contrast cancels a request-independent shift.

Confirmation model	Shared/opposite joint00 / 11, %	Shared/opposite own-only00 / 11, %	Joint order11−00, pp
Qwen3-8B	63.87 / 64.84	90.82 / 91.41	0.98 [-2.34, 4.10]
Mistral-7B-v0.3	59.18 / 56.25	96.88 / 96.09	-2.93 [-5.08, -0.78]
Llama-3.1-8B	53.52 / 52.93	98.05 / 97.85	-0.59 [-2.34, 1.17]
Qwen3-14B	63.09 / 60.55	97.46 / 97.85	-2.54 [-5.47, 0.20]

Table 24: Private functions and joint reading under conflicting conventions, n32 confirmation. Percent mean accuracy under coherent backgrounds and the paired order effect are reported separately.

label-interface KV reconstruction has raw r=.542 with shared-word B-position caches and .854 with separate-word B-position caches; both permit full-source query reading. Raw RMSEs are .632/.825 versus zero-prediction RMSEs .705/1.472. All K/KV/V targets remain reported: V has lower RMSE than KV in the first background, and K slightly exceeds the zero baseline there. Fixed-word answer relational influence has excess magnitude .237 [.116, .368] nats and a negative signed contrast −.229 [−.382, −.085]. This discovery result identifies contributions under the fixed interfaces; the complete cross-size conclusion awaits the unchanged confirmation tests.

A first TF32-accelerated implementation failed the existing cached/full .001-nat tolerance before any scientific row; those empty failed attempts and their resource use remain archived. Restoring full FP32 GEMM preserved all materials, operators, and analyses. Valid completed stages have errors at most 6.68e−5 nats with self/graph checks zero. Independent analysis reproduces the original raw readouts and actual graph/gold layouts; it is not a second GPU execution.

The 70B discovery native-operation batch (n16) has a signed full-source answer-removal effect of 1.542 [1.096, 1.970] nats. Restricted restoration is positive under A-only reading (1.011 [.637, 1.401]) and negative under full-source reading (−.440 [−.653, −.208]). Cross-versus-within removal is 2.560 [2.065, 3.059] nats. A single-source full-output accuracy is 100% in both private directions. This adds a native-operation contribution in the large checkpoint, with the same reading-dependent restoration signs as the confirmed 32B chain; it remains a discovery result pending the fixed new-context confirmation.

In the 70B discovery real-text function batch (n16), shared-word full-output accuracy increases from 86.72 [83.20, 90.23]% to 98.44 [96.48, 100.00]% under selective formation, a paired gain of 11.72 [8.01, 15.43] points. A gains 5.86 and B 17.58 points; opposite-convention gains are 23.44 points and same-convention gains are zero. Selective formation exceeds complete cross-source restriction by 23.44 [21.88, 25.00] points. Query-only restriction gains 10.55 [5.66, 15.04] points. Selective mean accuracy differs from joint correctness (97.27%). With separate words, complete cross-source formation restriction retains 99.02% candidate accuracy but gives 51.17% full-output accuracy and 3.91% joint correctness. The corresponding other-source-word rate is 48.24%, with source-dependent rates retained in the complete report. Single-source functions score 95.31–100% across sources, conventions, and namespaces. This completes the discovery chain; all three new-context confirmation stages remain fixed and pending.

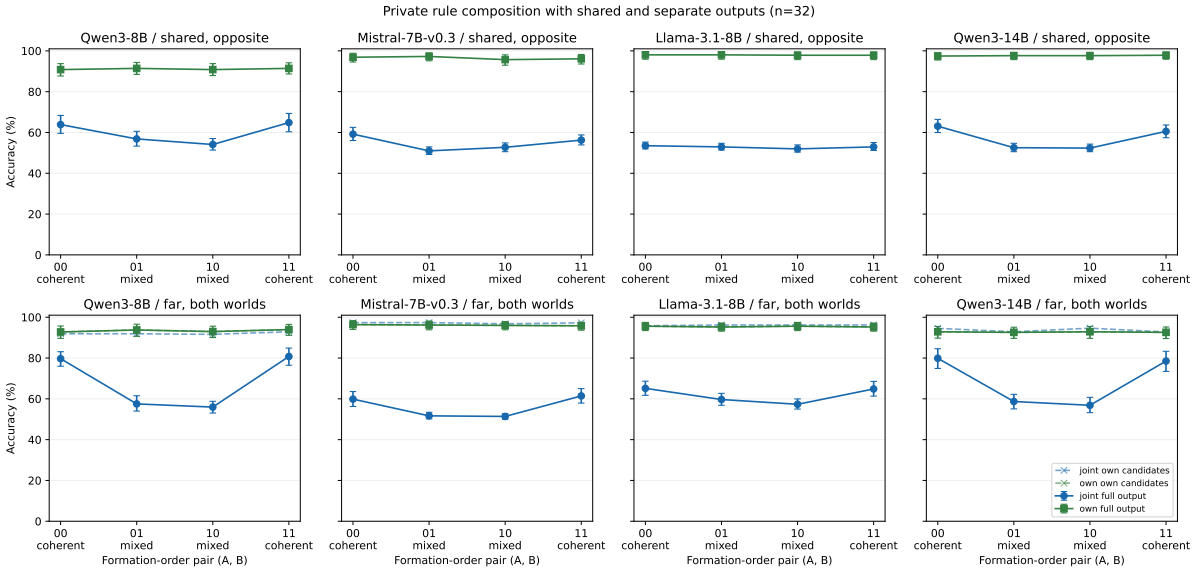


Figure 13: All cache pairs and private controls

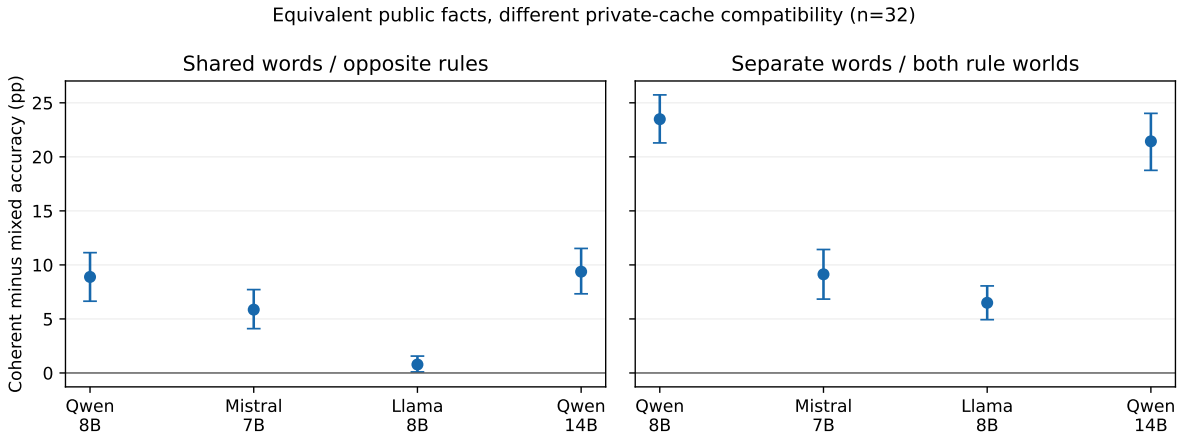


Figure 14: Shared-word rule use and separate-word use

Model / phase	Primary coherence, pp [95% CI]	Criterion-oriented margin coherence T, nats [95% CI]
Qwen3-8B / discovery	2.34 [0.39, 4.69]	0.03660 [0.01877, 0.05657]
Mistral-7B-v0.3 / discovery	-0.78 [-1.95, 0.00]	0.00853 [0.00412, 0.01442]
Llama-3.1-8B / discovery	1.56 [0.00, 3.52]	0.00767 [0.00501, 0.01085]
Qwen3-14B / discovery	0.00 [-1.95, 2.34]	0.04048 [0.02804, 0.05418]
Qwen3-8B / confirmation	1.56 [-0.59, 3.52]	0.04256 [0.03058, 0.05624]
Mistral-7B-v0.3 / confirmation	0.98 [-1.37, 3.32]	0.00502 [0.00303, 0.00714]
Llama-3.1-8B / confirmation	0.78 [0.00, 1.76]	0.00837 [0.00624, 0.01068]
Qwen3-14B / confirmation	2.93 [0.59, 5.47]	0.05037 [0.03974, 0.06198]

Table 25: Cache compatibility when input criteria differ. The shared-word conflicting-input primary and criterion-oriented margin are reported with context-bootstrap 95% CIs.

Confirmation model	Joint00 / 11, % [95% CI]	Own-only00 / 11, % [95% CI]	Single Food / Service on conflicting queries, % [95% CI]
Qwen3-8B	50.78 [48.05, 53.91] / 52.73 [49.61, 55.86]	56.25 [51.95, 60.55] / 58.20 [53.91, 62.89]	46.88 [40.62, 52.34] / 69.53 [62.50, 77.34]
Mistral-7B-v0.3	50.39 [48.05, 52.73] / 50.00 [48.05, 51.95]	54.69 [51.95, 57.81] / 55.47 [51.95, 59.38]	57.81 [53.12, 64.06] / 53.12 [49.22, 57.81]
Llama-3.1-8B	51.17 [50.00, 52.73] / 51.95 [50.39, 53.52]	70.70 [64.84, 76.95] / 70.31 [64.45, 76.56]	67.19 [60.16, 74.22] / 71.09 [63.28, 78.91]
Qwen3-14B	52.34 [48.83, 55.86] / 51.95 [49.61, 54.30]	75.00 [71.09, 78.91] / 73.05 [69.14, 76.95]	68.75 [61.72, 75.80] / 78.12 [70.31, 85.94]

Table 26: Actual criterion-function scores on conflicting inputs, n32 confirmation. Joint reading, own-source-only reading, and single-source food/service accuracy have distinct scopes.

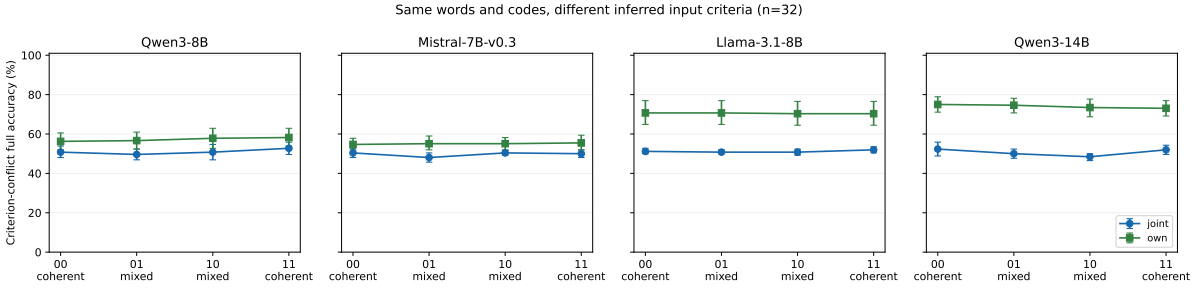


Figure 15: Criterion-conflict cache compatibility

Model / phase	CueB det, nats² [95% CI]	Source-specific det, nats² [95% CI]
Llama-3.1-8B / discovery	-5.914e-06 [-1.213e-05, 5.864e-07]	6.157e-06 [2.463e-06, 1.080e-05]
Qwen3-14B / discovery	-9.656e-06 [-1.251e-04, 8.155e-05]	5.551e-06 [-1.524e-05, 2.613e-05]
Llama-3.1-8B / confirmation	2.492e-06 [-5.264e-06, 1.097e-05]	-1.259e-06 [-3.316e-06, 4.967e-07]
Qwen3-14B / confirmation	3.740e-05 [-1.643e-05, 1.014e-04]	-1.148e-05 [-4.553e-05, 1.924e-05]

Table 27: Pre-input source-cue criterion transfer across discovery and confirmation. Context-wise determinants report the full CueB response and the source-specific difference.

Confirmation model	Single Food early / late conflict accuracy, % [95% CI]	Single Service early / late conflict accuracy, % [95% CI]	Native B early / late origin conflict accuracy, % [95% CI]
Llama-3.1-8B	68.75 [60.94, 76.56] / 69.53 [62.50, 77.34]	71.09 [64.84, 77.34] / 74.22 [67.97, 80.47]	60.94 [51.56, 69.53] / 63.28 [55.47, 71.09]
Qwen3-14B	78.12 [70.31, 85.16] / 78.91 [71.88, 85.94]	76.56 [68.75, 84.38] / 79.69 [71.88, 87.50]	67.19 [60.94, 74.22] / 64.84 [56.25, 73.44]

Table 28: Source-cue private-function controls, n32 confirmation. Early/late serialization and food/service criteria are evaluated on conflicting-polarity queries.

Confirmation model	Native CueB det, nats² [95% CI]	Native Source-specific det, nats² [95% CI]
Llama-3.1-8B	−1.751e−5 [−4.604e−5,3.566e−6]	2.940e−6 [−2.280e−5,3.356e−5]
Qwen3-14B	−1.821e−4 [−6.037e−4,2.459e−4]	−3.650e−4 [−9.933e−4,2.110e−4]

Table 29: Pre-input source-cue follow-up under native formation, n32 confirmation. CueB and source-specific determinants retain both weak and reverse effects.

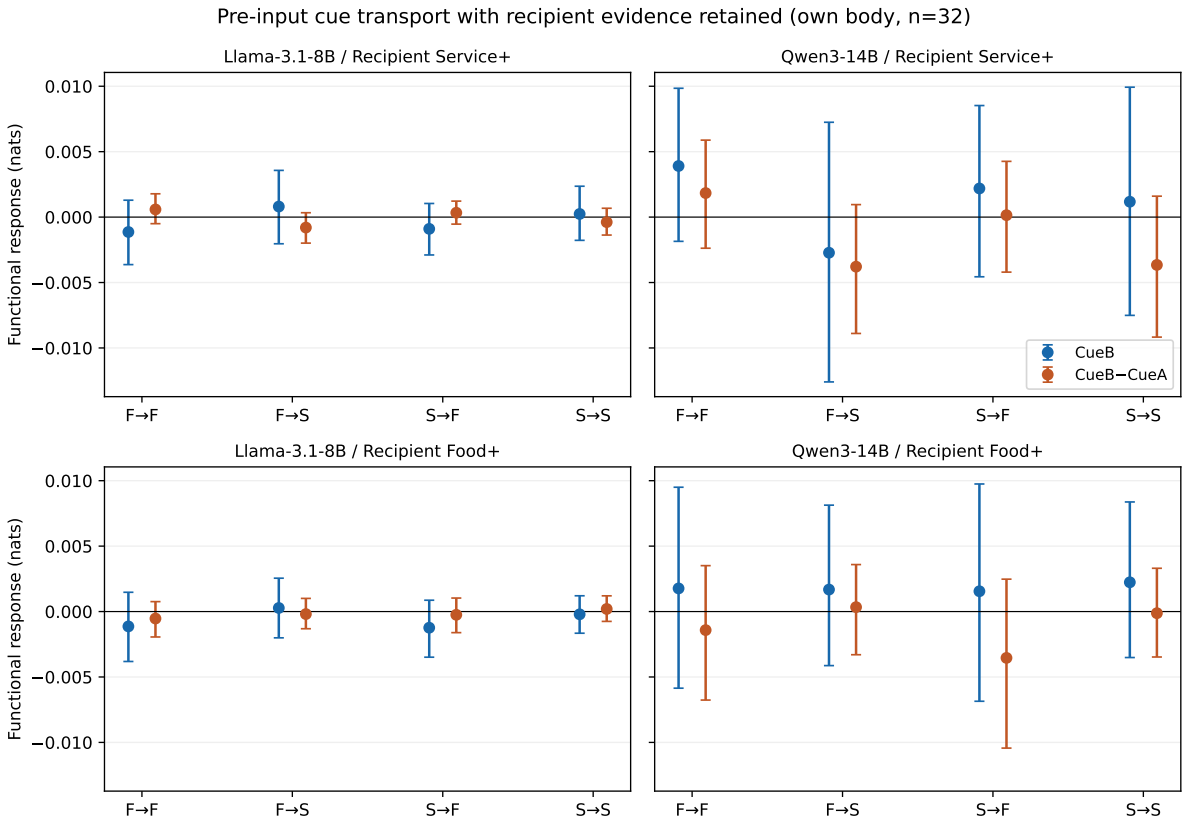


Figure 16: Functional cue information

Qwen3, fresh confirmation n32	Full-source answer removal, nats	A-only restricted restoration, nats	Full-source restricted restoration, nats
8B	1.762 [1.264, 2.294]	2.143 [1.426, 2.893]	−.650 [−1.269, −.080]
32B	1.972 [1.467, 2.500]	.775 [.321, 1.284]	−.544 [−.833, −.276]

Table 30: Fresh matched Qwen3 confirmation across sizes, n32. Signed native removal and restoration effects in nats distinguish A-only from full-source reading.