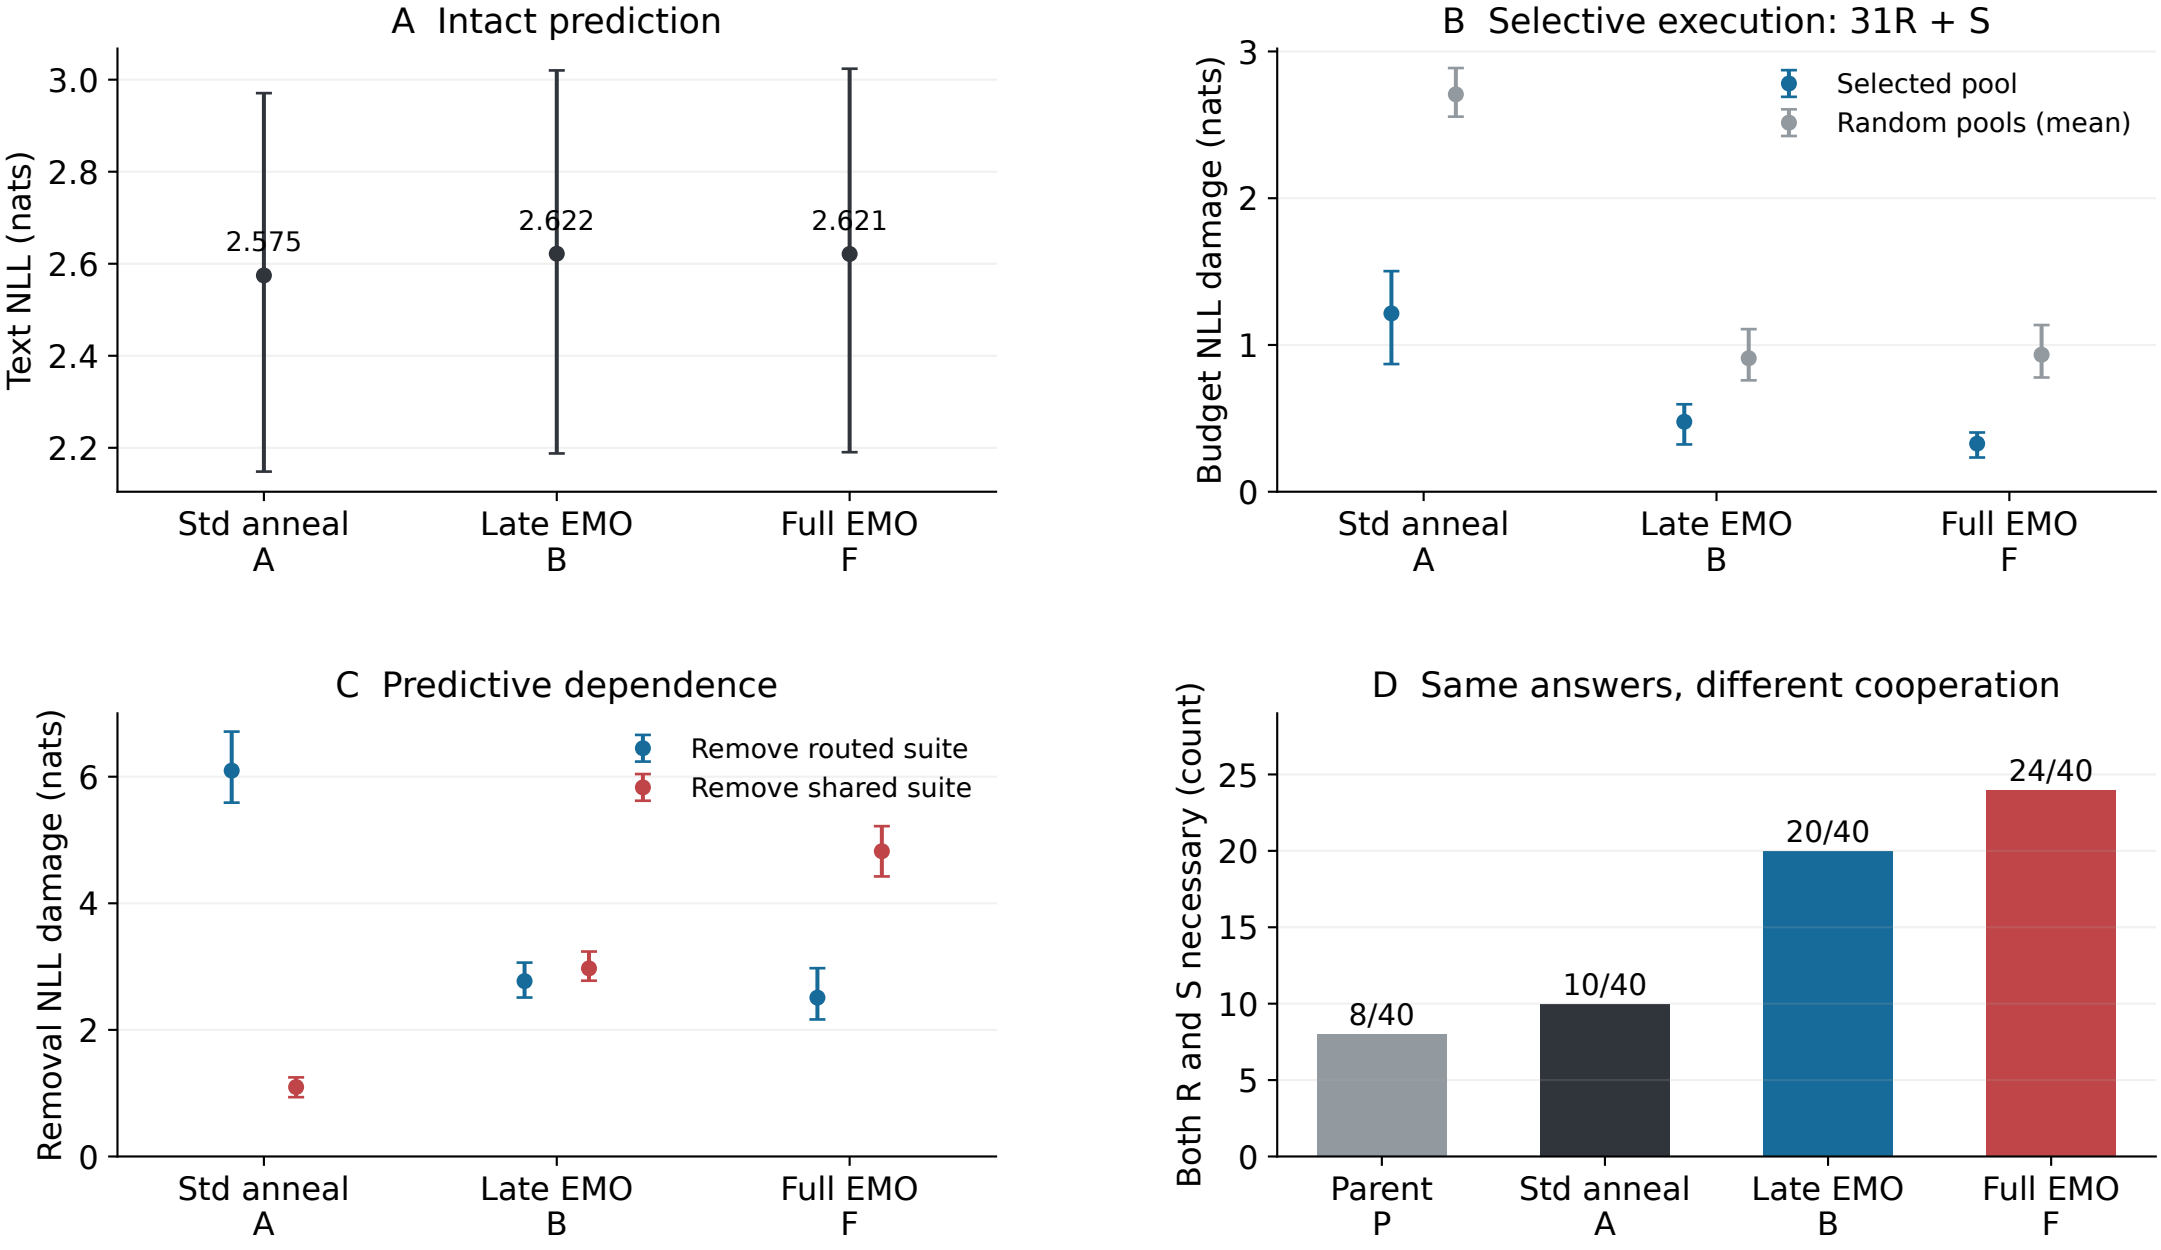

A-C: 64 language windows /20 file clusters, original 95% input CIs; B uses each model's 128-token prefix pool.
D: fixed 40 questions correct in all four models; selected-set counts, no CI. R/S cover all 16 layers.
B/F intact text NLL 2.6218 /2.6212; task raw accuracy 39.06% each.