Computational dependence changes around unchanged expert functions

Expert function suites span all 16 layers; weights are fixed across the compared contexts.

A Language modeling · shared suite

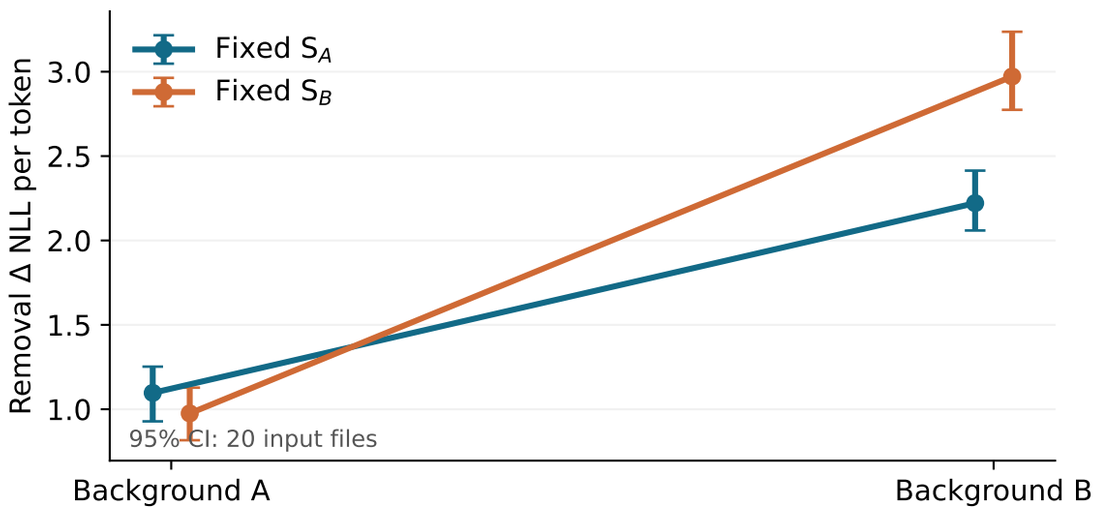

B Language modeling · routed suite

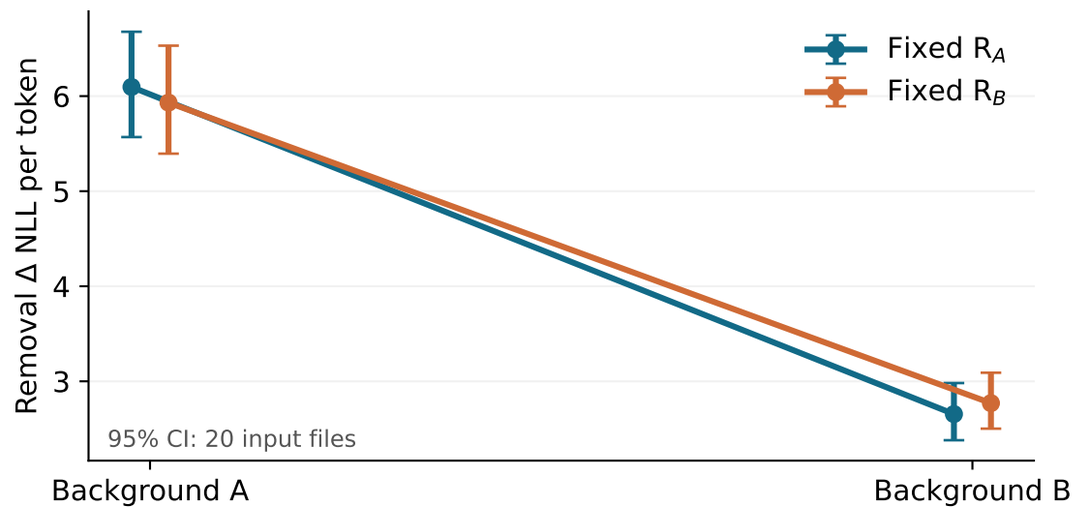

C Fixed RC tasks · same S_B; 33 full-correct cases

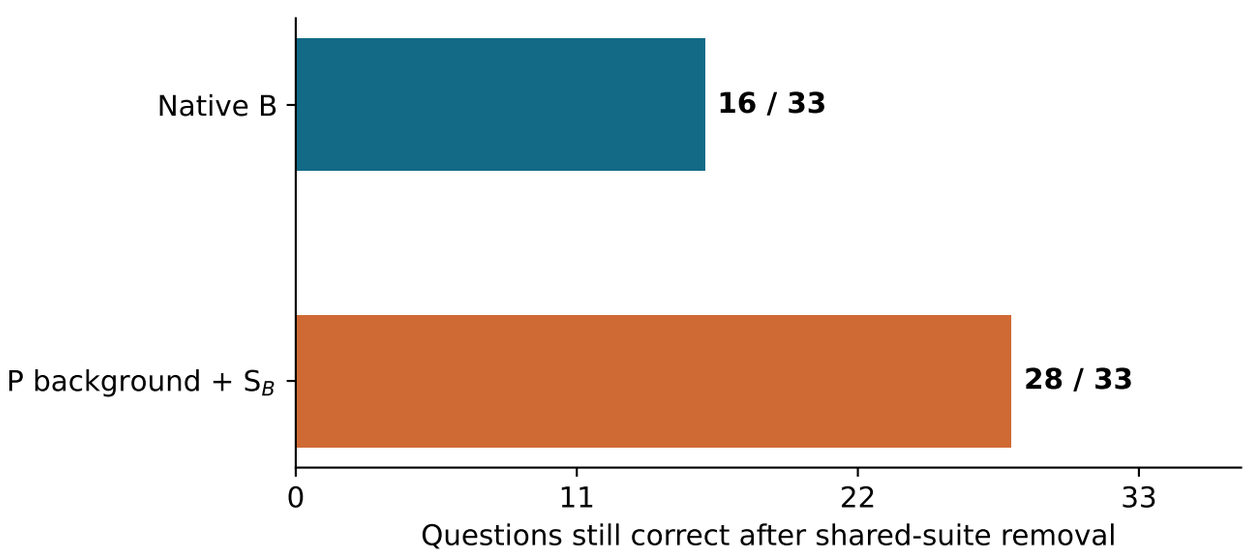

33 questions: 30 biology, 3 mathematics; no CI.
Full true-answer token NLL: native B 2.0078, P + S_B 2.1739.
Same correct decisions do not imply equal distributions.

D Fixed RC tasks · 128 fixed questions

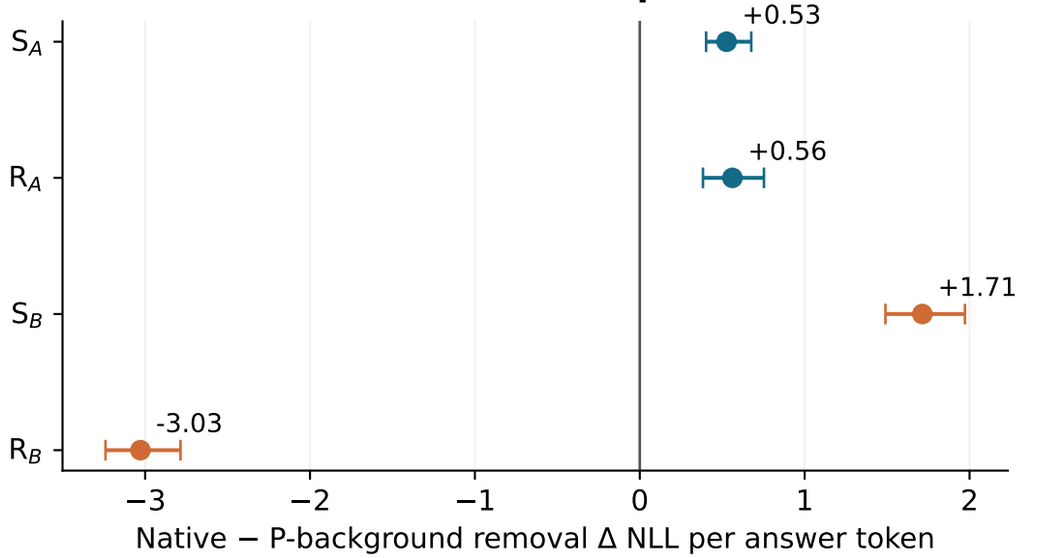

95% CI: paired question bootstrap, 64 per subject;
equal biology / mathematics macro.

POST-HOC: E150/E151/E155/E156 reuse; inherited zero-endpoint identity; 0 new forwards; one public lineage.
Context changes dense backbone, router and other expert branch jointly; no isolated partner effect.
JSON SHA-256: MMLU fec931d344cd | LM 6637e9696a37