Input-based selection across both normalization rules

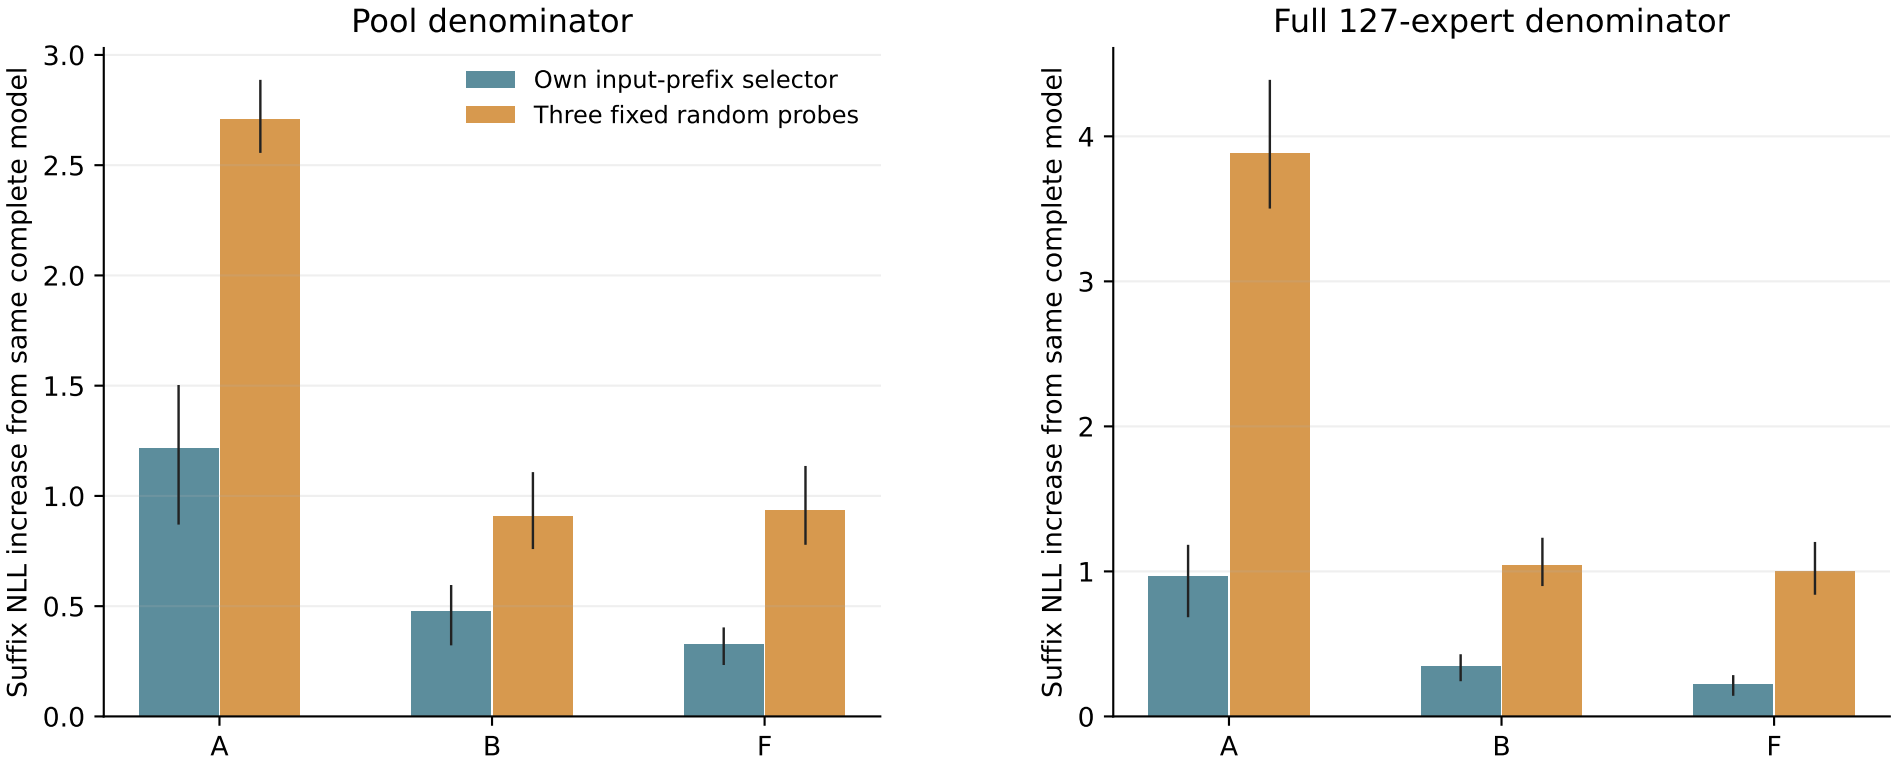

Fixed public checkpoints and inputs; no new training. Input CIs do not quantify training-seed uncertainty.
F and B share the declared final EMO recipe; original initialization and optimizer/batch chains are unverified.
32 slots = 31 routed + shared; prefix-only selection, suffix-only scores. Both rules and all three random probes retained.
P is absent from this common-rule figure because its full-denominator counterfactual was not measured.