Native task computation across mature training paths · True-answer token NLL

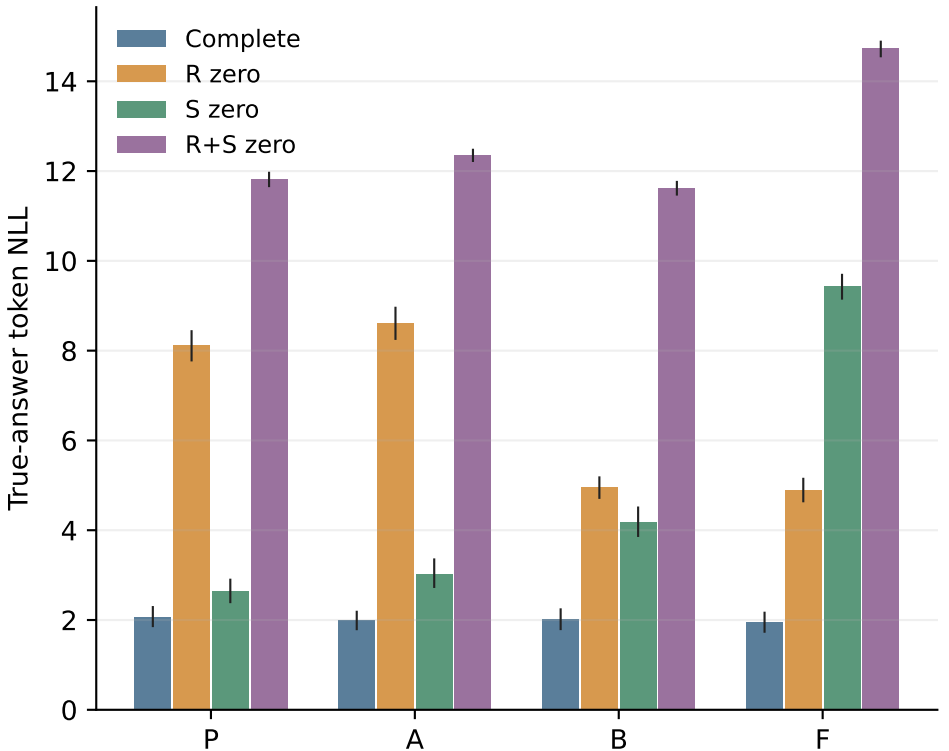

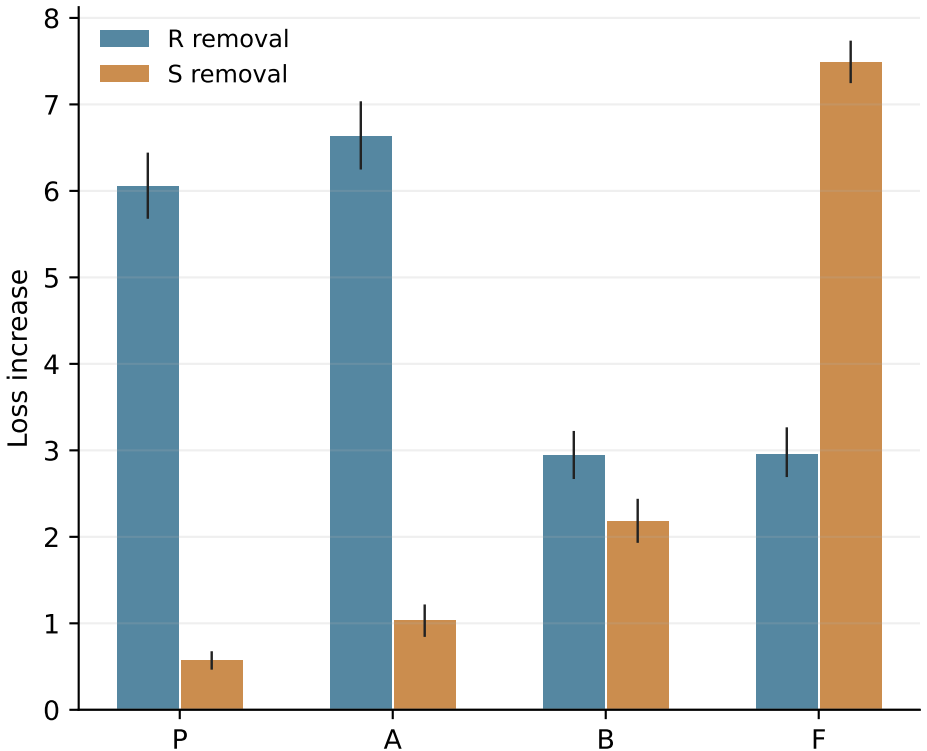

Fixed public checkpoints and inputs; no new training. Input CIs do not quantify training-seed uncertainty.
F and B share the declared final EMO recipe; original initialization and optimizer/batch chains are unverified.
128 questions, 64 per subject; paired within-subject bootstrap 95% CI. All four conditions retained.