

Which Head Takes Which Role? Initialization Leaves a Trace in Attention-Head Specialization

Anonymous ACL submission


Abstract

Attention heads in a trained language model specialize into roles such as induction, previous-token and retrieval heads, but which head takes which role differs between training runs. How much of this difference comes from the random initialization and how much from the training data is unknown, because studies of training variability change only one of the two. We separate them with public model suites that train the same initializations on different corpora: DataDecide (three initializations, up to 25 corpora, 4M–1B parameters) and Pythia (one initialization, two versions of the Pile, 70M–12B). For nine head roles, models broadly agree on the layer that hosts a role, while which head within that layer takes it follows the initialization rather than the corpus: models that share an initialization agree on it across natural-language corpora, and models that share only a corpus do not. In controlled pretraining, this assignment is made during a short early stretch of training: perturbing the weights reassigns roles only within the first few percent of training, while switching the corpus rewrites a share of the assignment that shrinks the later the switch comes. How much of it two corpora share follows the similarity of their content words and the learning rate per batch token. The initialization’s effect on head roles carries across corpora, but its effect on behaviour does not: no initialization performs consistently better across corpora on 10 benchmarks, and removing Flan instruction data (1–2% of the corpus) changes how models use a counterfactual context after the template Question: while leaving their head layout as intact as other data ablations do.

1 Introduction

Mechanistic interpretability describes a model through the roles of its components: an induction head completes a pattern seen earlier in the context (Olsson et al., 2022), a previous-token head attends to the preceding position, and retrieval heads

(a) Which head takes a role, 75 models at 1B



(b) Layer vs. head within the layer (9 roles)




Figure 1: **Models that share an initialization agree on which heads take each role.** (a) Agreement on *which* head in a layer takes a role (within-layer rank correlation of head scores, mean over induction, previous-token and retrieval maps) between all 75 one-billion-parameter DataDecide models, three initializations $\times$ 25 corpora, ordered by initialization. Bright blocks: models that share an initialization agree on every corpus; models that share only a corpus do not. $\blacktriangleright$: six runs initialized differently from their label (§3), excluded from all means. (b) The two levels of a role map, mean over nine roles: the layer profile is similar for all pairs, the head within the layer only for pairs that share an initialization.

copy a fact from the context into the prediction (Wu et al., 2025). When a model is trained again, its algorithms tend to recur while the components that implement them change (Chughtai et al., 2023; Gurnee et al., 2024; Tigges et al., 2024; Bali et al., 2026). Two sources of this change are confounded. Studies of seed variability retrain on a fixed corpus (Sellam et al., 2022; van der Wal et al., 2025;

Bali et al., 2026), and studies that relate mechanisms to data vary the data under one initialization or many unrelated ones (Chan et al., 2022; Singh et al., 2024; Chen et al., 2026). Neither design can tell whether the head that takes a role is chosen by the initialization, by the data, or by neither.

We separate the two with a natural experiment that public pretraining suites already contain. DataDecide (Magnusson et al., 2025) trains three initializations on up to 25 corpora at each of 14 sizes, and Pythia (Biderman et al., 2023) trains one initialization on the Pile and on the deduplicated Pile at every size up to 12B parameters. Comparing models that share an initialization but not a corpus with models that share a corpus but not an initialization shows how much each factor contributes. For nine head roles we ask two questions: which layer hosts the role, and which head within that layer takes it. Heads within a layer are interchangeable, so independently initialized models agree on the second question only by chance, which gives every comparison a baseline of zero (§2).

The two questions have different answers (Figure 1). All models largely agree on the layer. Only models that share an initialization agree on the head within the layer, and they do so across corpora: under one initialization, the same head of layer 7 is the strongest induction head on 14 of 25 corpora (chance: 1.6). The agreement holds for all nine roles, at all 14 DataDecide sizes and in Pythia up to 12B, and it is specific enough to tell which initialization a model came from. Comparing layouts without labels, we found six DataDecide checkpoints that, as their released weights confirm, were initialized differently from their label (§3).

We then ask when the assignment is made and what decides how much of it carries over to another corpus. In controlled pretraining, switching the corpus or perturbing the weights during the first 1% of training reassigns the roles; a few percent later, a model returns to its assignment even after noise as large as its weights, and a switch of corpus rewrites only part of it, less the later it comes (§4). How much of the assignment two corpora share follows the similarity of their content words, and how strongly the initialization decides it at all follows the learning rate per batch token, which also explains why larger DataDecide models agree more (§5).

Finally, we ask whether the initialization’s effect on heads carries over to behaviour (§6). How strong a circuit becomes, when it forms and how

a model scores on 10 benchmarks follow the corpus, and no initialization does consistently better across corpora. A case study on one data component makes the division concrete. Removing Flan, 1–2% of Dolma 1.7, changes how models use a counterfactual context introduced by Flan’s template Question:, while the head layout of these models stays as close to its seed-matched siblings as under any other change of corpus. The attention that Flan adds to the context after Question: is spread over the heads that already attend to it: the data changes what the existing retrieval heads do, not which heads they are.

Our contributions are:

- A seed $\times$ corpus design built from public suites at 21 sizes from 4M to 12B parameters, with a measurement that separates which layer from which head hosts a role and a permutation-symmetry baseline (§2).
- Evidence that which head within a layer takes each role follows the shared initialization rather than the corpus, at every size and for every role we measure, together with an audit of the suites’ initializations (§3).
- Controlled experiments that locate when the assignment is fixed and identify corpus content and the stochastic-gradient noise scale as what determines how much of it is kept (§4, §5).
- Evidence that this structural effect does not carry over to behaviour, with a case study of a template habit that 1–2% of the pretraining data strengthens while the head layout is preserved, and its consequences for knowledge-conflict evaluation (§6).

2 Setup and Measurement

Initializations crossed with corpora. DataDecide pretrains OLMo-style models at 14 sizes (4M–1B parameters) on 25 corpus recipes—C4, DCLM and its quality-filtered variants, Dolma 1.6/1.7 and their ablations, Falcon RefinedWeb variants, FineWeb-Edu/-Pro and mixtures—each with three seeds. Within a size, a seed’s initial weights are shared across recipes, so each initialization has been trained on up to 25 corpora. Pythia trains a standard and a deduplicated model from identical initial weights at every size; we use 70M–12B, with the independently initialized PolyPythias (van der Wal et al., 2025) as a reference. A DataDecide seed sets both the initial weights, which we check for ev-

ery run (Appendix A), and the data order. The order is not tied to any head and yields different batches on different corpora, so of what seed-matched runs share, only the initial weights can favour a particular head. Model pairs fall into three types: *same initialization, other corpus* (seed-matched pairs); *other initialization, same corpus*; and both different.

Head-role maps. For each model we compute a [layer $\times$ head] map for each of nine roles from fixed probe material (Appendix B): induction, previous-token, retrieval (attention from the answer position to the context entity in knowledge-conflict prompts), attention sink (Xiao et al., 2024), duplicate-token, current-token, two-back and delimiter heads, and each head’s OV copying score, which is computed from the weights alone (Elhage et al., 2021).

Two levels of agreement. A role map says which *layer* hosts a role and which *head* within the layer takes it. We compare two models on the first with the rank correlation of their per-layer profiles and on the second with the *which-head agreement*: the Spearman correlation of head scores within each layer, averaged over layers.

The chance baseline. Permuting the heads of a layer, together with their query, key, value and output weights, leaves the network’s function unchanged; the initialization distribution and gradient-based training respect this symmetry. Two independently initialized models therefore agree on which head takes a role only by chance, and their expected which-head agreement is exactly zero, whatever their corpora (Lemma 1, Appendix C). Agreement above zero between models trained on different corpora thus measures what their shared initialization carries through training.

Behaviour. We use DataDecide’s 11 benchmark evaluations, circuit strength and developmental timing, and *context trust*: on knowledge-conflict prompts that state a counterfactual fact, the log-probability margin of the counterfactual over the memorized answer, on facts the model knows (Li et al., 2025; Du et al., 2024).

3 Heads Follow the Initialization

Same initialization, same heads. In the 1B crossing, models with the same initialization agree on which head within a layer takes each role,

(a) DataDecide 1B, 3 seeds $\times$ 25 corpora (69 models)

Role	same init.	other init.	gap, 95% CI
Two-back	0.43	0.00	[0.40, 0.46]
Previous-token	0.42	−0.01	[0.40, 0.47]
Duplicate-token	0.36	−0.02	[0.36, 0.41]
Induction	0.33	−0.02	[0.32, 0.37]
Current-token	0.32	−0.02	[0.31, 0.36]
Retrieval	0.31	0.00	[0.28, 0.35]
Delimiter	0.22	−0.01	[0.20, 0.25]
Attention sink	0.18	−0.01	[0.16, 0.23]
OV copying (weights)	0.25	−0.01	[0.25, 0.29]

(b) Pythia, one initialization: Pile vs. deduplicated Pile

Size	induction	prev.-token	retrieval	null
70M	–	0.60	0.58	0.26
160M	0.49	0.58	0.53	0.14
410M	0.47	0.55	0.45	0.09
1B	0.56	0.72	0.58	0.16
1.4B	0.48	0.54	0.50	0.09
6.9B	0.53	0.60	0.49	0.06
12B	0.38	0.51	0.38	0.04
12B, at 2.1%	0.55	0.60	0.45	0.04

Table 1: **Models with the same initialization agree on which head takes each role.** Which-head agreement (within-layer Spearman of head scores). (a) All nine roles at 1B for models with the same initialization trained on different corpora and with different initializations trained on the same corpus; the gap’s interval is a recipe bootstrap. (b) Pythia’s two models from one initialization; null: 95th percentile after permuting heads within layers; no induction head at 70M.

while models with different initializations agree at chance, even on the same corpus (Table 1a, Figure 1). This holds for all nine roles (0.18–0.43 against −0.02 to 0.00), including the copying score computed from the weights alone, with five initializations at 1B, and on corpus pairs that share no data source. Agreement is not limited to the ranking of weak heads: in the three layers where a role is strongest, the strongest head of two seed-matched models is the same one 2–5 times as often as chance (12–30% against 6%; Appendix B). When every head’s within-layer rank is decomposed into initialization, corpus and residual components, the initialization is significant for 66–79% of heads, that is, an initialization favours the same heads across corpora; the corpus main effect, which the symmetry rules out, is at the false-positive rate (4–7%). The agreement holds at all 14 DataDecide sizes and grows with size (Figure 2a; Spearman with log parameters 0.81–0.92). In Pythia, the two models from one initialization agree at every size from 70M to 12B, far above a permutation null (Table 1b), while independently initialized PolyPythias agree at chance.




Figure 2: **The trace across sizes, as an identifier, and in other quantities.** (a) Which-head agreement (mean over induction, previous-token and retrieval maps) for DataDecide pairs with the same initialization and different corpora and with different initializations and the same corpus, at 14 sizes, and for Pythia’s two models from one initialization, up to 12B. (b) Identifying a model’s initialization from its head layout, leaving its corpus out, at the end of training and at 3–5% of training. (c) Similarity gained over unrelated models by sharing the initialization or the corpus: the initialization is shared in index-aligned quantities, the corpus in unit-invariant ones.

Depth is shared across models. The layer profile of each role agrees at 0.4–0.7 for every pair type (Figure 1b), and all 75 1B models implement in-context copying with the same two-step algorithm, a previous-token head composing with an induction head. Where in depth a role lives and how it is computed are largely common to all models; the initialization, not the corpus, biases which of the interchangeable heads in a layer carries it.

The layout is specific to the initialization. We hold out a model, compare its head layout with the models trained from each initialization on the *other* corpora, and assign it to the closest one. From 60M parameters on, this identifies the initialization with 98–100% accuracy (chance: 33%, or 20% with five seeds), including all 69 verified 1B models, and already at 3–5% of training (Figure 2b, Table 6); in Pythia, the true sibling ranks first among ten candidates at all three sizes with independently initialized alternatives. Applied without labels to the 75 1B models, the layout singles out six runs that agree with none of their nominal siblings but pairwise with each other (Figure 1a, ►). Their released weights confirm that these runs were not initialized from their seed’s weights (correlation 0.000) and that each pair shares an initialization of its own (Table 5).

What carries the trace. Seed-matched models on different corpora have drifted far apart: their weights correlate at only 0.01–0.04 (other pairs 0.000), enough to identify the initialization (Table 7), they are not linearly connected (loss barrier 7.5 nats, against 7.8 for unrelated models; Frankle

et al., 2020), and under unit-invariant measures, CKA (Kornblith et al., 2019) and best-match similarity, their representations are no closer than those of unrelated models (Figure 2c). What they share are index-aligned quantities: which head hosts a role, and which residual dimensions carry most variance (0.20 vs. 0.00) and the massive outlier activations (0.17 vs. 0.00). The head layout is the functional form of this trace.

4 When the Assignment Is Fixed

Intervening during training. We pretrain small Llama-style models on C4 and branch them at step k of 10,000, either switching the corpus to source code or adding Gaussian noise to every weight (Figure 3a, Table 2). Up to $k = 100$ (1% of training), both interventions reassign the roles: after a switch to code, agreement with the uninterrupted run is 0.09, and noise as large as the weights leaves 0.03 at step 0 and 0.36 at step 100. From $k = 250$ (2.5%) on, the model returns to its assignment after the same noise (0.85 at 2.5%, 0.95–0.98 from 5%, the level of an exact rerun). A switch to code still rewrites part of the layout, a smaller part the later it comes: agreement is 0.60 at 2.5%, 0.71 at 10% and 0.95 at 80%. A 12-layer model behaves alike: a switch to code at 1% keeps 0.40 of the layout and at 10% 0.84, while training on code from the start keeps none.

What closes the window. The early steps are also the warm-up of the learning rate (300 steps), and the branches inherit the optimizer’s state, so we repeat the noise branches with a 1,000-step

Step	Share	300-step warm-up				1,000-step warm-up		Optimizer reset
		Switch to code	Noise 100%	Noise 10%	Code, 12 layers	Switch to code	Noise 100%	Noise 100%
0	0%	−0.03	0.03	0.74	−0.03	—	—	—
100	1%	0.09	0.36	0.94	0.40	—	0.16	—
250	2.5%	0.60	0.85	0.97	—	0.41	0.70	0.81
500	5%	0.72	0.96	0.99	—	—	0.87	0.91
1000	10%	0.71	0.95	0.99	0.84	0.69	0.87	0.97
2000	20%	0.81	0.98	1.00	—	—	0.94	—
4000	40%	0.86	0.96	1.00	—	—	—	—
8000	80%	0.95	0.98	1.00	—	—	—	—
exact rerun					0.97			

Table 2: **Interventions during controlled pretraining.** Which-head agreement (mean of induction and previous-token maps) between the final layout of a run intervened on at a given step and the uninterrupted run, for the default 300-step learning-rate warm-up, a 1,000-step warm-up, and branches that start with a freshly initialized optimizer; means over two initializations (one for the 12-layer model). Step 0 for the code switch: the same initialization trained on code from scratch; noise is Gaussian, scaled to each tensor’s standard deviation.

warm-up and, separately, with a freshly initialized optimizer (Appendix G). With the longer warm-up the window closes later in steps (0.70 at 2.5%, 0.87 from 5%), and the two schedules align when agreement is plotted against the learning applied before the intervention, the cumulative learning rate (Figure 6). Resetting the optimizer changes little (0.81, 0.91 and 0.97 at 2.5, 5 and 10%). The window thus tracks the amount of early learning more closely than the number of steps, and the assignment it leaves is held in the parameters rather than in the optimizer state.

Public suites. The public models also settle early (Figure 3b). The same-initialization agreement of the DataDecide 1B models reaches its final level by 3.6% of training, Pythia’s previous-token heads stabilize by about 2%, and Pythia-12B’s two models already share their layout at 2.1% (Table 1b). Before training, none of the statistics of a head’s initial weights or initial gradient that we tested predicts which head will take a role, so the assignment is not readable at step 0; it is set by the first steps of training on top of the initialization.

5 What Decides How Much Is Kept

The agreement between same-initialization models varies: across corpus pairs, across sizes and across training setups. Two quantities track this variation.

Corpus similarity. Across DataDecide’s corpus pairs, the Jensen–Shannon distance between two corpora’s unigram distributions predicts how closely an initialization’s layout recurs on both.

The relation holds at every size from 60M on and strengthens with size (Spearman -0.20 to -0.33 at 60M, -0.65 to -0.83 from 300M on, -0.76 to -0.80 across the 300 corpus pairs at 1B; Mantel $p < 0.001$ from 90M; Table 6), on pairs that share no data source (-0.36 to -0.82), and after partialling out source overlap. Bigram distance predicts less (-0.55 to -0.62 at 1B). Splitting the vocabulary into function words (Yang et al., 2026) and content words shows which part of the distribution matters (Figure 4). Content-word distance predicts agreement at every size (ρ from -0.23 to -0.69), also when function-word distance is controlled (partial ρ -0.25 to -0.78), whereas function-word distance, once content words are controlled, no longer predicts lower agreement (its partial ρ turns positive, ≥ 0.09 , as two collinear distances do); the English web corpora in DataDecide use function words alike. Controlled pretraining extends the relation to more distant corpora: C4 and scientific papers share the layout (0.19 ± 0.06 over nine initializations), pairs involving source code share none (≈ 0), and Pythia’s two versions of the Pile, nearly identical in content, share it most.

Gradient noise. Which head takes a role is settled when the symmetry between heads breaks early in training, by the initialization together with the gradient noise of those steps. The noise of stochastic gradients scales with the learning rate per example in a batch (Smith and Le, 2018; Jastrzębski et al., 2018), and lowering it lets the initialization decide more. At fixed model size, runs that differ only in the order of their batches agree




Figure 3: **The assignment is fixed early.** (a) Controlled pretraining (4 layers $\times$ 8 heads, 10k steps): agreement of the final layout with the uninterrupted run after switching the corpus to code or adding Gaussian noise to every weight at a given step; hollow circles: a 12-layer model. (b) Agreement of the layout with its own final layout during training, in public suites and controlled runs; star: Pythia-12B at 2.1%. Shaded: 1–2.5% of training.

0.43 with 16-sequence batches and 0.90 with 512-sequence batches, and same-initialization agreement between C4 and scientific papers rises from 0.01 to 0.22 (Appendix H); tripling the learning rate acts like quartering the batch. This can account for the growth with size in DataDecide: its recipes lower the learning rate per batch token 140-fold from 4M to 1B, whereas a 30 $\times$ larger controlled model at a fixed rate keeps no more of its initialization, and Pythia, trained with large batches at every size, keeps it at 70M.

6 Behaviour Follows the Corpus

Strength, timing and performance follow the corpus. How strongly a circuit is expressed follows the corpus (significant corpus effect on 4 of 6 strength measures, initialization component ≈ 0), and so does when it forms: in controlled pretraining with six initializations and three corpora, the corpus explains 81% of the variance in the onset of induction heads and the initialization none ($p = 0.0001$ vs. 0.47). Across 10 benchmarks and their average at 14 sizes, the initialization has a significant main effect in 3.9% (accuracy) and 7.8%




Figure 4: **Similar content words, similar heads.** (a) 1B pairs of DataDecide corpora that share no data source: Jensen–Shannon distance between their content-word distributions vs. the which-head agreement of same-initialization models trained on them. (b) Negative partial Spearman correlation between distance and agreement for content words (controlling for function words) and for function words (controlling for content words), by size; bars span the three roles; †: 1B with five seeds at 10% of training.

Property	Follows	Key statistic
Which head takes a role	initialization	0.18–0.43 vs. 0.00
Residual coordinates	initialization	0.20 vs. 0.00
Layer that hosts a role	all models	0.4–0.7, all pairs
Copying algorithm	all models	75 / 75 models
Circuit strength	corpus	4 / 6 measures
Onset of induction heads	corpus	81% vs. 0%
Representations	corpus	CKA, best match
10 benchmarks, 14 sizes	corpus	init. sig. 4–8%
Knowledge conflicts	corpus	init. sig. 0 / 12
Question: habit	corpus	+2.4 nats

Table 3: **Which properties follow the initialization and which the corpus.** Key statistic: same- vs. different-initialization agreement; variance share of corpus vs. initialization; or share of cells with a significant initialization effect (5% is the false-positive rate).

(per-byte likelihood) of the 154 cells, about the 5% expected by chance, and the corpus in 75% and 98% (median variance components: initialization 0.000, corpus 0.49–0.83; Table 9); behaviour under knowledge conflict shows no initialization effect in any of 12 conditions. What the initialization does contribute is specific to the corpus it is trained on (with one run per cell, interaction and run-to-run noise are not separable): an initialization that does well on one corpus is no more likely to do well on another (Table 3).

6.1 Case Study: A Data Component and the Question: Template

The previous results concern where roles are placed and show that performance does not follow the




Figure 5: **Flan changes behaviour, not the head layout** (1B). (a) Flan minus no-Flan models: added trust in a counterfactual context relative to a declarative prompt, by template ($\pm$ SE over initializations). (b) Which-head agreement (mean over nine roles) of same-initialization pairs with and without Flan, of same-initialization pairs of other Dolma 1.7 ablations, and of different-initialization pairs. (c) The ten strongest retrieval heads of each Flan model: their share of all heads, of the attention to the context entity, and of the attention that Flan adds after Question: (relative to Q:).

placement. This case study asks the converse question for one data component that clearly changes behaviour: does it move the head layout, and where is the behaviour it teaches expressed? We compare seed-matched models trained with and without the component. Dolma 1.7 contains about 1–2% of Flan instruction data (Longpre et al., 2023; Soldaini et al., 2024), and DataDecide trains a Dolma 1.7 recipe with Flan removed from the same initializations. We compare the two on knowledge-conflict prompts from ParaConflict (Li et al., 2025): one sentence asserts a counterfactual fact (“The capital of France is Rome.”) and is followed by a declarative continuation or by a question in Flan’s format (Question: What is the capital of France? Answer:).

Flan adds trust after Question:. Removing Flan lowers the 1B models’ trust in the counterfactual context by 1.3–3.5 nats across four kinds of context when the prompt is a question, and leaves declarative prompts and the models’ knowledge of the facts unchanged. The share of questions on which a model ignores a one-sentence context and answers from memory rises from 1.9% to 14.3%. The effect is tied to Flan’s own wording (Figure 5a,

Ending after the context	count in Dolma 1.7	Flan effect	SE
Question: q Answer:	11.3M / 9.6M	+2.43	0.79
Question: q	11.3M	+2.04	0.83
Answer:	9.6M	+1.13	0.63
Q: q A:	38.4M / 6.9M	+0.39	0.89
Query: q Response:	0.15M / 0.75M	+0.52	0.83
Question: q' Answer:	as first row	+3.06	1.42
declarative (no question)	–	+0.20	0.50

Table 4: **The habit is keyed to Flan’s literal template.** Flan minus no-Flan models (1B, three initializations each): added trust in the counterfactual context relative to a declarative ending, in nats. Example: *The capital of France is Rome.* Question: *What is the capital of France?* Answer: *The capital of France is ...*; q' asks about another entity. Marker counts: occurrences in the pretraining corpus. The last row is the raw trust difference for the declarative ending.

Table 4): at 1B Question: alone carries most of it, and across the seven sizes Flan’s template exceeds the equivalent Q: ... A: by 0.93 ± 0.29 nats and Query: ... Response: by 0.98 ± 0.32 , although Q: is three times as frequent as Question: in Dolma 1.7; a question about a different entity triggers it as strongly. It is present from 3.6% of pre-training, at sizes from 60M on, on PopQA counterfactuals (Mallen et al., 2023) and NQ-Swap (Longpre et al., 2021), and in OLMo 2 (Team OLMo et al., 2024) after mid-training on a mixture with FLAN (Appendix J).

The layout is preserved. Seed-matched models with and without Flan agree on which head takes each role at 0.45 (mean over nine roles), as closely as seed-matched pairs of other Dolma 1.7 ablations (0.40), while different initializations agree at -0.03 (Figure 5b). Removing Flan changes the behaviour without lowering the agreement that the shared initialization produces.

Where the habit acts. To locate the habit, we measure each head’s attention from the last prompt token to the counterfactual entity, the read-out of the retrieval role, and take the difference between Question: and Q: prompts. Models trained with Flan attend more to the context after Question: (summed over heads, $+0.54$), models without Flan less (-0.21). The added attention is spread over the heads that already attend to the entity: the model’s ten strongest retrieval heads carry 30% of its attention to the entity and receive 29% of the addition (30% when the heads are chosen on held-out items;

Figure 5c), and models with the same initialization agree on where it falls (0.11, against -0.03 for different initializations). Flan changes how much the existing retrieval heads attend after its template, not which heads retrieve.

Consequences for evaluation. Context-faithfulness and knowledge-conflict benchmarks commonly pose questions in exactly this format, and every public base model we tested trusts a counterfactual context more after Question: than in a declarative prompt, by 1.0–4.6 nats, most of all the model pretrained with Flan (Table 11). A score on such benchmarks therefore mixes faithfulness with a habit that a small slice of pretraining data strengthens and a change of template removes, which gives prompt-template sensitivity (Sclar et al., 2024; Hua et al., 2025; Liu and Chu, 2026) a concrete source in the pretraining corpus.

7 Related Work

Seeds and training variability. Seed variation is a known source of irreproducibility in fine-tuning and pretraining (Reimers and Gurevych, 2017; Dodge et al., 2020; McCoy et al., 2020; Sellam et al., 2022), modelled as noise (Benito-Santos et al., 2025) or studied as convergence across training (Fehlauer et al., 2025; van der Wal et al., 2025); SeedPrints (Tong et al., 2026) identifies a seed from output biases. Crossing seeds with corpora separates the initialization from the data: its effect on head placement carries across corpora, its effect on behaviour does not.

Universality and stability of circuits. Hidden units of independently trained networks correspond only up to permutation (Li et al., 2016; Entezari et al., 2022; Ainsworth et al., 2023), and models fine-tuned from one pretrained checkpoint often stay in one basin (Neyshabur et al., 2020), though not always (Juneja et al., 2023). Algorithms recur across models while components vary (Chughtai et al., 2023; Gurnee et al., 2024; Tigges et al., 2024); attention heads differ across retraining (Bali et al., 2026), specialize along developmental trajectories (Wang et al., 2025), and can be read from their parameters (Elhelo and Geva, 2025). Without any alignment, a shared initialization makes heads correspond across corpora, although the models are not linearly connected.

Data and the formation of mechanisms. Data properties decide whether in-context learning

emerges (Chan et al., 2022; Singh et al., 2024), and individual training examples shape when a head forms (Chen et al., 2026). We find that the corpus sets whether, when and how strongly a mechanism forms, while the initialization sets where. Early training is known to be decisive for what a network can still learn (Achille et al., 2019) and for which basin it ends in (Frankle et al., 2020); we show the same for which head takes which role.

Knowledge conflicts and prompt format.

Whether a model follows a counterfactual context or its memory depends on the context’s form and the model’s confidence (Longpre et al., 2021; Xie et al., 2024; Du et al., 2024), is mediated by competing mechanisms (Ortu et al., 2024), and shifts with the prompt’s wording (Zhou et al., 2023; Sclar et al., 2024; Hua et al., 2025; Liu and Chu, 2026); synthetic pretraining shows which data properties make models use both sources (Kim et al., 2026). We trace one such format effect to a 1–2% slice of natural pretraining data.

8 Discussion and Conclusion

Interpretability. A circuit found in one model generalizes to other models in its algorithm and its layers; its head labels transfer only partly, and only to models with the same initialization. Ablating the ten heads that matter most for induction in a same-initialization sibling trained on another corpus does 19.5% of the damage of ablating the model’s own, against 10% for random heads in the same layers and 2.0% for heads found in a different-initialization model trained on the *same* corpus (1B, six corpora). Comparisons of circuits across training data should therefore match initializations, and seed-matched siblings in public suites provide such comparisons: their heads correspond, so differences between their circuits can be attributed to the data.

Training and evaluation. Larger batches or smaller learning rates make head assignments more reproducible across corpora; seed selection does not carry over between recipes.

Conclusion. Models trained from one initialization on different corpora agree on which heads take each role; the agreement is set by a small amount of early learning and grows with corpus similarity and lower gradient noise, while what the heads do follows the corpus.

Limitations

We study decoder-only transformers with multi-head attention, pretrained on English corpora, up to 12B parameters, and define head roles through attention patterns and OV circuits; with grouped-query attention, heads are interchangeable only within and across groups, which changes the baseline. Multilingual seed-crossed suites would test the corpus-similarity relation where function words differ across corpora, and other architectures can be added as such suites appear. Our interventions during training use models of up to 12 layers; intervening on larger models in their first percent of training is the natural next step. The template habit is one data component; why models key on Question: rather than an equivalent marker, and which other instruction-data families install habits of their own, are open. Which property of the initialization decides which head will take a role is likewise open.

Ethical Considerations

We analyse publicly released models and datasets under their licences and train small models on public corpora; no human subjects are involved. We report our audit of public checkpoints to support their correct use.

References

- Alessandro Achille, Matteo Rovere, and Stefano Soatto. 2019. [Critical learning periods in deep neural networks](#). In *International Conference on Learning Representations*.
- Samuel K. Ainsworth, Jonathan Hayase, and Siddhartha Srinivasa. 2023. [Git re-basin: Merging models modulo permutation symmetries](#). In *International Conference on Learning Representations*.
- Karan Bali, Jack Stanley, Praneet Suresh, and Danilo Bzdok. 2026. [Quantifying LLM attention-head stability: Implications for circuit universality](#). arXiv preprint arXiv:2602.16740.
- Alejandro Benito-Santos, Adrian Ghajari, and Víctor Fresno. 2025. [Robust estimation of population-level effects in repeated-measures NLP experimental designs](#). In *Proceedings of the 63rd Annual Meeting of the Association for Computational Linguistics (Volume 1: Long Papers)*, pages 33076–33089, Vienna, Austria. Association for Computational Linguistics.
- Stella Biderman, Hailey Schoelkopf, Quentin Anthony, Herbie Bradley, Kyle O’Brien, Eric Hallahan, Mohammad Aflah Khan, Shivanshu Purohit, USVSN Sai

- Prashanth, Edward Raff, Aviya Skowron, Lintang Sutawika, and Oskar van der Wal. 2023. [Pythia: A suite for analyzing large language models across training and scaling](#). In *Proceedings of the 40th International Conference on Machine Learning*.
- Stephanie C. Y. Chan, Adam Santoro, Andrew K. Lampinen, Jane X. Wang, Aaditya Singh, Pierre H. Richemond, Jay McClelland, and Felix Hill. 2022. [Data distributional properties drive emergent in-context learning in transformers](#). In *Advances in Neural Information Processing Systems*.
- Jianhui Chen, Yuzhang Luo, and Liangming Pan. 2026. [Mechanistic data attribution: Tracing the training origins of interpretable LLM units](#). In *Proceedings of the 43rd International Conference on Machine Learning*.
- Bilal Chughtai, Lawrence Chan, and Neel Nanda. 2023. [A toy model of universality: Reverse engineering how networks learn group operations](#). In *Proceedings of the 40th International Conference on Machine Learning*.
- Jesse Dodge, Gabriel Ilharco, Roy Schwartz, Ali Farhadi, Hannaneh Hajishirzi, and Noah Smith. 2020. [Fine-tuning pretrained language models: Weight initializations, data orders, and early stopping](#). arXiv preprint arXiv:2002.06305.
- Kevin Du, Vésteinn Snæbjarnarson, Niklas Stoehr, Jennifer White, Aaron Schein, and Ryan Cotterell. 2024. [Context versus prior knowledge in language models](#). In *Proceedings of the 62nd Annual Meeting of the Association for Computational Linguistics (Volume 1: Long Papers)*, pages 13211–13235, Bangkok, Thailand. Association for Computational Linguistics.
- Nelson Elhage, Neel Nanda, Catherine Olsson, Tom Henighan, Nicholas Joseph, Ben Mann, Amanda Askell, Yuntao Bai, Anna Chen, Tom Conerly, Nova DasSarma, Dawn Drain, Deep Ganguli, Zac Hatfield-Dodds, Danny Hernandez, Andy Jones, Jackson Kernion, Liane Lovitt, Kamal Ndousse, and 6 others. 2021. [A mathematical framework for transformer circuits](#). Transformer Circuits Thread.
- Amit Elhelo and Mor Geva. 2025. [Inferring functionality of attention heads from their parameters](#). In *Proceedings of the 63rd Annual Meeting of the Association for Computational Linguistics (Volume 1: Long Papers)*, pages 17701–17733, Vienna, Austria. Association for Computational Linguistics.
- Rahim Entezari, Hanie Sedghi, Olga Saukh, and Behnam Neyshabur. 2022. [The role of permutation invariance in linear mode connectivity of neural networks](#). In *International Conference on Learning Representations*.
- Finlay Fehlaue, Kyle Mahowald, and Tiago Pimentel. 2025. [Convergence and divergence of language models under different random seeds](#). In *Proceedings of the 2025 Conference on Empirical Methods in Natural Language Processing*, pages 32982–32991,

653	Suzhou, China. Association for Computational Lin-	Shayne Longpre, Le Hou, Tu Vu, Albert Webson,	709
654	guistics.	Hyung Won Chung, Yi Tay, Denny Zhou, Quoc V. Le,	710
655	Jonathan Frankle, Gintare Karolina Dziugaite, Daniel M.	Barret Zoph, Jason Wei, and Adam Roberts. 2023.	711
656	Roy, and Michael Carbin. 2020. Linear mode connec-	The Flan collection: Designing data and methods for	712
657	tivity and the lottery ticket hypothesis . In <i>Proceed-</i>	effective instruction tuning . In <i>Proceedings of the</i>	713
658	<i>ings of the 37th International Conference on Machine</i>	<i>40th International Conference on Machine Learning</i> .	714
659	<i>Learning</i> .		
660	Wes Gurnee, Theo Horsley, Zifan Carl Guo, Tara Rezaei	Shayne Longpre, Kartik Perisetla, Anthony Chen,	715
661	Kheirkhah, Qinyi Sun, Will Hathaway, Neel Nanda,	Nikhil Ramesh, Chris DuBois, and Sameer Singh.	716
662	and Dimitris Bertsimas. 2024. Universal neurons in	2021. Entity-based knowledge conflicts in question	717
663	GPT2 language models . <i>Transactions on Machine</i>	answering . In <i>Proceedings of the 2021 Conference</i>	718
664	<i>Learning Research</i> .	<i>on Empirical Methods in Natural Language Process-</i>	719
665	Andong Hua, Kenan Tang, Chenhe Gu, Jindong Gu,	<i>ing</i> , pages 7052–7063, Online and Punta Cana, Do-	720
666	Eric Wong, and Yao Qin. 2025. Flaw or artifact?	minican Republic. Association for Computational	721
667	rethinking prompt sensitivity in evaluating LLMs .	Linguistics.	722
668	In <i>Proceedings of the 2025 Conference on Empiri-</i>	Ian Magnusson, Nguyen Tai, Ben Bogin, David Heine-	723
669	<i>cal Methods in Natural Language Processing</i> , pages	man, Jena D. Hwang, Luca Soldaini, Akshita Bha-	724
670	19889–19899, Suzhou, China. Association for Com-	gia, Jiacheng Liu, Dirk Groeneveld, Oyvind Tafjord,	725
671	putational Linguistics.	Noah A. Smith, Pang Wei Koh, and Jesse Dodge.	726
672	Stanisław Jastrzębski, Zachary Kenton, Devansh Arpit,	2025. DataDecide: How to predict best pretraining	727
673	Nicolas Ballas, Asja Fischer, Yoshua Bengio, and	data with small experiments . In <i>Proceedings of the</i>	728
674	Amos Storkey. 2018. Three factors influencing min-	<i>42nd International Conference on Machine Learning</i> .	729
675	ima in SGD . In <i>Artificial Neural Networks and Ma-</i>	Alex Mallen, Akari Asai, Victor Zhong, Rajarshi Das,	730
676	<i>chine Learning – ICANN 2018</i> .	Daniel Khoshnab, and Hannaneh Hajishirzi. 2023.	731
677	Jeevesh Juneja, Rachit Bansal, Kyunghyun Cho, João	When not to trust language models: Investigating	732
678	Sedoc, and Naomi Saphra. 2023. Linear connectivity	effectiveness of parametric and non-parametric mem-	733
679	reveals generalization strategies . In <i>International</i>	ories . In <i>Proceedings of the 61st Annual Meeting of</i>	734
680	<i>Conference on Learning Representations</i> .	<i>the Association for Computational Linguistics (Vol-</i>	735
681	Minsung Kim, Dong-Kyum Kim, Jea Kwon, Nakyeong	<i>ume 1: Long Papers)</i> , pages 9802–9822, Toronto,	736
682	Yang, Kyomin Jung, and Meeyoung Cha. 2026. How	Canada. Association for Computational Linguistics.	737
683	training data shapes the use of parametric and in-	Nathan Mantel. 1967. The detection of disease cluster-	738
684	context knowledge in language models . In <i>Proceed-</i>	<i>ing and a generalized regression approach</i> . <i>Cancer</i>	739
685	<i>ings of the 64th Annual Meeting of the Association for</i>	<i>Research</i> , 27(2):209–220.	740
686	<i>Computational Linguistics (Volume 1: Long Papers)</i> ,	R. Thomas McCoy, Junghyun Min, and Tal Linzen.	741
687	pages 23242–23257, San Diego, California, United	2020. BERTs of a feather do not generalize together:	742
688	States. Association for Computational Linguistics.	Large variability in generalization across models with	743
689	Simon Kornblith, Mohammad Norouzi, Honglak Lee,	similar test set performance . In <i>Proceedings of the</i>	744
690	and Geoffrey Hinton. 2019. Similarity of neural	<i>Third BlackboxNLP Workshop on Analyzing and In-</i>	745
691	network representations revisited . In <i>Proceedings</i>	<i>terpreting Neural Networks for NLP</i> , pages 217–227,	746
692	<i>of the 36th International Conference on Machine</i>	Online. Association for Computational Linguistics.	747
693	<i>Learning</i> .	Behnam Neyshabur, Hanie Sedghi, and Chiyuan Zhang.	748
694	Gaotang Li, Yuzhong Chen, and Hanghang Tong. 2025.	2020. What is being transferred in transfer learn-	749
695	Taming knowledge conflicts in language models . In	ing? In <i>Advances in Neural Information Processing</i>	750
696	<i>Proceedings of the 42nd International Conference on</i>	<i>Systems</i> .	751
697	<i>Machine Learning</i> .	Catherine Olsson, Nelson Elhage, Neel Nanda, Nicholas	752
698	Yixuan Li, Jason Yosinski, Jeff Clune, Hod Lipson, and	Joseph, Nova DasSarma, Tom Henighan, Ben Mann,	753
699	John Hopcroft. 2016. Convergent learning: Do differ-	Amanda Askell, Yuntao Bai, Anna Chen, Tom Con-	754
700	ent neural networks learn the same representations?	erly, Dawn Drain, Deep Ganguli, Zac Hatfield-Dodds,	755
701	In <i>International Conference on Learning Representa-</i>	Danny Hernandez, Scott Johnston, Andy Jones, Jack-	756
702	<i>tions</i> .	son Kernion, Liane Lovitt, and 7 others. 2022. In-	757
703	Yang Liu and Chenhui Chu. 2026. Understanding the	context learning and induction heads . <i>Transformer</i>	758
704	prompt sensitivity . In <i>Proceedings of the 64th An-</i>	<i>Circuits Thread</i> .	759
705	<i>nual Meeting of the Association for Computational</i>	Francesco Ortu, Zhijing Jin, Diego Doimo, Mrinmaya	760
706	<i>Linguistics (Volume 1: Long Papers)</i> , pages 44362–	Sachan, Alberto Cazzaniga, and Bernhard Schölkopf.	761
707	44388, San Diego, California, United States. Associ-	2024. Competition of mechanisms: Tracing how	762
708	ation for Computational Linguistics.	language models handle facts and counterfactuals .	763
		In <i>Proceedings of the 62nd Annual Meeting of the</i>	764
		<i>Association for Computational Linguistics (Volume 1:</i>	765

Long Papers), pages 8420–8436, Bangkok, Thailand. Association for Computational Linguistics.

Nils Reimers and Iryna Gurevych. 2017. [Reporting score distributions makes a difference: Performance study of LSTM-networks for sequence tagging](#). In *Proceedings of the 2017 Conference on Empirical Methods in Natural Language Processing*, pages 338–348, Copenhagen, Denmark. Association for Computational Linguistics.

Melanie Sclar, Yejin Choi, Yulia Tsvetkov, and Alane Suhr. 2024. [Quantifying language models’ sensitivity to spurious features in prompt design or: How i learned to start worrying about prompt formatting](#). In *International Conference on Learning Representations*.

Thibault Sellam, Steve Yadlowsky, Jason Wei, Naomi Saphra, Alexander D’Amour, Tal Linzen, Jasmijn Bastings, Iulia Turc, Jacob Eisenstein, Dipanjan Das, Ian Tenney, and Ellie Pavlick. 2022. [The MultiBERTs: BERT reproductions for robustness analysis](#). In *International Conference on Learning Representations*.

Aaditya K. Singh, Ted Moskovitz, Felix Hill, Stephanie C. Y. Chan, and Andrew M. Saxe. 2024. [What needs to go right for an induction head? a mechanistic study of in-context learning circuits and their formation](#). In *Proceedings of the 41st International Conference on Machine Learning*.

Samuel L. Smith and Quoc V. Le. 2018. [A bayesian perspective on generalization and stochastic gradient descent](#). In *International Conference on Learning Representations*.

Luca Soldaini, Rodney Kinney, Akshita Bhagia, Dustin Schwenk, David Atkinson, Russell Authur, Ben Bogin, Khyathi Chandu, Jennifer Dumas, Yanai Elazar, Valentin Hofmann, Ananya Jha, Sachin Kumar, Li Lucy, Xinxu Lyu, Nathan Lambert, Ian Magnusson, Jacob Morrison, Niklas Muennighoff, and 17 others. 2024. [Dolma: an open corpus of three trillion tokens for language model pretraining research](#). In *Proceedings of the 62nd Annual Meeting of the Association for Computational Linguistics (Volume 1: Long Papers)*, pages 15725–15788, Bangkok, Thailand. Association for Computational Linguistics.

Team OLMo, Pete Walsh, Luca Soldaini, Dirk Groeneveld, Kyle Lo, Shane Arora, Akshita Bhagia, Yuling Gu, Shengyi Huang, Matt Jordan, Nathan Lambert, Dustin Schwenk, Oyvind Tafjord, Taira Anderson, David Atkinson, Faeze Brahman, Christopher Clark, Pradeep Dasigi, Nouha Dziri, and 24 others. 2024. [2 OLMo 2 furious](#). arXiv preprint arXiv:2501.00656.

Curt Tigges, Michael Hanna, Qinan Yu, and Stella Biderman. 2024. [LLM circuit analyses are consistent across training and scale](#). In *Advances in Neural Information Processing Systems*.

Yao Tong, Haonan Wang, Siquan Li, Kenji Kawaguchi, and Tianyang Hu. 2026. [SeedPrints: Fingerprints](#)

[can even tell which seed your large language model was trained from](#). In *International Conference on Learning Representations*.

Oskar van der Wal, Pietro Lesci, Max Muller-Eberstein, Naomi Saphra, Hailey Schoelkopf, Willem Zuidema, and Stella Biderman. 2025. [PolyPythias: Stability and outliers across fifty language model pre-training runs](#). In *International Conference on Learning Representations*.

George Wang, Jesse Hoogland, Stan van Wingerden, Zach Furman, and Daniel Murfet. 2025. [Differentiation and specialization of attention heads via the refined local learning coefficient](#). In *International Conference on Learning Representations*.

Wenhao Wu, Yizhong Wang, Guangxuan Xiao, Hao Peng, and Yao Fu. 2025. [Retrieval head mechanistically explains long-context factuality](#). In *International Conference on Learning Representations*.

Guangxuan Xiao, Yuandong Tian, Beidi Chen, Song Han, and Mike Lewis. 2024. [Efficient streaming language models with attention sinks](#). In *International Conference on Learning Representations*.

Jian Xie, Kai Zhang, Jiangjie Chen, Renze Lou, and Yu Su. 2024. [Adaptive chameleon or stubborn sloth: Revealing the behavior of large language models in knowledge conflicts](#). In *International Conference on Learning Representations*.

Xiulin Yang, Heidi R. Getz, and Ethan Gotlieb Wilcox. 2026. [Function words as statistical cues for language learning](#). In *Proceedings of the 64th Annual Meeting of the Association for Computational Linguistics (Volume 1: Long Papers)*, pages 16042–16058, San Diego, California, United States. Association for Computational Linguistics.

Wenxuan Zhou, Sheng Zhang, Hoifung Poon, and Muhao Chen. 2023. [Context-faithful prompting for large language models](#). In *Findings of the Association for Computational Linguistics: EMNLP 2023*, pages 14544–14556, Singapore. Association for Computational Linguistics.

A Auditing the Public Suites

DataDecide 1B recipe	seed	step 0 vs. seed’s step 0	step 2500 vs. seed’s step 0	step 2500 vs. partner
DCLM QC 7% FW2	1	0.00	0.00	0.11
DCLM QC 7% FW3	1	0.00	0.00	0.11
Dolma 1.7 no math/code	1	0.00	0.00	0.08
Dolma 1.7 no Reddit	1	0.00	0.00	0.08
Falcon+CC QC orig. 10%	3	0.00	0.00	0.09
Falcon+CC QC Tulu 10%	3	0.00	0.00	0.09
genuine siblings (C4 vs. Dolma 1.7, same seed), step 2500				0.07–0.08

Table 5: **Six 1B runs were not initialized from their seed’s weights.** Correlation of one attention weight matrix (block 5) between checkpoints. The other 69 runs share their seed’s step-0 weights exactly.

Step-0 identity. Pythia’s standard and deduplicated models have tensor-identical step-0 weights at 70M, 160M, 410M, 1.4B and 2.8B, rounding-level identical weights at 1B (correlation 0.99993) and identical files at 6.9B and 12B. For DataDecide we compare one weight matrix of every released step-0 checkpoint with the step 0 of the same seed’s C4 run, reading the tensor directly from the hosted files.

All sizes. Below 1B, we use for each size the recipes whose three seeds share their published step-0 weights (by file hash) and the checkpoint listed in Table 6: the final checkpoint up to 90M, and at larger sizes a checkpoint available for all selected runs (43–98% of training; 11% for the five-seed 1B set). For every one of these 686 runs we compare the earliest training checkpoint (step 1,250; 2,500 at 1B) with the step-0 weights of every seed. Each correlates with its own seed’s step 0 (0.05–0.41, rising with size) and with no other seed’s (≤ 0.009), except one 1B run at 11% of training (FineWeb-Pro, seed 3), which matches no seed and is excluded.

Training start. A published step-0 checkpoint need not be the state a run started from, so we compare the earliest training checkpoint with every candidate initialization. The DataDecide 1B runs of C4, Dolma 1.7 and DCLM correlate at step 2,500 only with their own seed’s step-0 weights (0.23 vs. 0.000). PolyPythias’ “weight-seed” variants at 160M, in contrast, were trained from the standard initialization: their step-1 weights equal the standard model’s (correlation 1.0) and are unrelated to their own published step 0; we label them accordingly.

Six 1B runs with an unlisted initialization. Six 1B DataDecide runs disagree with all their nominal siblings in Figure 1a: the DCLM-QC 7% FineWeb-2/3 recipes and the Dolma 1.7 no-math-code / no-Reddit ablations of seed 1, and the Falcon-and-CC QC original / Tulu recipes of seed 3 (Table 5). Their published step-0 weights differ from their seed’s, their step-2,500 weights correlate with no seed’s step 0, and each pair’s step-2,500 weights correlate with each other as strongly as those of two genuine siblings: each pair was trained from an initialization of its own. All other 69 runs share their seed’s step-0 weights exactly, so the layout flagged exactly the six runs that differ. We exclude them from every same- and different-initialization

comparison at 1B.

A near-duplicate Pythia pair. The released final checkpoint of pythia-2.8b-deduped is nearly identical to that of pythia-2.8b (attention-weight correlation 0.993, embedding correlation 0.985), whereas models trained from one initialization on different data retain a correlation of only 0.03–0.09 with it; we exclude the 2.8B pair.

B Head Roles

Induction: attention from the second copy of a random 128-token block to the token after its first occurrence (200 blocks). **Previous-token:** attention to position $t - 1$ on 50 natural texts (split-half reliability 0.99–1.00). **Attention sink:** attention to position 0 on the same texts. **Retrieval:** attention from the last prompt token to the counterfactual entity in the knowledge-conflict prompts. **Duplicate-token, current-token, two-back and delimiter:** attention to an earlier copy of the current token, to itself, to $t - 2$ and to punctuation, on natural text. **OV copying:** the share of positive eigenvalue mass of the head’s OV circuit composed with the embedding and unembedding (Elhage et al., 2021), computed from weights alone. An induction map is used only when the model has an induction head (maximum score > 0.3). Split-half reliabilities are 0.99–1.00 for the previous-token, sink, current-token, two-back and delimiter maps, so differences in agreement between roles are not differences in measurement noise.

Agreement on the strongest head. Which-head agreement averages a rank correlation over all layers. Restricted to the three layers where a role is strongest at 1B, it is as high or higher (0.19–0.47 for seed-matched pairs, -0.07 to 0.03 for other pairs) than in the remaining layers (0.18–0.43). In those three layers, the strongest head of two seed-matched models is the same in 12–30% of cases across the nine roles (different initializations: 4–9%; chance: 6%), and their three strongest heads overlap by 27–48% (chance: 19%). The example of Section 1 is the most consistent of the three initializations: under the other two, the most frequent strongest induction head wins on 11 of 21 and 9 of 23 corpora with an induction head.

Ablation-defined roles. Defining a head’s role by the effect of ablating it rather than by its attention pattern gives the same picture, more weakly: maps of each head’s contribution to induction and

to context use agree at 0.10–0.12 between seed-matched 1B models trained on different corpora and at 0.00 between models with different initializations (six corpora, three initializations).

C Head Permutation Symmetry and the Chance Baseline

Lemma 1. *Let a layer have H heads and let π permute them together with their query, key, value and output blocks. Assume (i) the initialization distribution is invariant under π ; (ii) training is π -equivariant, i.e. the update applied to $\pi\theta$ is π applied to the update of θ (true for gradient descent, Adam/AdamW, weight decay and global-norm clipping, as the loss is π -invariant); and (iii) the data and their order are independent of the initialization. Let $w(\theta)$ be the head with the largest score for a role, computed from the network’s internals, so that $w(\pi\theta) = \pi(w(\theta))$. Then for every corpus D , $P(w(\theta_T) = h \mid D) = 1/H$, and two independently initialized models trained on any corpora D, D' agree on the winning head with probability $1/H$; the expected within-layer rank correlation of their head scores is 0.*

Proof sketch. By (i)–(ii), θ_T and $\pi\theta_T$ have the same distribution given D , so $w(\theta_T)$ is uniform over heads. Independent initializations make w and w' independent, so they agree with probability $\sum_h (1/H)^2 = 1/H$; averaging over permutations sends any rank correlation to zero in expectation. $\square$

The lemma fixes the baseline for pairs with independent initializations; it makes no prediction for pairs that share one, whose agreement is the quantity we measure. Different-initialization pairs agree at -0.02 to 0.00 for all nine roles (Table 1a).

D Results at Every Size

Table 6 lists, for every DataDecide size, the recipes and seeds used, same- and different-initialization agreement, initialization identification at the end of training and in the earliest checkpoints, the corpus-distance relation and the recipe’s learning rate per batch token. Agreement is weak up to 20M parameters, where most models have not yet formed an induction head, and strong from 150M on.

E Identifying the Initialization

Table 7 compares identification from the head layout with two baselines on the same leave-one-

corpus-out task. Final weights identify the initialization at every size: same-initialization pairs keep a weight correlation of 0.01–0.04, other pairs 0.000. SeedPrints (Tong et al., 2026), a rank test on final hidden states for random token sequences, identifies it from 60M on; its statistic is also raised for pairs that share only a corpus, which the head layout and the weights are not. On the six 1B runs with an unlisted initialization, all three methods agree: each run is closest to its partner run and to none of its nominal siblings.

F Controlled Pretraining

Llama-style models (RMSNorm, SwiGLU, RoPE, untied embeddings, OLMo tokenizer): 4 layers $\times$ 8 heads, $d = 256$; and 12 layers $\times$ 12 heads, $d = 768$. AdamW (learning rate 10^{-3} , $\beta = (0.9, 0.95)$, weight decay 0.1), 300 warm-up steps, cosine decay to 10%, batches of 64×512 tokens, 10,000 steps. Corpora: the first 400M tokens of single DataDecide sources for web text (C4), code (StarCoder), scientific papers (peS2o) and books (Gutenberg). Branches continue from the parent’s saved model and optimizer state at step $k \in \{100, 250, 500, 1000, 2000, 4000, 8000\}$, either switching the corpus or adding $\epsilon \cdot \text{std}(W) \cdot \mathcal{N}(0, 1)$ to every tensor (Table 2). At initialization the basin is finite: perturbations up to 1% of the weight scale leave the layout unchanged (0.91–0.99), 10% partly reassign it (0.74; 0.57 and 0.91 for the two initializations), and 100% reassign it fully (0.03).

G Warm-up and Optimizer State

We repeat the noise branches of Appendix F with two parents whose learning rate warms up over 1,000 instead of 300 steps, branching at $k \in \{100, 250, 500, 1000, 2000\}$, and with branches that discard the parent’s AdamW state at $k \in \{250, 500, 1000\}$. A parent continued from its own step-2,000 state reproduces its final layout (0.995–1.00), so the comparisons are not limited by run-to-run noise. With the longer warm-up, a switch to code keeps 0.41 of the layout at 2.5% and 0.69 at 10% (0.60 and 0.71 with the default warm-up). Changing only the warm-up, with the same initialization and data order, changes part of the final layout (agreement 0.21–0.59 for induction and 0.62–0.67 for previous-token maps), another sign that the first steps of training help decide which head takes a role.

Parameters	Checkpoint (share)	Recipes	Seeds	Same init., other corpus	Other init., same corpus	Init. ID (%)	Init. ID, early (%)	ρ (corpus distance, agreement)	LR per batch token (10^{-9})
4M	5,725 (100%)	18	3	0.03	0.00	56	—	—	213.6
6M	9,182 (100%)	23	3	0.05	−0.01	77	—	—	183.1
8M	13,039 (100%)	17	3	0.02	0.00	65	—	—	167.8
10M	15,117 (100%)	16	3	0.04	0.00	81	72	—	152.6
14M	21,953 (100%)	18	3	0.08	−0.01	83	—	—	140.4
16M	24,432 (100%)	13	3	0.09	0.00	85	—	—	135.8
20M	14,584 (100%)	20	3	0.02	−0.01	63	61	—	64.1
60M	29,042 (100%)	17	3	0.08	0.00	98	100	−0.20 to −0.33	29.5
90M	29,901 (100%)	20	3	0.11	0.01	100	—	−0.43 to −0.50	15.0
150M	37,500 (98%)	15	3	0.23	−0.03	100	100	−0.47 to −0.64	10.7
300M	45,000 (98%)	15	3	0.22	−0.01	100	100	−0.69 to −0.78	5.0
530M	51,250 (89%)	10	3	0.34	0.01	100	—	−0.65 to −0.71	3.1
750M	27,500 (43%)	10	3	0.37	0.02	100	100	−0.66 to −0.77	2.1
1B [†]	7,500 (11%)	10	5	0.36	−0.01	100	—	−0.75 to −0.83	1.5
1B	69,369 (100%)	25	3	0.35	−0.01	100	—	−0.76 to −0.80	1.5

Table 6: **DataDecide at every size.** The checkpoint used (DataDecide revision step<N>-seed-<seed>; share of the default seed’s training), the recipes whose three seeds share their step-0 weights, which-head agreement (mean over induction, previous-token and retrieval maps; induction only where models have induction heads), leave-one-corpus-out identification of the initialization from the head layout at that checkpoint and at the earliest available one (3–9% of training; chance 33%, 20% with five seeds), the Spearman correlation between unigram corpus distance and same-initialization agreement over corpus pairs (range over roles), and the recipe’s learning rate per batch token. †: five seeds at 11% of training; last row: the verified three-seed crossing at the end of training. Every run starts from its labelled initialization (Appendix A).




Figure 6: **The window tracks cumulative learning.** Agreement of the final layout with the uninterrupted run after noise as large as the weights at 1, 2.5, 5, 10 and 20% of training, against the cumulative learning rate before the intervention (in units of the peak rate), for warm-ups of 300 and 1,000 steps; crosses: branches with a freshly initialized optimizer (300-step warm-up). Means over two initializations and the induction and previous-token maps.

H Gradient Noise

At fixed size we vary the batch (16, 64, 512 sequences) and the learning rate ($3 \cdot 10^{-3}$, 10^{-3} , $3 \cdot 10^{-4}$) and read the layout at step 2,000, after the assignment is fixed (Table 8). Order-only agreement is ordered by the learning rate per batch token, and tripling the learning rate acts like quartering the batch.




Figure 7: **Less gradient noise, more agreement.** (a) Controlled models at step 2,000: which-head agreement between runs that differ only in batch order, and between C4 and scientific-paper runs from one initialization, against the learning rate per batch token. (b) DataDecide: same-initialization agreement against the recipe’s learning rate per batch token.

I Initialization Effects on Benchmarks

Table 9 breaks the benchmark analysis of §6 down by task. The corpus is a significant factor on almost every task and size; the initialization reaches significance in 0–2 of 14 sizes per task, as expected by chance, and its median variance share is zero on every task.

J The Template Habit in Detail

Prompts. Context S = the first sentence of a Para-Conflict substitution conflict; declarative: S + an-

Size	Identified (%)			Weight correlation same init.	SeedPrints z	
	head layout	weights	SeedPrints		same init.	same corpus
4M	56	98	94	0.022	9.5	1.0
6M	77	100	91	0.016	8.7	0.7
8M	65	100	94	0.013	8.7	1.6
10M	81	100	88	0.011	8.4	1.5
14M	83	100	94	0.009	7.8	3.6
16M	85	100	82	0.009	7.1	3.4
20M	63	100	97	0.016	8.1	2.3
60M	98	100	100	0.010	11.9	5.1
90M	100	100	100	0.012	13.3	3.9
150M	100	100	100	0.013	24.0	5.1
300M	100	100	100	0.016	30.1	5.9
530M	100	100	100	0.019	46.7	6.1
750M	100	100	100	0.022	46.4	4.1
1B [†]	100	100	100	0.043	54.7	1.7
1B	100	100	100	0.016	63.8	9.8

Table 7: **Identifying the initialization three ways.** Leave-one-corpus-out accuracy from the head layout, from the correlation of final weights (a fixed sample of 2M attention and MLP weights) and from SeedPrints (Tong et al., 2026); mean weight correlation of same-initialization pairs, and mean SeedPrints statistic of same-initialization and of same-corpus, different-initialization pairs. [†]: five seeds at 11% of training.

Batch	LR	LR per batch token (10^{-6})	other batch order	C4 vs. papers
16	$1 \cdot 10^{-3}$	0.12	0.43	0.01
64	$3 \cdot 10^{-3}$	0.09	0.57	0.07
64	$1 \cdot 10^{-3}$	0.03	0.61	0.21
64	$3 \cdot 10^{-4}$	0.01	0.82	0.19
512	$1 \cdot 10^{-3}$	0.00	0.90	0.22

Table 8: **Gradient noise.** Same-initialization which-head agreement of previous-token maps at step 2,000 between runs that differ only in batch order, and between runs on C4 and on scientific papers.

swer stem; QA: S + Question: q Answer: + stem; variants replace the markers (Table 4).

Items. Facts known to every compared model (the memorized answer is more likely than the counterfactual without context): 1,680 ParaConflict items over six relations, 1,299 PopQA items, and 3,533 (DataDecide) and 3,750 (OLMo 2) NQ-Swap items. Standard errors use the seed-to-seed variance pooled over four 1B recipes with three seeds each.

Across sizes, training and datasets. The habit is present at six of seven sizes from 60M to 1B (Figure 8a, Table 10), appears by 3.6% of pretraining and grows to +2.4 nats (Figure 8b), and replicates on PopQA counterfactuals (+1.91) and on NQ-Swap, where on natural reading-comprehension passages it is strongest after Question: and extends to Q: (Figure 8d). In OLMo 2, mid-training

on a mixture containing FLAN raises it from ≈ 0 to +1.4–2.3 nats in each of three runs, while trust under declarative prompts falls (Figure 8c); the mixture contains other data as well, so this replication shows the habit in a second training stack rather than isolating FLAN. Continued pretraining of a Flan-free 60M model on 90M Flan tokens does not install it ($+0.04 \pm 0.04$).

Public models. Table 11 gives the trust of public base models with and without the question format.

K Artifacts, Compute and Software

Artifacts. We use the DataDecide models (Apache 2.0) and data recipes (ODC-BY), Pythia and PolyPythias (Apache 2.0), OLMo 2 (Apache 2.0), ParaConflict (Apache 2.0), NQ-Swap (MIT) and PopQA as released by their authors, all for research use consistent with their intended use; the corpora are English.

Compute. All experiments ran on NVIDIA A100 80GB GPUs, about 150 GPU-hours in total, including exploratory runs; most of it went to the controlled pretraining runs, while measuring the public checkpoints takes a few minutes per model.

Software. PyTorch 2.8, Hugging Face Transformers 4.57 and Datasets 4.8, NumPy 1.26 and SciPy 1.17; models are loaded in bfloat16 with eager attention to read attention patterns. Benchmark scores are DataDecide’s released OLMES evalua-

Benchmark	Accuracy				Correct-answer likelihood per byte			
	init. sig.	corpus sig.	init. share	corpus share	init. sig.	corpus sig.	init. share	corpus share
ARC-Easy	0/14	14/14	0.00	0.95	0/14	14/14	0.00	0.91
ARC-Challenge	1/14	10/14	0.00	0.49	0/14	14/14	0.00	0.92
BoolQ	0/14	7/14	0.00	0.25	2/14	13/14	0.00	0.55
CSQA	0/14	12/14	0.00	0.55	1/14	14/14	0.01	0.56
HellaSwag	2/14	13/14	0.00	0.60	2/14	14/14	0.00	0.96
MMLU	1/14	13/14	0.00	0.87	1/14	14/14	0.00	0.95
OpenBookQA	1/14	10/14	0.00	0.40	1/14	14/14	0.00	0.85
PIQA	0/14	14/14	0.00	0.56	1/14	14/14	0.00	0.76
Social IQa	0/14	10/14	0.00	0.30	2/14	13/14	0.00	0.48
WinoGrande	1/14	3/14	0.00	0.07	1/14	13/14	0.00	0.79
OLMES average	0/14	10/14	0.00	0.46	1/14	14/14	0.00	0.78

Table 9: **Initialization effects on benchmarks.** For each benchmark, the number of the 14 model sizes at which the initialization or the corpus has a significant main effect ($p < 0.05$, two-way ANOVA over recipes and seeds) and the median unbiased variance share of each factor across sizes.




Figure 8: **The Question: habit across sizes, training and datasets.** Added trust in a counterfactual context (log-probability margin of the counterfactual over the memorized answer, relative to a declarative prompt; nats, $\pm$ SE). (a) Flan minus no-Flan DataDecide models at every size, for the Flan template and two equivalent templates. (b) During pretraining of the 1B models: present at the first checkpoint (3.6%); declarative prompts unaffected. (c) OLMo 2 1B before and after mid-training on a mixture with 16.6% FLAN (three mid-training runs). (d) PopQA counterfactuals and NQ-Swap (Question: vs. Q:).

tions (Magnusson et al., 2025); marker counts in Dolma 1.7 come from the infini-gram index.

L Statistics

Recipe-bootstrap 95% intervals for same- minus different-initialization agreement; unbiased two-way variance components from mean squares with F tests; permutation tests with recipe labels permuted; Mantel tests for pair-level regressions (Mantel, 1967); partial Spearman correlations by rank regression.

Size	Question: <i>q</i> Answer:	Question: <i>q</i>	Answer:	Q: <i>q</i> A:	Query: <i>q</i> Resp.:	other entity
60M	+1.71 (0.58)	+0.73 (0.35)	+1.74 (0.59)	+0.48 (0.51)	+0.52 (0.83)	+2.14 (0.51)
90M	+1.39 (0.37)	+1.23 (0.55)	+1.51 (0.17)	+0.20 (0.44)	+0.40 (0.52)	+1.89 (0.47)
150M	+1.20 (0.57)	+1.09 (0.43)	+1.45 (0.70)	+1.16 (0.62)	+0.85 (0.59)	+1.48 (0.41)
300M	+1.52 (0.44)	+1.60 (0.60)	+1.54 (0.50)	+0.58 (0.48)	+0.19 (0.65)	+1.37 (0.41)
530M	+0.94 (0.79)	+0.96 (0.79)	+0.33 (0.70)	−0.14 (0.72)	+0.08 (0.72)	+1.51 (1.14)
750M	+1.05 (0.59)	+1.72 (0.35)	+1.35 (1.40)	+0.68 (0.50)	+0.40 (0.61)	+2.17 (0.50)
1B	+2.43 (0.79)	+2.04 (0.83)	+1.13 (0.63)	+0.39 (0.89)	+0.52 (0.83)	+3.06 (1.42)

Table 10: **The Question: habit at every size.** Flan minus no-Flan DataDecide models: added trust in the counterfactual context relative to a declarative ending, in nats (SE). Flan’s template carries the effect at every size, and more than the equivalent Q:/A: and Query:/Response: templates at every size (pooled over sizes by 0.93 ± 0.29 and 0.98 ± 0.32 nats); below 1B, Answer: alone often carries much of it.

Model	Instruction data in pretraining	declarative	question	gain
OLMo-1B (July 2024)	Dolma 1.7, with Flan	4.33	8.92	+4.60
Qwen2.5-1.5B	synthetic data from instruct models	−0.47	2.70	+3.18
Pythia-1B	none (the Pile)	5.47	6.89	+1.41
Pythia-1.4B	none (the Pile)	5.41	6.78	+1.38
SmolLM2-1.7B	not documented	2.89	3.93	+1.04
Gemma-2-2B	not documented	3.48	4.51	+1.03
OLMo-2-1B, before mid-training	none			−0.26 to +0.08
OLMo-2-1B, after mid-training	16.6% FLAN			+1.45 to +2.28

Table 11: **The habit in public base models.** Trust in a one-sentence counterfactual context (nats) with a declarative ending and with a Question:/Answer: ending, and the gain from the question format. The model pretrained with Flan trusts a question most.